

arrayForth® 3

User's Manual

Rev 03C for G144A12 chips

Running on saneFORTH™/Win32
and polyFORTH® for GA144

This manual is designed to prepare you for using arrayForth 3 (aF-3) in designing, implementing and testing applications of our chips.

aF-3 is a complete, interactive software development , debugging and installation environment for GreenArrays Chips. It includes an F18 Assembler, example source code including all ROM on each chip, a full software-level simulator for each chip, an Interactive Development Environment for use with real chips, and utilities for creating boot streams and burning them into flash memory. As of aF-3, colorForth is no longer used in this system.

aF-3 is written to run on polyFORTH in G144A12 environments with sufficient resources, and on saneForth for Win32 environments. These versions complement each other and each has a different emphasis of tools, reflecting their differing purposes. The principal purpose of the Win32 environment is cross-compiling for new chips, commission new boards, and simulate at high speeds. The G144A12 environment is intended for interactive development and testing, with F18 and polyFORTH source code in a single base that can be maintained by either system. In most cases, we intend that you will be using an EVB instead of a PC as the principal host for software development.

Although it is configured to support the GreenArrays EVB002 Evaluation Board, it may easily be used to program and debug our chips on the EVB001 in your own designs.

Along with the above tools, including complete source code for the Virtual Machine environments, this release incorporates the source code for our Automated Testing systems as well as that which has been used in taking the characterization measurements reflected in the G144A12 Data Book.

Your satisfaction is very important to us! Please familiarize yourself with our Customer Support web page at <http://www.greenarraychips.com/home/support>. This will lead you to the latest software and documentation as well as resources for solving problems and contact information for obtaining help or information in real time.

Contents

1.	Introduction to this Manual	6
1.1	<i>Related Publications</i>	6
1.2	<i>Status of Data Given</i>	6
1.3	<i>Documentation Conventions</i>	7
1.3.1	Numbers	7
1.3.2	Node coordinates	7
1.3.3	Cardinal directions	7
1.3.4	Register names	7
1.3.5	Bit Numbering	7
1.4	<i>Getting Started</i>	7
2.	Introduction to arrayForth 3	8
2.1	<i>The Common Environment</i>	9
2.1.1	Nature of the Environment	9
2.1.2	Layering of arrayForth	10
2.1.3	Caveats	10
2.1.4	Installation Tips	11
2.2	<i>Environment in saneFORTH (sF)</i>	12
2.2.1	Installation on Win32 Platforms	12
2.2.2	Suggested Usage	12
2.3	<i>Environment in polyFORTH/144 (pF/144)</i>	13
2.3.1	Installation on Evaluation Boards	13
2.3.2	Suggested Usage	13
2.4	<i>Common arrayForth 3 User Vocabulary</i>	14
3.	Mass Storage	15
3.1	<i>Disk Layout</i>	15
3.1.1	Gross Disk Organizations	16
3.2	<i>Tools for Managing Disk</i>	17
3.2.1	Concordance (sF Only)	17
3.2.2	<i>Printing Listings</i>	17
3.3	<i>Configuration Blocks</i>	18
3.4	<i>The Bonus Materials</i>	19
4.	Programming the F18	20
4.1	<i>Object Code</i>	20
4.1.1	EXAMINE - Object Code Auditing	21
4.1.2	Getting Object Bins from colorForth	21
4.1.3	Assembling Stock Code	22
4.2	<i>Assembler Syntax and Semantics</i>	23
4.2.1	Assembler State Variables	23
4.2.2	Assembler Directives for Code Bounding	23
4.2.3	Location Counter	25
4.2.4	Control Structure Directives	27
4.2.5	F18 Opcodes	28
4.2.6	Dictionary Labels	29
4.2.7	Other Useful Words	29
4.3	<i>Module Organization</i>	30

4.3.1	Load Block.....	30
4.3.2	Boot Descriptors.....	30
4.3.3	Residual Paths.....	30
4.3.4	Organization of Larger Projects.....	30
4.4	Methods of Loading Code.....	31
5.	Interactive Testing.....	32
5.1	External IDE.....	32
5.1.1	Terminology.....	32
5.1.2	Using the External IDE.....	32
5.1.3	External IDE Vocabulary.....	33
5.1.4	Advanced External IDE Uses.....	34
5.1.5	Configuring IDE for Target Environments.....	35
5.1.6	Working with Two Chips using External IDE.....	36
5.1.7	Default External IDE Paths.....	37
5.2	Internal IDE.....	38
5.2.1	Terminology.....	38
5.2.2	Using the Internal IDE.....	38
6.	Preparing Boot Streams.....	39
6.1	The Streamer Utility.....	39
6.1.1	Boot Descriptor Language (BDL).....	40
6.1.2	Stream Structure.....	42
6.2	Transmitting Boot Streams.....	44
6.2.1	From saneFORTH.....	44
6.2.2	From pF/144.....	44
6.3	Burning Flash.....	46
6.3.1	Burning Flash from saneFORTH.....	46
6.3.2	Burning flash from pF/144.....	46
6.3.3	Erasing Flash.....	46
6.4	Auditing Streams.....	47
6.4.1	Getting Streams from Flash pF/144 ONLY.....	47
6.4.2	Getting Streams from colorForth.....	47
7.	Simulation Testing with SOFTSIM.....	49
7.1	The Display.....	49
7.1.1	Meaning of the Data Displayed.....	50
7.1.2	"Time" and Cycle Count.....	50
7.1.3	The Overview Section.....	50
7.1.4	The Focus Node Section.....	52
7.2	Loading Code to Simulate.....	52
7.3	Operating the Simulator.....	53
7.3.1	About Breakpoints.....	53
7.3.2	Initialization.....	53
7.4	Testbeds.....	54
7.5	Interactive Testing with Softsim.....	54
8.	Practical Example.....	55
8.1	Selecting resources.....	55
8.2	Wiring.....	55
8.3	Writing the code.....	56

8.4	<i>The IDE script</i>	58
8.5	<i>Output Observations</i>	59
8.6	<i>Further Study</i>	60
9.	Commissioning a New G144A12 Board	61
9.1	<i>G144A12 Standalone, No Flash or SRAM</i>	61
9.1.1	sF as Host	61
9.1.2	pF/144 as Host	61
9.2	<i>G144A12 Embedded, Flash Only</i>	61
9.2.1	sF as Host	61
9.2.2	pF/144 as Host	61
9.3	<i>G144A12 polyFORTH Capable</i>	61
9.3.1	sF as Host	61
9.3.2	pF/144 as Host	62
10.	Reference Material	63
10.1	<i>F18A Code Library</i>	63
10.1.1	Perfect Wire (any node).....	63
10.1.2	64-word Delay Line (any node)	63
10.1.3	Synchronous Port Bridge (300 and others).....	64
10.1.4	Ganglia Mark 1 (any node)	64
10.1.5	Ganglia Mark 2 (any node).....	64
10.1.6	SRAM Mark 1 Cluster (7, 8, 9, 107, more)	64
10.1.7	Snorkel Mark 1 (any SRAM client node)	64
10.1.8	IDE Components (708, 300 + any nodes)	64
10.1.9	Boot Stream Components.....	65
10.1.10	polyFORTH Cluster (specific nodes)	65
10.1.11	Ethernet NIC Cluster (specific nodes)	66
10.1.12	ATS Mark 1 (Serial IDE based)	66
10.1.13	ATS Mark 2 (On-chip based)	66
10.1.14	eForth (specific nodes)	66
10.1.15	Documented Examples	66
10.1.16	Additional Test Code	66
10.2	<i>Bin Assignments for Rev 03b4</i>	67
10.3	<i>Examples to hang onto</i>	70
11.	Maintenance	71
11.1	<i>saneFORTH Maintenance</i>	71
11.1.1	Bootstrapping	71
11.1.2	Nucleus Maintenance	71
11.2	<i>polyFORTH/144 Maintenance</i>	71
12.	Appendix: Microsoft Windows® Platform	72
12.1	<i>Windows arrayForth Requirements</i>	72
12.2	<i>Installation</i>	72
12.2.1	Identifying and Configuring COM Ports	73
12.2.2	Windows 8 and 10 Installations	76
12.3	<i>Running arrayForth 3</i>	76
12.4	<i>Installing arrayForth 3 on an EVB002 Board</i>	76
12.5	<i>Installing arrayForth 3 on an EVB001 Board</i>	76
12.6	<i>Customizing for a Project</i>	77

12.6.1	Managing Multiple Major Projects	77
13.	Appendix: Apple Mac® Platform with Windows	79
13.1	Mac Requirements	79
13.2	Installation	79
13.3	Running arrayForth.....	79
14.	Appendix: unix Platform including Mac OS X	79
14.1	Installation	79
14.1.1	Identifying and Configuring COM Ports	80
14.2	Running arrayForth.....	81
15.	Data Book Revision History	83

1. Introduction to this Manual

This is the primary reference manual for the arrayForth programming environment. It should be read and understood in its entirety. In the interest of avoiding needless and often confusing redundancy, it is designed to be used in combination with other documents.

1.1 Related Publications

- **DB005 *polyFORTH Reference Manual*** is the foundation for understanding the polyFORTH model, tools, and development methods.
- **DB006 *G144A12 polyFORTH Supplement*** assumes you understand the material in the above *Reference*, documenting only implementation-specific details. We recommend that you familiarize yourself with the *Reference* before studying this manual.
- **SFW32PR *Programmer's Reference, saneFORTH for Win32 x86 Platforms*** fulfills the role of the above Supplement for this system (ATHENA Programming, distributed by permission.)
- **DB001 *F18A Technology Reference*** serves as the Programmer's Reference for the F18 computers and I/O architecture.
- **DB002 *G144A12 Chip Reference*** documents the configuration of this specific chip. Both DB001 and DB002 should be understood before you begin programming F18 code.
- **DB003/DB014 *Evaluation Board Reference Manual*** for EVB001 and EVB002, respectively, contain information with which every user of these boards should be familiar. The polyFORTH release may be adapted for other hardware configurations but is shipped with configuration settings suitable for running on the EVB002 Evaluation Board.
- **Application Notes** exist, or are planned, for many software modules that are either intrinsic to polyFORTH (such as the SRAM Control Cluster, Snorkel, and Ganglia) or optional (such as the Ethernet NIC and networking packages.) These and other reference materials for our chips, such as the boot protocols supported by ROM code in their boot nodes, may be found on our website at <http://www.greenarraychips.com>. The most comprehensive list of these will be found in the *Index of Downloads* tab. It is always advisable to ensure that you are using the latest documents before starting work.

1.2 Status of Data Given

The data given herein are *released* and *supported*. However, if a section heading is **highlighted in yellow**, that heading and all text and subheadings subordinate to it are not yet converted/updated to reflect current hardware or software. **The same is true of single text passages** unless otherwise specifically noted otherwise near the beginning of the section in question. The subject applications are under continual development; thus the software and its documentation may be revised at any time.

Supplemental status information is available on our website at <http://www.greenarraychips.com/home/support>. This page is updated frequently and we recommend that you visit it regularly. You may also wish to subscribe to the RSS feed on our technical blog; changes in software or documentation are announced there.

1.3 Documentation Conventions

1.3.1 Numbers

Numbers are written in decimal unless otherwise indicated. Hexadecimal values are indicated by explicitly writing "hex" or by preceding the number with the lowercase letter "x". This is true in source code as well.

1.3.2 Node coordinates

Each GreenArrays chip is a rectangular array of *nodes*, each of which is an F18 computer. By convention these arrays are represented as seen from the top of the silicon die, which is normally the top of the chip package, oriented such that pin 1 is in the upper left corner. Within the array, each node is identified by a three or four digit number denoting its Cartesian coordinates within the array as *yx*x or *yyxx* with the lower left corner node always being designated as node 000. This convention is expanded to *cyxx* on multi-chip boards such as the EVB001 and 2, in which case *c* is chip number (0 for host, 1 for target, etc). All functions herein accepting *yx*x notation also recognize *cyxx* appropriately.

1.3.3 Cardinal directions

When it's necessary to refer to directions in the chip geometry without reference to node-specific port directions such as left or up with, we use cardinal compass directions such as North for positive Y axis and East for positive X. Cardinal directions are also defined in the Assembler as **north south east west** for use in code, and in SOFTSIM and STREAMER as **NORTH SOUTH EAST WEST** primarily for use in Boot Descriptor Language (BDL). The cardinal directions are supported in these tools to facilitate assigning code to geometrically dissimilar nodes.

1.3.4 Register names

Register names in prose may be used with or without the word "register" and are usually shown in a bold font and capitalized where necessary to avoid ambiguity, such as for example the registers **T S R I A B** and **IO** or **io**.

1.3.5 Bit Numbering

Binary numbers are represented as a horizontal row of bits, numbered consecutively right to left in ascending significance with the least significant bit numbered zero. Thus bit *n* has the binary value 2^n . The notation P9 means bit 9 of register **P**, whose binary value is x200, and T17 means the sign (high order) bit of 18-bit register **T**.

1.4 Getting Started

If you are a new user of Forth based systems, particularly of those respecting the polyFORTH model, you can save yourself a great deal of frustration by familiarizing yourself with the environment and tools prior to using them. You will be speaking to a text based Interpreter that is vastly more powerful than any other "shell" you may have used. You may extend the environment including the Forth language and the tools. This entire system is vastly simpler than anything else you may be familiar with; for example, many of us have proven that one person can create and maintain a Forth system, either hosted on some other operating system or standing alone as its own.

The first three documents in the Related Publications section above are the key reading for the two Forth systems included in arrayForth-3. You might browse the next chapter first for context, but these three documents will inform you of the basic tools and of their use. You might not wish to read every word of all three initially but at the minimum familiarize yourself with their contents and with basic tools such as disk management and text editing. It's best to start with DB005 because both DB006 and SFW32PR assume you are familiar with that Reference Manual.

You will find that this simple, consistent programmer interface empowers a programmer greatly with no burden of having to read shelf-feet of library manuals. Of course, to achieve this simplicity, it does not have, nor do you need, many of the "tools" required by conventional systems. If all of this does not suffice, please contact hotline@greenarraychips.com for further assistance.

2. Introduction to arrayForth 3

aF-3 is a complete, interactive software development, debugging and installation environment for GreenArrays Chips. It includes an F18 Assembler, example source code including all ROM on each chip, a full software-level simulator for each chip, an Interactive Development Environment for use with real chips, and utilities for creating boot streams and burning them into flash memory, as well as auditing changes to both source and binary object code. As of aF-3, colorForth is no longer used in this system.

aF-3 is written to run on polyFORTH in G144A12 environments (polyFORTH or pF/144 herein) with sufficient resources, and on saneFORTH for Win32 environments (saneFORTH or sF herein). These versions complement each other and each has a different emphasis on selection of tools from the common set, reflecting their differing purposes. The principal purposes of the Win32 environment are cross-compiling for new chips, commissioning new boards, and simulating at high speeds. The G144A12 environment is intended for interactive development and testing of nodes on its own chip, or on other chips by internal or external means, with F18 and polyFORTH source code in a single base that can be maintained by either system. In most cases, we intend that you will be using an EVB instead of a PC as the principal host for software development. Here is the checklist of features supported by each system:

Capability Supported	sF	pF/144
F18 native code Assembler	YES	YES
Source and object code auditing	YES	YES
Concordance and fast FIND in source code	YES	---
Support for 2-chip targets	YES	YES
External IDE	YES	tbd
Internal IDE	---	tbd
Boot stream generation	YES	YES
Boot stream auditing	YES	YES
Internal boot stream delivery (Node 207)	---	YES
Serial boot stream delivery	YES	tbd
Flash boot stream burning in target	tbd	tbd
Flash boot stream burning internal to own system	---	YES
Target Compiler for pF/144 Nucleus	---	YES
Chip test via External IDE	YES	tbd
External SRAM assembly test via external IDE	YES	tbd
ATS Chip Testing	YES	tbd
Software Simulator	YES	---

Wherever practical we have kept the human interface to the programming tools consistent between the sF and pF/144 environments. In some cases this is inconvenient; for example, even though sF is a 32-bit system, values that must be double precision on pF/144 must also be double in sF. Experience has taught that this is a case in which a little inconvenience buys much portability between the environments and helps avoid nasty programmer traps.

Where there are differences between the two systems, we note them with the markers **sF ONLY** or **pF/144 ONLY**.

Although it is configured to run on, and support, the GreenArrays EVB002 Evaluation Board, aF-3 may easily be used to program and debug our chips on the EVB001 or in your own designs.

2.1 The Common Environment

Both systems use a simple, consistent human interface in which your communications, command and control are normally done in words, employing user-extensible language and phrasing. This has always been the intrinsic human interface provided across systems and platforms by Forth, and is commonly implemented as a 24x80 (we prefer 25x80) ASCII text-based display with standard keyboard input.

We create our systems for the purposes of efficiently accomplishing useful work and of writing and testing reliable code. When working on a variety of platforms and environments, and when human psychophysiological capabilities differ, a simple, consistent human interface that can be adapted as necessary has proven to be a great solution. For example, we have worked with at least one successful Forth programmer who was completely blind; supporting such a customer our way was proven to require almost nothing special from us.

During the 1990s, we had the choice of adopting emerging commercial human interface artifacts such as GUIs with pointing devices. Although we often implemented and used GUIs in appropriate applications long before that, we concluded, correctly, that for the basic tasks of programmer self-expression these artifacts contributed few useful advantages while introducing significant liabilities. Therefore, when programming systems whose human interfaces are provided by tools such as X or as Microsoft Windows, we simply draw our own 25x80 terminals. When communicating with one of our systems over the Internet, we use Telnet emulating (you guessed it) a 25x80 terminal. And, by interfacing with an appropriate terminal device, that blind man can still work effectively with our systems. Meanwhile, those fully embracing GUIs for basic programming tasks have had the privilege of spending a great deal of time learning and adapting to a long series of human interface changes that were made largely for capricious, stylistic reasons as contrasted with rigorous human engineering considerations; their hands and eyes are busy with overhead activity, and their blind users need to find something other than programming to do for a living.

Incidentally, GreenArrays no longer uses or supports the experimental colorForth programming environment, because its advantages proved, in practice, insufficient to overcome its disadvantages.

2.1.1 Nature of the Environment

Our systems use a standard keyboard supporting printable 8-bit ASCII characters along with SPACE, ENTER (Carriage Return) and BACKSPACE keys. When communicating with the Forth Interpreter, one may type and edit with backspace at will; nothing actually happens until ENTER is pressed, at which time the phrase is parsed and the specified activities are performed. The upper and lower-case versions of a given character are treated as distinct from each other. The basic system requires only on the 7-bit characters listed above (applications may use any additional capabilities of the keyboard environment, of course.)

Unless you have experience with polyFORTH systems, it's worth your effort to study the polyFORTH Reference Manual (DB005). This is worthwhile because it covers the common ground across more computer and operating system types than can be said of, we think, any other programming system / human interface in existence.

Based on that prior experience and/or study, you should then read the Supplements to that reference manual. DB006 is the Supplement for the G144A12 implementation. SFW32PR, Programmer's Reference, fulfills that purpose for saneFORTH running on the x86 platform.

In the basic polyFORTH model, the system boots with a precompiled/assembled nucleus capable of interacting with basic terminal and mass storage and of editing and compiling/assembling more code to extend itself. Normally the extension process is accomplished by loading block 9. This is done by typing the word HI that does two things: HOME, which performs any initialization required for mass storage access, and 9 LOAD to interpret block 9. If you will not be using the system at the default offset (zero), you might say something like 20 DRIVE HI .

2.1.2 Layering of arrayForth

A principal design goal of arrayForth is that in both environments there shall be a resident F18 assembler with the full ability to build instructions and reference labels defined in object code bins. Our experience has been that forcing programmers to write F18 instructions, or addresses in F18 routines, as hexadecimal literals in their application code is not merely an impediment to documentation. It has also, even in software released by GreenArrays, proven to embed un-auditable future bugs in that software. Hence, with arrayForth, we return to fully symbolic representation of code. This has, after all, proven over the last half century to be a very effective way of protecting oneself from many programmer traps. The problem with writing a simple number like "5" is that without context the meaning of that number is not evident, and that meaning can only be determined by holographic recollection or sometimes exhaustive study of the code. Most of us learned early in life that our ability at recollection is anything but holographic.

However, in a 16-bit machine a single compiled block of high level source can consume a substantial fraction of the available address space, and having decided to make that block resident the scope of what can be done with what is left of the machine is diminished. Many of us have not lived in such a resource-constrained world for forty years or so, and we've found the habits that requires have faded. Because environmental conventions are painful to change after one has started building things respecting them, we have tried to make the best decisions we could in factoring and layering arrayForth. Of course, in any undertaking of that sort one will inevitably make mistakes. The conventions underlying arrayForth can, and will, be changed if we find serious design flaws. Such changes will be documented.

2.1.2.1 AFORTH Capsule

This is a resident CAPSULE. It is accessed by interpreting its name, AFORTH. (Immediately after HI the word AFORTH means to build that capsule by loading its resident code, and redefining the word AFORTH as a CAPSULE. Thus the code does not exist until you need it.)

You must be running in the AFORTH capsule to have access to any of its capabilities. The code that is physically resident on both platforms is the F18 Assembler, configuration data, and at least the names of all the other tools provided by arrayForth. At present, none of these tools are resident on either system. The reason for this is that we cannot countenance making anything beyond the F18 Assembler resident in the pF/144 environment, and we believe that requiring different steps and procedures on each platform to accomplish things supported on both would be a *very bad thing*. The vocabulary provided with AFORTH on each system is specifically documented in later sections.

2.1.3 Caveats

Along with many similarities, there are significant environmental differences between the two platforms. One immediately visible difference is that the size of the basic integer and address data in sF/x86 is 32 bits while on pF/144 it is 16 bits. Another is that all addresses in sF/x86 are byte addresses while on pF/144 there are both byte and cell addresses. Further, and most significantly for programming purposes, the resources of the GA144 hardware are limited (basic addressing covers 64 kWords or 64 kBytes in the EVB002, with a maximum extended memory of 2 Mbytes, onboard mass storage of 16 Mbytes, and the relatively low performance of a virtual machine implementation; whereas the sF environment has effectively gigabytes of memory, orders of magnitude more mass storage, and the performance of a dedicated x86 processor for running the high level language.

However, in the context of arrayForth, you will not be using saneFORTH to write x86 code but rather to run our tools for programming the GA144 chips. Therefore, from your perspective, the main differences will lie in resources and performance, with the resulting differences in which tools we have chosen to implement for each platform.

Another thing which comes as a surprise to newcomers first working with Forth systems is that, in the basic Interpreter context, *every word and function compiled in the system is accessible from your keyboard*. Indiscriminate typing of words whose meanings you do not understand and are not employing correctly can lead to surprising results! Moreover, again without specific contextual modifications to prevent such, you have full power to crash the system if you wish. On machines without memory protection hardware, you merely need say something like **0 10000 ERASE** and in most cases you won't get an answer to that statement. So, before hitting ENTER, make sure what you typed is what you meant to say!

2.1.4 Installation Tips

Here are a few suggestions that might make your life more pleasant when using arrayForth 3.

2.1.4.1 Console Shortcuts

The installer tool we use does not give us any control over the characteristics of the shortcuts it creates during the installation process on the x86 platform. Unfortunately, key parameters for console windows are stored in these shortcuts. The defaults in the installer-created shortcuts are unacceptable.

In the directory structure created for arrayForth, you will find several shortcuts in the default location

C:\GreenArrays\EVB002\sF

Those relevant to you are called **arrayForth 3 PANIC**, **arrayForth 3 G144A12**, and **arrayForth 3 SOFTSIM**. These are configured more appropriately in terms of console window configuration and command line content. It will require some acquaintance with the Windows platforms to make these a convenient part of your desktop; if you have difficulty with that, contact our customer support.

We have used the white characters on moderately dark blue background for years because, of the default colors available on these consoles (which derive from the original EGA/VGA colors on PCs), this combination has worked well in terms of human engineering; good for eye fatigue, no false color [stereopsis](#) and so on. You might start out with our conventions before changing them.

There are three sF executables provided with aF3: sF5b2-panic.exe, sF5b2-af3.exe, and sF5b2-glow.exe. The first two are "small" memory configurations while the third has much more space available. The first of the three simply comes up and says hi while the latter two automatically load block 534, which by default loads block 9 (equivalent to saying HI) and then interprets whatever is in the command line following the name of the executable file. This additional interpretive content is editable in each shortcut and may be customized as you see fit to save you steps, at the price of maintaining situational awareness of which environment you are entering.

If you need to maintain a number of differing environments but do not wish to edit the basic system source code for each of them (or maintain several copies of the entire system), you may do a great deal of customizing from the shortcut by using INCLUDE to interpret a text file from the shortcut's command line after the 9 LOAD has completed. An example of this is provided in the aF-3 distribution; see [12.6 Customizing for a Project](#) below.

2.1.4.2 Setting up your EVB002 Flash

As part of testing an EVB002, we copy blocks 0-4800 of the pF/144 serial disk onto its SPI flash and use that to configure and test the Ethernet interface. By the time you receive your EVB002 kit, the system we installed on it is probably out of date. If you plan to use the SPI flash as your main mass storage for work with the board, you will probably want to copy more down there. At minimum you should copy the current source at 0 DRIVE, the reference object code at 8 DRIVE, and the partial section of output object bins at 12 DRIVE (although you may produce the latter using the phrase

EXAMINE LOAD -ALL STOCK LOAD

Simplest is just to copy everything:

BULK LOAD 0 DRIVE 24000 0 16256 BLOCKS

It's your choice whether to use 4 DRIVE on your flash, or 20 DRIVE on your serial disk, for working backup. If you choose to do the latter, then the 4800 blocks starting at 4 DRIVE on the flash are completely uncommitted. As the system is shipped starting with the versions released after July 2019, 4 DRIVE is by default configured as a "Project" file to simplify keeping your own code separated from the system.

2.2 Environment in saneFORTH (sF)

On the saneFORTH/Win32 platform, aF-3 is organized in a CAPSULE. Some of the environment is resident in the capsule, such as the Assembler and state variables for the External IDE. Other parts are overlays loaded as needed. The capsule may be loaded by block 9 or it may be loaded when needed. One word does it all:

AFORTH This word, defined in **GOLD** by block 9, loads the aF-3 environment when it's first invoked, defining a resident capsule **AFORTH** whose name is also defined in **GOLD** thus overloading the prior word which created the environment. Thus the first and any subsequent uses of the word **AFORTH** will yield the same environment, except that state variables are persistent until the system is next booted.

The resident **AFORTH** environment includes all the words and functionality documented in 4.2 Assembler Syntax and Semantics, below, as well as those in 2.4 Common arrayForth 3 User Vocabulary.

Within the **AFORTH** environment the following high level things may be done:

HELP displays reminders for these functions.

CONFIG LIST shows the main configuration block for arrayForth.

HOST TARGET or **BRIDGE LOAD** loads the External IDE for indicated mode. See 5.1. External IDE below.

EXAMINE LOAD loads tools for examining object bins and comparing them.

SOFTSIM LOAD loads the 2-chip instruction level simulator for testing and visualizing code execution.

STREAMER LOAD loads utility for building port executable streams.

WHO identifies which system (sF/x86 or pF/144) you're talking to. Replies "sF on x86" in this system.

SERIAL LOAD loads terminal emulator for talking to polyFORTH/144 (or to eForth).

SELFTEST (*n*) Runs chip tests appropriate for EVB001/2 Host chip using COM port *n* . **sF ONLY**

AUTOTEST (*n*) Runs chip tests appropriate for EVB001/2 Target chip using COM port *n* . **sF ONLY**

RAMTEST (*n*) Runs external SRAM tests on EVB001/2 Host chip using COM port *n* . **sF ONLY**

' ' (two consecutive tics) runs production chip testing on a TB001 board via COM port A-COM. **sF ONLY**

CONC refreshes the Concordance data base by re-parsing all of the source code. **sF ONLY**

EVB (-*n*) returns start block of pF/144 serial disk relative to current OFFSET in sF. **sF ONLY**

2.2.1 Installation on Win32 Platforms

Procedures and tips for installing and managing aF-3 on several common Win32 platforms may be found in the Appendices of this document, supplemented by FAQ items on our website.

2.2.2 Suggested Usage

The sF environment is useful for all modes of host-target development and debugging with a suitable umbilical connection. It is essential for initial commissioning of any GA144 hardware, such as an EVB, when a commissioned pF/144 system is not available. You may prefer it when the higher speed of host operations is more important than the higher bandwidth and lower latency umbilical connections achievable with pF/144 hosts, and when you do not intend to use pF/144 at all. Some utilities, such as SOFTSIM, are only implemented in this system.

2.3 Environment in polyFORTH/144 (pF/144)

On the polyFORTH/144 platform, aF-3 is organized in a CAPSULE. Some of the environment is resident in the capsule, such as the Assembler and state variables for the External IDE. Other parts are overlays loaded as needed. The capsule may be loaded by block 9 or it may be loaded when needed. One word does it all:

AFORTH This word, defined in **GOLD** by block 9, loads the aF-3 environment when it's first invoked, defining a resident capsule **AFORTH** whose name is also defined in **GOLD** thus overloading the prior word which created the environment. Thus the first and any subsequent uses of the word **AFORTH** will yield the same environment, except that state variables are persistent until the system is next booted.

The resident **AFORTH** environment includes all the words and functionality documented in 4.2 Assembler Syntax and Semantics, below, as well as those in 2.4 Common arrayForth 3 User Vocabulary.

Within the **AFORTH** environment the following high level things may be done:

HELP displays reminders for these functions.

CONFIG LIST shows the main configuration block for arrayForth.

HOST TARGET or **BRIDGE LOAD** loads the External IDE for indicated mode. See 5.1. External IDE below.

EXAMINE LOAD loads tools for examining object bins and comparing them.

STREAMER LOAD loads utility for building port executable streams.

WHO identifies which system (sF/x86 or pF/144) you're talking to. Replies "pF on GA144" in this system.

BULK LOAD loads a special version of **DISKING** that is much faster at writing flash. Use with care.

2.3.1 Installation on Evaluation Boards

Please see DB006, *G144A12 polyFORTH Supplement*, or the internal GreenArrays commissioning document PCBCOM for installation procedures.

2.3.2 Suggested Usage

The pF/144 environment is useful for all modes of host-target development and debugging with a suitable umbilical connection. In addition, only the pF/144 system permits intimate F18 code development and testing on the development system itself; for example, one can add micro-instructions to extend the polyFORTH Virtual Machine in vivo and even to alter and test them interactively without rebooting.

This platform, like sF, may be used for initial commissioning of any GA144 hardware, such as another EVB. You may prefer it when the higher bandwidth, lower latency umbilical connections achievable with pF/144 hosts are more important than is the higher speed of host operations in sF. Certain utilities, such as the pF/144 target compiler, are only implemented in this system.

2.4 Common arrayForth 3 User Vocabulary

This section lists resident words (those accessible after **EMPTY** or after naming **AFORTH** on the sF system) that are published for general use in high level programming. These are all high level FORTH functions, not F18 functions; for F18 Assembler syntax see section 4.2 below.

Chip Configuration Words

nnx (-n) node columns/chip.
nnny (-n) node rows/chip.
nns (-n) number of nodes/chip.
nnc (-n) number of object bins.
nn-b (nn-n) convert ccyyxx to linear node (or bin) number.
b-nn (n-nn) convert linear node number to ccyyxx form.

AFORTH Config Words

0bin (-n) starting absolute block of the array of object bins. Size of this region is **nnc** blocks.
CFORG (-n) starting absolute block of a reference array of object bins.

Bin Manipulation Words

)BIN (nn-a) return start address of given bin in block buffer.
@BIN (nn) retrieve given bin to bMEM array.

Named Blocks

STOCK (-n) assembles all the standard object bins that are supported by GreenArrays.
SRAM (-n) assembles the bins for the SRAM cluster. SRAM 1+ is BDL for the cluster. SRAM 2+ defines residual paths for boot nodes after SRAM is installed.
PFVM (-n) assembles the bins for the polyFORTH virtual machine. PFVM 1+ is BDL for the full virtual machine environment including ganglia. PFVM 3+ is an example of 1-chip pF boot stream for flash.
ENIC (-n) assembles the bins for the Ethernet cluster. ENIC 1+ is BDL for the Ethernet cluster.
SRAMIF (-n) is loaded at, normally, an org of x39 to provide the interface routines needed by an SRAM client.

Persistent Parameters

A-COM A-BPS are variables for the "A" (host port). **sF ONLY**.
C-COM C-BPS are variables for the "C" (target port). **sF ONLY**.
U-COM U-BPS are variables for the port currently in use by the external IDE. **sF ONLY**.
DH (-dh) returns currently open IDE file handle or -1. **sF ONLY**.

Unnamed Load Blocks

1662 Serial boot of 2-chip target system for polyFORTH, with or without Ether NIC.
1665 Serial boot of 2-chip target system for polyFORTH, with or without Ether NIC.
1668 makes 2-chip Flash boot stream for polyFORTH, with or without Ether NIC.
1671 makes 1-chip Flash boot stream for polyFORTH, with or without Ether NIC.
1581 generates example stream to locate one node internally using the snorkel through node 207.
1582 generates a similar example for the second chip assuming bridge has been built.
1583 generates another example of using the snorkel to add the Ethernet NIC to a running polyFORTH system. Requires editing block 9 and rebooting before ETHER LOAD..
1584 generates an example serial boot to wiggle a pin.

3. Mass Storage

3.1 Disk Layout

Both systems organize source in 4800 1k-byte blocks, subdivided into 2400 of source followed by 2400 of associated shadows, and further subdivided into 60-block index pages with the following organization as shipped (the green highlighted pages are redacted in systems configured for public release and available for application use):

Block	sF/Win32 (aF-3 and GLOW)	pF/144
0	9 LOAD options, tools, utilities	<--- Same
60	x86 Utilities	pF/144 Nucleus Source
120	x86 Extensions, tools, tests	Extensions, tests, documentation
180	x86 Nucleus Source	Novix arithmetic for conversion
240	x86 Target Compiler & build components	Benchmarks, memory tests, utilities
300		
360	x86 Arithmetic	
420	x86 Data base support	
480	x86 utilities, benchmarks, term emulator	
540	x86 PE (32-bit) EXE generation	Ethernet and TCP/IP Networking Support Partially converted from x86 to pF/144 Eventually, much of this space will be freed for other use.
600	x86 Win32 APIs	
660	x86 Win32 Window Management	
720		
780	F18 Nucleus Source for x86 Targeting	
840	F18 Target Compiler for x86	
900	GLOW	These 600 blocks are reserved for customer code on both systems. Alternatively, user code may be placed in the "Project" area starting at relative block 4800. You may find that this makes it easier to accept system updates.
960	GLOW	
1020	GLOW	
1080	GLOW 180nm IC	
1140	GLOW 28nm HPP	
1200	reserved for GLOW	
1260	reserved for GLOW (Temp misc tools)	
1320	reserved for GLOW (Temp Win32 iF)	
1380	reserved for GLOW (Temp Win32 iF)	
1440	x86 Cryptography	
1500	x86 Trace, dumps, concordance	
1560	F18 Assembler & tools	This code implements the arrayForth 3 development environment on both systems. These 840 blocks are allocated the same on both systems, although not all code works on both, and that which does work on both will typically have differences in implementation. DO NOT put your own code here, and be careful to audit any changes you make. This code is subject to change and reorganization in each release.
1620	IDE and streams	
1680	Softsim	
1740	JTAG Studies	
1800	ATS EXATRON support	
1860	ATS Creeper Runners (External)	
1920	ATS Creepers	
1980	pF VM and Ethernet NIC (Temporary)	
2040	Misc Test code	
2100		
2160		
2220	G144A12 ROM and F18 system pieces	
2280	Stock F18 code	
2340	Target output (x86 and F18)	Copy of some 35-block boot stream.
4800	PROJECT 2340/2340 (GLOW or app code)	PROJECT, Backup or other use as you wish.

3.1.1 Gross Disk Organizations

On a large scale, each operating system organizes its disk suitably to match its resources.

3.1.1.1 saneFORTH Disk Layout

When using saneFORTH, mass storage is mapped as is shown in block 8 and its shadow. This is currently as follows, with the 4800 blocks of 0 DRIVE, and with 40, 44, and 52 DRIVE, write enabled (+WRT, see SFW32PR manual):

DRIVE	Block	Size	Content
0	0	4860	Working system as diagrammed above. Includes a dummy index page at 4740 so that the Project file may be used in both sF and pF environments.
4	4800	4740	Project file. The last index page (4740) does not exist.
8	9600	4800	Base image for this system; ultimate backup.
12	14400	4800	Another image, typically the previous base.
16	19200	4800	--- not assigned ---
20	24000	9600	Backups of working system and Project file.
40	48000	24000	Serial disk for pF use. 0 DRIVE EVB OFFSET +! to see this perspective in sF.

DISKING is set by default to compare 0 with 24000 for a scope of 9540 blocks.

3.1.1.2 polyFORTH Flash and Serial Disk Layouts

When using polyFORTH, mass storage is mapped as is shown in block 11. This is currently as follows; note that in the standard nucleus, flash is mapped at 0, serial disk at 24000, and flash boot at 48000.

DRIVE	Block	Size	Content
0	0	4860	Working system as diagrammed above. The index page at 4740 is fully usable.
4	4800	4740	Project file. The last index page (4740) should not be used if the same file is being used in the saneFORTH environment. Could also be backup of working system if you don't mind coping with future reorganization of the working system that may force you to move your own code, or any other use you like
8	9600	4800	Reference (colorForth) object images. Other use if you wish.
12	14400	1856	Object output bins. 1856 usable! Must be used for this purpose if you intend to assemble F18 code in polyFORTH.
---	16526	7744	Void area.
20	24000	4800	Working system. Serial backup if using flash as primary.
24	28800	4800	Project file. Serial backup if using flash as primary.
28	33600	4800	OBJ-REF reference object images. Serial backup if using flash as primary.
32	38400	4800	OBJ-AF3 object output bins. Serial backup if using flash as primary.
36	43200	4800	Base image for this system; ultimate backup.
40	48000	128	Flash boot (absolute blocks [0..128])

DISKING is set by default to compare 0 with 24000 for a scope of 9540 blocks, optimized for use in flash. As noted above, serial disk at 20, 24 and 32 DRIVE is write enabled. 28 DRIVE may be enabled by 48 52 +WRT or use of COMMIT in sF; 36 DRIVE may be enabled by 56 60 +WRT in sF.

You are of course free to reorganize the polyFORTH disk as your application requires. We will eventually reduce the mass storage allocation for object images, once everything brought over from the colorForth environment has been verified to reproduce object exactly.

3.2 Tools for Managing Disk

Source code is maintained using the standard polyFORTH character and line EDITOR . Large scale examination is supported by the word set `ax nx bx qx` . Reconciliation and management of blocks is done with the standard polyFORTH `DISKING` utility set by default to cover two 4800-block source chunks, and the special BULK utility ([pF/144 ONLY](#)) for large scale writing to the flash.

3.2.1 Concordance (sF Only)

The concordance utility scans the source in specified ranges of blocks, building a database containing a comprehensive source concordance of all words found in those blocks. The utility ignores a short list of ubiquitous words such as colon and semicolon, but otherwise basically parses space delimited strings and sorts them. When the librarian code is loaded, the following resident functions are added, and the EDITOR block listing is enhanced to show the results of the current search. Current search results are not instantiated per user; there is no restriction against concurrent use of the librarian by multiple terminals, but the result set in effect for *all* terminals will always be the most recent produced by *any* terminal.

CONC Runs utility to re-parse source and rebuild the data base. Takes a handful of seconds.

LIB Displays Librarian help screen.

FIND (`_`) Composes a result set of all references to the given word or string.

NEAR (`_`) Composes result set of all words or strings beginning with the given string.

HIT (`_`) Composes result set of all words or strings that would collide with the given string in a 3 character plus length dictionary.

NN and **BB** move forward and backward within the blocks containing the result set. The current block number must be within `[0..38400[` relative to `OFFSET` .

This tool is invaluable when researching changes or corrections for large applications. The data base is not updated automatically as a result of editing but only when you do so manually. For definitive research the advantages of holographically accessing commented and conditionally compiled code, overlays, nonresident utilities, comments, and arbitrary strings such as block numbers in `LOAD` statements, are of overriding importance; hence we use this method of parsing as contrasted with other possibilities.

3.2.2 Printing Listings



3.3 Configuration Blocks

We will now discuss several configuration blocks of particular interest.

Blocks [0..12] The first twelve blocks are reserved for the native system's colorForth kernel and boot.

Blocks [12..18] These six blocks hold the binaries for the fonts used on the colorForth display.

Block 18 This load block defines the configuration of the colorForth system and application. It is loaded on cold or **warm** boot, compiling various extensions to colorForth, giving names to utilities, and loading blocks 144, 202 and 204 to complete the definition of the arrayForth environment. *The three magenta variables at the start of this block (**ns nblk nc**) must never be moved. Use great caution if altering block 18 or anything it loads since if this sequence is aborted early enough the editor is not available to fix the problem.* The yellow word **qwerty** in this block configures the system for standard keyboard layout and semantics; the yellow word **seeb** immediately following **qwerty** configures the system to display blue words. These are the recommended settings.

Block 144 This load block extends colorForth to support the arrayForth tools with global variables, extra functions, names of further utilities, resident capability to generate png files of screen graphics, and the principal chip configuration parameters from the three blocks 190, 192, and 194.

Block 202 This block defines application tools and is maintained by GreenArrays. When you have identified the COM port numbers corresponding to the three USB ports on the Eval Board, you will need to edit the values for **a-com** and **c-com** to be the COM port numbers for USB ports A and C, respectively. The baud rates for these ports are specified here as well.

Block 204 This block is yours, for any definitions you wish to make global so that they survive empty. Before adding a word to this block, make sure it does not redefine anything else that lies under the empty! Use **def** and **' (tick)** to verify no redefinitions.

Block 200 This is also yours, for compiling F18 code as described in section **Error! Reference source not found.**

Block 148 is the load block for softsim. At present it's necessary to alter this block to control what code is loaded and executed in each node.

By reading the system load blocks starting with 18 in load order you will be able to follow the process of booting arrayForth and become familiar with its components.

To find the ROM code for SPI node 705 for example, type the phrase **705 @rom list**. **@rom** takes a node's cyxx coordinates and returns the number of its ROM load block.

3.4 The Bonus Materials

Some of the "study material" supplied on this disk image may be executed to do useful things, as follow:

SELFTEST (n) this phrase will run all relevant ATS tests on the chip whose IDE COM port number is given. Note that in order to test the Host chip in this way, its no-boot jumper J26 must be installed before running **SELFTEST** or the IDE will hang.

AUTOTEST (n) this will cause the Host chip to run ATS tests on the Target chip using the same general procedure as does the factory chip tester, by running tests through the synchronous connection between each chip's node 300. In addition, it runs SERDES test between Host node 701 and Target node 10001, transferring a quarter million words of known pseudorandom values each way and verifying correct receipt with line turn-arounds.

RAMTEST (n) this phrase runs tests to detect assembly problems with external SRAM on the host chip of an EVB001/2 compatible board using the IDE COM port number given.

450 load Host polyFORTH IDE boot procedure. Install no boot jumper before loading this.

460 load Burns polyFORTH into flash.

1140 load Host eForth IDE boot procedure. Install no-boot jumper before loading this.

1190 load Burns eForth into flash, for example when updating eForth.

4. Programming the F18

If you have not read DB001, the F18A Technology Reference, please do so to familiarize yourself with the computers and their instruction sets before reading this section.

Both environments include a resident Assembler for the F18 instruction set. As usual, the F18 assembler is a "vocabulary engine" meaning that when the appropriate vocabulary is selected the F18 opcodes, directives, and user defined symbols are available. On the Win32 saneFORTH system this is vocabulary number 9 while on pF/144 it is number 5. The full Assembler environment also uses a special Interpreter to access user defined labels which are stored with the object code rather than in the local dictionary.

The assembler may be used in two distinct ways. The most common is to generate code to reside in and execute from a node's RAM or ROM. In this case code is laid down at consecutive addresses based on a location counter, and is assumed to execute at the addressing for which it was compiled, and in the node stated for things like the "warm" multiport execution address and for ROM content.

The second way is to generate an instruction word as a data item. The usual reason for doing this is to make an instruction word that will be used in port execution. In the past, it was often necessary to hand code such instructions and write them in hex; on the other hand, when it was possible to generate the instructions symbolically, the destination address checking for jumps and calls was based on memory location counter for the instruction stream in which the literal instruction was being stored, not on the port in which it would actually be executed. This could generate instructions that would not work as expected when executed in the port. This new method solves both problems. By default, no location counter is assumed and the assembler will generate jumps and calls that don't depend on the execution address of the instruction word. Directives may be used to specify an execution address (normally a port) for greater range of possible destinations at the expense of tailoring the instruction word for a particular port.

Other than the common essentials of Forth compilation technology, this environment includes the mechanism to compile code for hundreds of nodes, each of which has ROM that may be used by the application.

4.1 Object Code

The assembly process was monolithic in earlier versions of arrayForth. As of aF-3 we return to incremental assembly, with object code stored in a binary file area from which it may be accessed to load nodes or chips. This file area survives reboots so it is no longer necessary to reassemble all of the source code every time one needs any part of it.

A set of binary images, called **bins**, reside on pF/144 mass storage (serial disk or flash in pF) starting at block **Obin** (4800*3 or 14400 by default). Each bin is one BLOCK (1024 bytes long) and contains 256 bytes of RAM image (stored as 32-bit numbers), 256 of ROM, and a table of up to 50 labels for RAM and ROM code. One bin exists for each physical node, and there are additional virtual bins for code that may be used by utilities or placed in different nodes for different applications. The constant **nnc** gives the number of bins for which space is allocated. By typing **nnc** you will see on the stack the value 576 which is 144*4; the area is sized for programs using up to two chips, with an extra 288 nodes' worth of virtual bins into which object (for library code distributed by GreenArrays) may be stashed as noted later.

At present, sF's Obin is at pF's Obin+2400 on the serial disk image. This is done to facilitate comparison of results between the two platforms. When arrayForth 3 is released for public use, this will cease and both systems will use a single set of object bins.

Because the label tables are persistent in object bins, a programmer with a twisted mind can create much confusion by defining primary Assembler directives or FORTH words as labels, giving the appearance of a broken system. If this has happened to you, either reassemble ROM for this node or fix it with the tools in the **EXAMINE** utility.

4.1.1 EXAMINE - Object Code Auditing

This minor utility for working with object code is useful during development. It can examine bin label tables and bin content in binary or dis-assembled, including comparison with reference binaries. Utilities are provided for initializing a set of bins and for moving quantities of bins.

Two sets of `nnc` bins are available to this utility. The first, starting at block `Obin`, is the object code produced by this machine's Assembler. The second, starting at `CFORG`, is a reference set produced by this or another system (such as `colorForth`). Facilities are provided for comparing these and for displaying differences.

- ALL** erases all RAM and ROM object code for all `nnc` bins. Additional words `-CH0 -CH1 -XTR -XTR2` erase the four chip-size sections of the bin array.
- ?ROMS** compares ROM for chips 1 and 0 with reference and displays an array of status indicators for all nodes with with "." indicating empty node in both images, "=" indicating identical ROM content and a highlighted "?" indicating disagreement.
- ?RAMS** does the same for RAM object for pseudo-nodes 1600 thru 2317, chip 1 and chip 0.
- USING** (nn) selects a particular bin for examination.
- .RAM** dumps current RAM for selected bin, highlighting each word that differs from reference.
- .CRAM** dumps reference RAM for selected bin, highlighting each word that differs from current.
- .ROM** dumps current ROM for selected bin, highlighting each word that differs from reference.
- .CROM** dumps reference ROM for selected bin, highlighting each word that differs from current.
- ASM** (a n) displays disassembly of the given address range (covering RAM and ROM) in the selected bin and annotating instruction words with labels. Differences from the reference are highlighted and the corresponding reference disassembly is shown in a second column.
- .SYM** displays the label table for the current bin.
- COMMIT** burns current RAM for the selected bin into the corresponding reference image. **Use carefully!**

In a normal development cycle where accumulated source changes are reconciled in the process of updating the backup, object can be audited in the same way. To set the current object as reference, use `DISKING` to copy `nnc` blocks from `Obin` to `CFORG`.

4.1.2 Getting Object Bins from colorForth

When converting an existing body of `colorForth` based code to `aF-3`, it is prudent to use the `colorForth` object as a reference. The following steps accomplish this:

1. In `colorForth`, **compile** all of the code of interest.
2. Say **bnamed obj-cf** (or other name of your choice) to output file name.
3. Say **32768 4800 wback** to write all the bins to a 4800-block file of the above name. Only the first 864 blocks are relevant, but this simplifies file mapping into the `sF` system.
4. Move this file into the `pf` directory
5. On the `sF` system, in block 149 temporarily replace **OBJ-AF3** with your file's name, **FLUSH** and **RELOAD**.
6. See block 98. This is a utility for converting your incoming file, mapped as 11 UNIT, into the `aF-3` bin format as one of the sections of **OBJ-REF** which is mapped as 10 UNIT. To see a text directory of these sections, list block **10 UNIT 60** - Note that loading block 98 write enables 10 UNIT
7. Select one of the 600-block reference slots for your object and maintain the directory accordingly.

8. Load block 98 and say **n WAM** where **n** is the starting relative block (multiple of 600) for your data in the **OBJ-REF** file at 10 UNIT .
9. In block 149, restore the **OBJ-AF3** file name, **FLUSH** and **RELOAD** .
10. in aF-3, find the definition of **CFORG** and aim it at your new reference image. **FLUSH** , **RELOAD** and proceed.

4.1.3 Assembling Stock Code

While incremental assembly is generally a good thing, it does make object code persistent and as a result a lapse in attention or a simple mistake in numbering can deposit garbage in a bin with lasting effects.

Any time the integrity of the object bin array is in doubt you may re-populate it with known good content on either platform as follows:

11. Using the **DISKING** utility, verify the integrity of the aF-3 source code for tools and for all stock F18 code.
 1. Erase all bins by saying **EXAMINE LOAD -ALL** .
 2. Reassemble all GreenArrays code by saying **STOCK LOAD**
 3. Assemble any application code you have added.
 4. Verify object integrity with **EXAMINE LOAD ?RAMS ?ROMS**

4.2 Assembler Syntax and Semantics

F18 coding may be interspersed with high level polyFORTH coding; indeed, it may even be used to generate instruction words as literals within high level FORTH definitions. This requires clear and explicit bounding between F18 coding and host system compilation / interpretation regimes. In this section, words shown in red are deprecated from colorForth while those shown in green are new in aF-3.

To assemble F18 code simply interpret Forth that includes one of the code bounding words `ASM[` or `A[` when necessary. The F18 assembler is resident in pF/144 and loaded as needed in saneFORTH.

4.2.1 Assembler State Variables

The Assembler's state is represented by five VARIABLES:

- IWD** 18-bit instruction word being built.
- 'SLOT** The next unused slot [0..3] of the instruction in **IWD**. If 4, further ops are forbidden (inline).
- 'IW** holds the F18 address at which the instruction in **IWD** will be stored. If negative, bits 10 and up are set, the low order bits [0..9] are the same as they are in **'IP**. In this case the instruction is being generated inline and will not be stored into a bin directly by the Assembler.
- 'IP** holds the address of the next word to store into target memory, corresponding with **p** register during execution. If negative, bits 10 and up are set, and the value of **p** is unknown for inline assembly. When an address is specified for inline assembly, the address is stored into the low order 10 bits of both **'IP** and **'IW**, and bits 10 and up of **'IP** are zero.
- 'CL** holds the slot number of call opcode in preceding word. Negative if there isn't one. Used in tail optimization. Second volatile cell is slot of last opcode stored.
- '##** Volatile flag set zero by **#** and true after parsing a number, a named literal, or a named call. When true these things perform their normal behaviors of generating literals or calls. When false, each of them leaves a number on the Assembler's stack.

4.2.2 Assembler Directives for Code Bounding

Source code compatibility between 32-bit (Win32) and 16-bit (pF/144) host platforms when generating code for an 18-bit computer is, fortunately, simple to assure. On both platforms anything being prepared to store into memory must be double precision. All other values are single precision on both platforms.

A special Forth interpreter provides the dictionary searching mechanism to allow access to **labels** defined in the current or in other nodes' **bins** in both Win32 and pF/144 implementations. In addition, this interpreter generates F18 literal references when naked numbers are encountered.

- ASM** Selects the Assembler's vocabulary without invoking the full environment, used when necessary in tool building.
- # (_ - n)** When running in the Assembler environment, as described below, most numbers occurring in the source code are intended to generate literals for the F18. When it is necessary to provide a number as an argument to one of the Assembler directives, preceding it by **#** will leave the number on the stack rather than generating a literal. **#** also conditions named values, such as the port names, to push values on the stack, and the same is true for named calls. This usage is shown where appropriate in the following examples.

Directives are used to control the Assembler's assumptions and behavior in generating inline instructions or bins for code. This section shows appropriate usage for each type of assembly. In the examples, curly brackets { } surround optional words or phrases, and the vertical bar | should be read as "or but not both".

4.2.2.1 Inline Instructions

To build an F18 instruction and leave it on the stack, the following pattern is used:

```
A[ {# <port> =P} <opcodes> ]]
```

By default, the instruction will be generated with no foreknowledge of the value in **P** during execution. A port (or other) address may be specified using the **=P** directive. The stack value produced will be double precision. When used within F18 code the value may be laid down in the memory image being built using **,** while when used in high level Forth code the value may be laid down in host memory using **I,** . Each of these words is defined appropriately in each version of the Assembler to compensate for host cell size. The formal vocabulary is as follows:

- A[** Enters the Assembler environment, saving its state and setting up for generation of a single instruction word. Both interpretable and IMMEDIATE, usable in both FORTH and Assembler environments to generate instruction words in-line for use as literals or in building tables.
-]]** (-d) Exits that environment, returning a single instruction word (double precision) and restoring the Assembler's state as it was before **A[** was encountered. *When used in a high level FORTH definition, generates a 2LITERAL with the value of the instruction word.*
- =P** (a) Sets the assumed address in register **P** when this instruction is executed. May only be used within an inline instruction definition.
- ,** (d) lays the double instruction or other value down at the next available address in F18 memory, used when the phrase **A[...]]** is nested within F18 code (see below).
- I,** (d) lays the double instruction or other value down in host memory as a four-byte number, used when the phrase **A[...]]** occurs in high level FORTH code or data construction.

4.2.2.2 Building a bin

The bin mechanism allows ROM and RAM source code to be maintained and assembled separately. Unlike the colorForth environment, aF-3 does not reassemble everything whenever a changes is made; instead we ideally assemble the ROM code only once, and incrementally assemble application RAM code as needed. To build a bin of F18 code, the following general pattern is used:

```
ASM[ # ccyxxx NODE {# ccyxxx BIN} { {-SYM}ERS}{-ROM} }
# <addr> org                                always necessary
<coding>
FORTH ... ASM                               simple vocabulary switch
]ASM <host code> ASM[                       complete environment switch
A[ ... ]] {,|lit}                            Generate instruction for table/lit
{>ROM} >BIN ]ASM
```

ASM[Enters the full Assembler environment for F18A *without initializing any of its state*. Visible in FORTH.

]ASM Exits that environment.

*Within **ASM[]ASM** the ASM vocabulary is selected and a special interpreter runs. As noted above in 4.2.2, this special interpreter recognizes the words in this section and by default generates opcodes in instruction words, literals and so on. To put numbers on the Assembler's stack it is necessary to use **#** because if you simply write numbers down they will be assembled as literals, and if you write a label it will be called. All numbers used with Assembler directives are of the width documented herein. Numbers to be assembled as literals **must** be written in double precision form.*

- NODE** (nn) Specifies which node's position, ROM and RAM to assume, and by default set the target bin to the same. Reads that node's bin into the local working image. Required.
- BIN** (nn) Overrides the target bin. Optional. *Bin assignments are currently controlled by GreenArrays; at this time we recommend you compile directly for the intended nodes and not use `bin` yourself.*
- SYM** Wipes ROM label table, used before `ERS` in nodes with many unnecessary ROM labels to recover space for application labels.
- ERS** Wipes the object memory and label table for the RAM portion of the working image. Does not touch ROM.
- ROM** Wipes the entire object memory and label tables in the working image.
- >ROM** Secures the current label table as ROM labels in the working image.
- >BIN** Writes the working image to the current target bin on mass storage. To write identical content to multiple bins, use `BIN >BIN` phrases repeatedly before `]ASM`.
- reclaim** This deprecated word was used in most colorForth code to prevent crashes due to filling the colorForth dictionary. Dictionary management in aF-3 uses conventional polyFORTH methods, applicable to any host code or data structures you might build. The host dictionary is not used for F18 labels in this system.

Later on, when loading code into nodes with the IDE "by hand" or when specifying boot conditions for tools such as the automated IDE loader or the stream generator, object code is identified by its *bin number*. By default that is simply the node number unless you have used **BIN** to stash the code elsewhere.

4.2.2.3 Special Interpreter Considerations

There is no problem with encapsulating multiple bins in a single source block, and in fact the source code distributed with aF-3 includes examples of doing this. However, the special interpreter activated by `ASM[` requires slight changes in common sF/pF practices.

4.2.2.3.1 Use of EXIT

If `ASM[` interprets an `EXIT` the block continues being processed by the normal FORTH interpreter. To achieve the usual effect of `EXIT` you must write it twice, as in `EXIT EXIT`.

4.2.2.3.2 Bins Needing more than 1 Block

This can be done in two ways, and the code distributed with aF-3 includes examples.

For the first method, see the code for Async Boot ROM. The first block loads the second within `ASM[]ASM` and the second block must begin with `ASM[` to activate the special interpreter. Note that the second block does not require any closing bracketing other than an explicit or implicit pair of `EXIT`.

For the second, see the Master DMA Nexus (bin 110) of the Ethernet NIC. The first block begins with all of the normal boilerplate to start a capsule's source, but has no closing bracketing. The second block begins with `ASM[` to resume using the special interpreter, and ends with the normal closing bracketing.

4.2.3 Location Counter

This is an incremental Assembler. Opcodes and literals are written into an image of target node memory as they are encountered; the only retroactive action it performs is to store into the destination fields of words containing forward referencing jumps (such as `if`) or calls (`leap`) when those forward references are resolved, or to change a call opcode to a jump when a call is followed by semicolon (tail optimization). The Assembler keeps track of the current position in target F18 memory with three variables documented earlier: `'IW`, `'SLOT` and `'IP`.

As opcodes are assembled, they are added to the instruction word being built at the address in 'IW and 'SLOT is maintained until the word is full or until the next opcode will not fit into the word. At that time the word being built is padded with . (nop) opcodes if necessary, and a new instruction word is started at the address in 'IP . This leaves 'IW pointing at the new instruction word and 'IP pointing at the following location in memory. The duality of 'IW and 'SLOT is necessary to support multiple instructions per word; the duality of 'IW and 'IP is necessary to support literals.

When literals are assembled, a @p opcode is generated and then the literal value is stored at 'IP , advancing 'IP . Alternatively you may write @p yourself and lay the following words down using , (comma), such as values calculated interpretively or instruction words generated by the assembler's inline nesting feature.

Jumping, calling, and memory operations address a word, not an opcode. When encountering a colon label, or any other assembler directive defining a place that may be addressed in memory, any code under construction is padded with . (nop) opcodes if necessary to align the location counter on a word boundary. This means that in absence of explicit control transfer opcodes, execution continues across alignment boundaries including the start of a colon labeled definition, a technique we have learned to use often.

The location counter 'IP is incremented in the same way the hardware increments P : The low order seven bits increment without changing the remaining bits of P . If you are generating code in RAM you will stay in RAM, wrapping its address space at x80 in terms of the value of P and also at x40 for actually addressing the memory image in the sense that the hardware ignores bit 6 of the address. The compiler will also generate code for ROM, in which case the wrapping points are at x100 and xC0. The P9 bit (x200) may be set as you wish to specify addressable destinations which will run in Extended Arithmetic Mode.

The location counter is managed using these words:

- .. forces word alignment; if an instruction word has been started, fills the rest of the word with nops. Equivalent to 1 <s1 .
- org** (a) forces word alignment then sets the compiler's location counter to a given address at which following code will be compiled into the current node's bin. 'IW and 'IP are initially equal but will separate after the first opcode has been compiled.
- #** Instructs the Assembler that the following number (or label, named literal, named call, or use if its [reference to a label in another bin]) should leave a number on the stack instead of assembling a literal or call as appropriate. For all but named literals the number will be single precision. If you wish, for example, to org to a given address you need to write a phrase like # x20 org or # joe org . See also 4.2.2 above.
- here** (-a) forces word alignment and returns the current aligned location.
- ,** (d) forces word alignment and lays the double instruction or other value down at the next available address in F18 memory, advancing location counter.
- +cy** forces word alignment then turns P9 on in the location counter. Places in memory subsequently defined will be run in Extended Arithmetic Mode if reached by jump, call, execute or return to those places.
- cy** forces word alignment then turns P9 off in the location counter.
- <s1** (n) Ensures that the next slot to be assembled will be *less than* the given number, forcing alignment in a new word if that is not the case. Used to solve problems with forward jumps or calls.

Slot assignment is a microscopic aspect of location counter management but is important to an F18 programmer for reasons such as optimizing forward references, aligning code on word boundaries for generation of literal instruction words to send another node through a com port, and the like. The words { . .. <s1 } are your main tools for managing slot allocation.

4.2.4 Control Structure Directives

These are used like those in classical Forth. The stack effects shown in square brackets reflect the host Assembler's stack; those shown in regular parens reflect the F18 stack at execution time.

4.2.4.1 Simple Forward Transfers

Forward transfer opcodes are assembled into the next available slot and may, unless you are controlling slot allocation yourself, be placed into slots 0, 1 or 2 with 10, 8 or 3-bit destination fields. The stack "handle" shown as **sa** for the unresolved forward transfer identifies both the location and the slot of the transfer opcode as well as the adjustment for **P** due to any preceding **@p** or **!p** opcodes in the same word. When **then** resolves a forward transfer, it will abort with error message "Range!" if the transfer is unable to reach the position at which **then** occurs without a larger destination field; when this occurs you must alter the code to resolve the problem.

if [-sa] If **T** is nonzero, program flow continues; otherwise jumps to matching **then** .

-if [-sa] If **T** is negative, program flow continues; otherwise jumps to matching **then** .

zif [-sa] If **R** is zero, pops the return stack and program flow continues; otherwise decrements **R** and jumps to matching **then** .

ahead [-sa] jumps to matching **then** .

leap [-sa] assembles a call to matching **then** .

then [sa] forces word alignment and resolves a forward transfer.

4.2.4.2 Count-controlled Looping

The F18 hardware supports looping under control of a count in **R** . The number in **R** is zero-based so the number of iterations such a loop makes is one greater than the initial value in **R** and that value will be zero during the last iteration of a loop. No forward transfers are used by these words and there are no issues with slots; all directives that generate backward transfers will pad the code if needed so that an opcode with the necessary size destination field may be assembled. The directives are as follow:

for [-a] (n) pushes **n** onto the return stack, forces word alignment and saves **here** to be used as a transfer destination by the directive that ends the loop. There are times when it is useful to decompose this directive's actions so that the pushing of the loop count and the start of the loop itself may be separated by such things as initialization code or a label. In this case you may write a phrase like **>r <other things> begin** .

next [a] ends a loop with conditional transfer to the address **a** . If **R** is zero when **next** is executed, the return stack is popped and program flow continues. Otherwise **R** is decremented by one and control is transferred to **a** .

unext [a] ends a micronext loop. Since the loop occurs entirely within a single instruction word, the address is superfluous; it is present only so that the form **<n> for ... unext** may be written. The micronext opcode may be compiled into any of the four slots.

4.2.4.3 Arbitrary Control Structures

As with ANS Forth, any desired control structure may be generated based on a few simple directives and flexible semantics; see the ANS Forth standard, or more to the point see the F18 code supplied with arrayForth, for many examples of composite control structures. The following directives are provided; the same stack notation (**a** for destinations and **sa** for handles to forward references) is employed here. If necessary you may code **SWAP** to affect the host Assembler's stack. New words introduced in aF-3 are shown in green.

begin [-a] forces word alignment and saves **here** to be used as a transfer destination.

while [x - sa x] equivalent to **if SWAP** . Typically used as a conditional exit from within a loop.

-while [x - sa x] equivalent to **-if SWAP** . Typically used as a conditional exit from within a loop.

until [a] If **T** is nonzero, program flow continues; otherwise jumps to a . Typically used as a conditional exit at the end of a loop.

-until [a] If **T** is negative, program flow continues; otherwise jumps to a . Used like **until** .

again [a] unconditionally jumps to a . The old **colorForth** spelling **end** may also be used.

repeat [sa a] unconditionally jumps to a and resolves the forward jump at sa . equivalent to **again then** .

else [sa - sa] jumps to matching **then** and resolves preceding forward transfer.

***next** [sa x - x] equivalent to **SWAP next** .

4.2.5 F18 Opcodes

The preferred opcode names, as shown in the *F18A Technology Reference*, each compile an opcode into a slot:

```
ex  @p @+ @b @ !p !+ !b !
+* 2* 2/ - + and or drop dup pop over a . push b! a!
```

The **call** opcode is compiled when an F18 label is referenced.

lit [d] generates a literal of the given value by inserting a **@p** opcode and laying the value down in memory.

alit [u] generates a literal of the given *unsigned* single precision value by inserting a **@p** opcode and laying the value down in memory.

<a valid number> encountered during assembly also generates a literal (*in which case it must be written as double precision!*) except when preceded by **#** in which case it leaves single or double precision on the Assembler stack as written.

S>D [n - d] converts a signed number to double precision, sometimes necessary with other directives.

Tail optimization is performed by **;** if immediately preceded by a **call**. In this case, the **call** is converted into a jump, conserving return stack space and leading to other useful techniques. Other jumps are generated by control structure directives, described later.

Four of the original opcode names assigned in **colorForth**, have proven to be poor human factors decisions because their names conflict with standard Forth usage and therefore create "programmer traps." In **aF-3** we have deprecated the four opcodes shown in red above to eliminate these "traps" and renamed them as follows:

inv replaces **-** for a bitwise ones complement of T.

xor replaces **or** for a bitwise exclusive OR of S and T.

r> Replaces **pop** for moving a word from the return to the data stack.

>r replaces **push** for moving a word from the data to the return stack.

A block of code is provided to add the old **colorForth** names to the Assembler if you truly wish to use them, however it is not resident by default to save memory.

By default, unused slots are set to return (**;**) opcodes for object compatibility with **colorForth** (and with the code in ROM on the chip.) When the Assembler concludes that it must start a new instruction word within a code stream such that execution may continue from the preceding word into the new one, it must fill any unused slots in the preceding word with nops (**.**). You may do this yourself, for example to lay down an inline machine code literal not encapsulated by **A[** and **]** with the **..** directive that's described earlier.

4.2.6 Dictionary Labels

During assembly, a table of labels is built inside each bin image. Each label is stored as a counted string with maximum of the first seven characters actually saved for uniqueness. Each label has a 16-bit value. The following words manage this dictionary:

- equ** (n _) creates a new label whose name follows, assigning it the given value.
- :** (__) forces word alignment and defines a label at here as did red words in colorForth.
- its** (nn _ - n) normally assembles a **call** to the following label as defined in the given bin. However when preceded by **#** the *single precision* value is placed on the Assembler's stack.
- <valid label>** when encountered during assembly, a valid label defined in the working image assembles a **call** to that label. However when preceded by **#** the *single precision* value of the label is left on the stack.

Writing a valid <label> is equivalent to writing **# <label> call**.

4.2.7 Other Useful Words

Several additional resident definitions facilitate writing source code for F18s:

4.2.7.1 Named Literals

The following words, naming registers, normally assemble literals in the F18 instruction stream, however they simply leave their values (double precision) on the stack when preceded by **#**.

io right down left up data ldata

4.2.7.2 Named Literals for Cardinal Directions

When developing systems for our chips, the "floorplan" of assignments for nodes and clusters of nodes often needs to be changed, resulting in moving nodes or clusters from one place to another. When doing this, one nuisance is the need to change port names when moving a node(s) between odd and even rows or columns. In addition to being a nuisance it's an excellent way to introduce bugs. We address this at the levels of both the Assembler and of the Boot Descriptor Language.

The following words, naming registers by cardinal direction, normally assemble literals in the F18 instruction stream, however they simply leave their values (double precision) on the stack when preceded by **#**.

east south west north

4.2.7.3 Named Calls

Each of the 15 valid multiport addresses has a named word that normally assembles a **call** to that address. However when preceded by **#** they simply leave their *single precision* addresses on the stack:

---u --l- --lu -d-- -d-u -dl- -dlu
r--- r--u r-l- r-lu rd-- rd-u rdl- rdlu

await generates a call to the default multiport execute for a node based on its position in the array, also may be conditioned by **#** to return address on the stack.

4.3 Module Organization

A preliminary convention has been adopted for organizing and packaging code in a modular way, from a single node to a cluster. This convention allows use of a single block number or name as a "handle" for the module. The first two or more blocks of the module have fixed functions at fixed offsets. We recommend that you apply these principles in packaging your own code. The SRAM Control Cluster Mark 1 may be studied as an example while reading the following sections; the constant `sram` is the number of the first block of that module.

4.3.1 Load Block

The first block in a module is a single block which when loaded will cause all of the source code needed by a module to be compiled. There may be arguments to this block. When completed the necessary object code will be stored in appropriate bins (which may simply be those belonging to the nodes programmed.) If by its design a module needs to export addresses or other identifiers, these will be available in the dictionary after it has been compiled for use in code compiled later. *In a degenerate case this block may actually contain all of the module's source code.* See block `sram` for an example.

4.3.1.1 Identifier Scope Control

F18 compilation is done incrementally, which means that identifiers may not be "forward referenced" symbolically. Any identifier of any sort must be defined earlier in the compilation sequence than its references.

The scope of identifiers, including intentional or inadvertent overloading of system or compiler vocabulary, is controlled using the `remember` word `reclaim`. In the normal case, identifiers are not exported at all, and if they are it is usually not far in the compilation sequence. As a general practice, we recommend starting each node's code with `reclaim` unless you are forced to do otherwise. Elaborate scope management structures may be implemented by defining and using your own `remember` words within the scope of `reclaim`.

4.3.2 Boot Descriptors

The second block in a module (load block number plus two) defines the loading requirements for the module, using the high-level language described in section 6.1.1 below. This language lists the nodes that have to be loaded, indicates what code to load into their RAM, provides for initialization of registers and/or stacks, and provides a starting execution address for each node. The data are recorded in tables so that the order of declaration does not need to match the actual order in which the nodes are loaded. See block `sram 2 +` for an example.

4.3.3 Residual Paths

In rare cases there may be modules that need to be loaded, activated, and used during a boot sequence; SRAM and SDRAM control clusters are examples of these. Once started it may be difficult or impossible to stop them without resetting the chip and, in the case of SDRAM, perhaps losing the data in the device. If an application's boot process has this sort of complication then the module in question should provide a new path for use by stream or IDE loaders that can reach and boot all nodes remaining in the chip that are not part of the module which is running and hence now "in the way." See block `sram 4 +` for an example, showing suitable paths for boot nodes 708 (async serial) and 705 (SPI flash.)

4.3.4 Organization of Larger Projects

The appropriate structure for a larger program depends on its size and intended use. The load block should still be first. Boot descriptors should start in the second block but may require several. If IDE operation is appropriate after the program has been loaded, an IDE personalization with path(s) adjusted to access the remainder of the chip may be useful. Scripts may be necessary to specify boot stream generation, IDE loading, and/or `softsim` setup as appropriate. You might wish to include a script to produce HTML listings. Not all of these elements are appropriate for every application, but if they are small you might wish to place them in this area before the actual source code. These are merely suggestions and the recommendations may change as the system evolves.

4.4 Methods of Loading Code

After all the desired code has been compiled into object bins, it may be loaded for execution or simulation. There are presently several basic ways in which this may be done:

Manual, interactive loading into the real chip with External or Internal IDE: Insert code into RAM, do any necessary initialization manually, and call routines to observe behavior, test boundary conditions, look for side effects and so on. It is impossible to over-emphasize the importance of this practice. Simple, straightforward unit testing is a fundamental entitlement of a Forth programmer and its use dramatically simplifies debugging later on. Don't deny yourself this advantage by combining a cluster of nodes full of untested code unless you enjoy coping with serious problems either immediately or some day in the future.

Automated loading into softsim: Initialize nodes as directed by Boot Descriptors.

Automated loading into the real chip using boot streams: Use the **STREAMER** utility to construct a boot stream as directed by Boot Descriptors, then present the resulting stream to a boot node or inject it into the chip using the Snorkel.

Fully automated loading into the real chip upon RESET- signal: Use the **STREAMER** utility to construct a boot stream as directed by Boot Descriptors, then burn it into the front of SPI Flash.

Each of the automated methods uses the common Boot Descriptor Language to define what code, if any, shall be loaded and what additional initialization shall be performed before starting each node.

Which methods are appropriate depends on which host system you are running on and what connection you have to the target system, if any. sF may use one or two serial COM ports connected to one or two ("host" and "target") chips. pF/144 may use its 1.8V async interface from node 708 (requires jumpers and cables), or its 1.8V sync interface from node 300 either to the target chip on the same eval board (default jumpers), or to another chip (requires jumpers and cables). pF/144 uses node 500 to drive the reset pin of the chip and this also requires attention to jumper and cable depending on the target chip.

pF/144 may also inject streams into the chip it's running on via the Snorkel in node 207, and by building a bridge these streams may continue into another chip. For example, if you're running a basic polyFORTH system that doesn't have the Ethernet NIC nor the external clock code booted into the chip's nodes, you can add this by loading block 1583. This is a simple example of using STREAMER to build a boot stream that loads the Ethernet and crystal support, then injects that stream into the chip through the snorkel in node 207.

5. Interactive Testing

Interactivity is a critical, cardinal virtue of Forth as a programming system. Turning Forth into a batch mode programming system would be a travesty, sacrificing one of its most potent properties and setting the clock back nearly half a century. When using Forth, a programmer becomes accustomed to having her fingers literally on the fabric of the computer, directly manipulating its memory, registers, and other resources; exercising hardware and software without necessarily having to write (and debug) software simply to exercise those things.

The IDE was implemented as a natural and obvious extension of the umbilical methods used with embedded Forth systems for many decades. The method used is as non-invasive of hardware and software as is practical in this architecture, requiring no RAM or ROM in a node being tested and touching its registers as lightly as feasible.

With aF-3 there are two different environments. External IDE allows a host system to manipulate a target chip using synchronous or asynchronous serial interfaces as appropriate; this environment has been in use with arrayForth since the dawn of GreenArrays. Internal IDE is new with aF-3 and employs the Snorkel/Ganglia mechanism to permit a pF/144 system running on a G144A12 chip to interact in analogous ways with other nodes on the same chip as pF/144 is running on (*even some nodes that are part of the pF Virtual Machine!*) and, if the chip is bridged to another, on a second chip as well. There is a technical App Note AN019 on the IDE internals. This section covers the use of both environments.

5.1 External IDE

This utility is implemented for both sF and pF/144.

5.1.1 Terminology

Debugging is done on one or more *target nodes*, using an umbilical connection which will involve one or more intermediate nodes of one to three kinds (*root*, *wire*, and *end*), depending on the physical location of the target node within the chip and on the *path* taken to reach it. For IDE into the asynchronous interfaces of a GA144 including either of the chips on an EVB001 board, the root node for either chip will always be 708 because that is the one and only asynch boot node on a G144A12. The asynch IDE may thus be used to debug code in any node *except* 708.

The inner load block for the asynchronous IDE is **SER** but this is intended to be used as a factor of other utilities. To work with the host chip, say **HOST LOAD** and to work with the target chip say **TARGET LOAD** after remembering to edit **A-COM A-BPS C-COM** and **C-BPS** to define the correct COM ports and line speeds.

5.1.2 Using the External IDE

1. Assemble the necessary code into bins and make any necessary connections to the target system. Set serial port numbers as needed. **sF ONLY:** Connect at least one COM port (usually high speed FTDI) to the target system. If this is a new port, configure it correctly as described in the Appendix applicable to your platform. Edit the main configuration block, **AF-3 2 +**, to reflect this port number as host or target chip, and select a baud rate that works for the electrical interface in use. Repeat this process if you are using two connections.
2. Examine and, if necessary, change the path lists to make sure you can reach all the places you need to. See the definitions of **0PA 1PA 2PA** after the block named **SER** for examples.
3. Make sure nothing else in the windows environment (like dialog boxes) has any open handles for these device(s).
4. To work with the host chip, say **HOST LOAD** on either platform. **sF ONLY:** To work with the target chip say **TARGET LOAD**.
5. Say **TALK** to reset the target chip (may be prevented) and download root node talker code into the chip. **On some target boards, you may need to cycle power or press a reset button before this step.** If the serial port is already set up and you do not wish to reset the chip, use **Talk** instead.

6. Say **.PTH** for a concise display of IDE state. You should at this point have three defined paths, one to each of the nodes immediately adjacent to the root node, for example:

```
.PTH
 2  709    1  709
 1  608    1  608
 0  707    1  707  <---  SELECTED
```

The first column is path number; the second shows which direction it goes (first node away from the root). Third column shows the number of COM port boundaries between the root and the target node, and the fourth column is the target node to which the path is currently connected.

7. To see more about the target node use **SEE** which displays the path state as well as the stack and RAM in the target node.
8. Define, tear down, and select paths as you wish, and operate on the nodes in question using the words described below. *Note: If you wish to prove that all other bootable nodes are completely idle, begin your activities by saying **2 <root> HOOK** (for example, **2 708 HOOK** when using async IDE) and then **2 -HOOK** to tear that path down. The default path 2 can reach all 143 nodes accessible to the IDE.*
9. Operate on the selected target node as you wish using the basic or enhanced vocabulary.
10. If you wish to change the code you are downloading into any node(s), change the F18 source and load it to update the object bin(s). IDE state is not affected by normal assembly. Any paths you had wired up, and whichever you had most recently selected, will still be in effect as you can see by displaying **PANEL** again. *Assembly does not communicate with the chip under test.*
11. When you are finished you may depart the IDE by **.....**

sf ONLY As long as you do not say **bye** to your arrayForth session, you may later load the IDE again and resume communicating with whichever paths you had left connected.

5.1.3 External IDE Vocabulary

Words come in two classes: Those which wire, rip and select paths, and those that operate on the target node.

5.1.3.1 Path Routing Control

To see how the path routing is currently set up, say **.PTH**. Node numbers used with these words are always in the **cyxx** form.

PATH (i) Selects path **i** (0, 1, or 2) so that all subsequent target operations will apply to the target node for that path. *This is necessary when more than one path leads to the same target node.*

NODE (nn) Selects whichever path presently has node **n** as its target, if any. Note that it is possible to set up more than one path to different ports of the same node; if you have done this, **PATH** will permit you to select the desired port. Leaves a default path selected if none of them targets the node given.

HOOK (i nn) Wires path **i** to target node **n** after ripping out any existing wiring for that path, leaving path **i** selected. If the node in question is not accessible in the route list for that path, leaves the path set for the appropriate adjacent node as target.

-HOOK (i) Forces ripping out of any wiring for path **i**. Each node that had previously been part of this path is normally left in its default **warm** state (multiport execute); the target node *is not affected*.

Path Node Hook -Hook These words do the same things, but silently.

PTB (i-a) an array of starting addresses indexed by path number. To adopt a new path, store its address into this table at a time when that path index has been unhooked.

5.1.3.2 Target Operations

Each of these operations applies to the target node of the currently selected path.

UPD retrieves all ten words of the data stack into the local array **STACK** which is indexed (T, S, and the eight elements in the F18 stack, first element being the one that would next be popped into S). Displays the stack conditionally (using **?STK**.)

.STK displays the contents of **STACK** in two rows, in the current **BASE**. First row is the deepest four elements of the F18 stack, followed by the shallower four elements, then S and T. The shallowest element is immediately left of S; “deeper” moves to the left on that row, then to the right on the row above, with the deepest element on the right side of the top row. Thus, the eight words of the stack array seem to move circularly. This is isomorphic with the stack array and makes its behavior clearer.

```
.STK                                vv--Next Push from S
      2AAAA 2AAAA      15      15      (S)  (T)
      2AAAA 2AAAA 2AAAA 2ABAA      15      15
                                ^^--Next Pop to S
```

?STK displays the stack as does **.STK** if **?MUTE** is zero.

?RAM displays all of RAM from the selected node (in the current **BASE**).

?ROM displays all of ROM from the selected node (in the current **BASE**).

-ASM (a n) disassembles **n** words of code starting at **a** in ROM or RAM. Labels are integrated from whatever bin is current.

LIT (d) removes a double number from the host computer's stack, pushing it onto the data stack of the target node and displaying the updated stack if not muted.

R@ (a-d) reads RAM, ROM, register or port word at **a** in the target node, pushing the data read, **d** onto the host computer's stack. *In this release, uses and does not restore register **a** in the target node.*

R! (d a) writes a double number from the host computer's stack into target memory (RAM, register, port) word at address **a**. *In this release, uses and does not restore register **a** in the target node.*

RINS (d) Executes a given instruction word in the port of the target node. Usually built with **A[. . ;]]**

CALL (a) calls, from the port, the code at address **a** in the target. If that code returns to the port or re-establishes the normal rest state of the target node, interactive use may continue once the code completes. No interlocking is done so this method may also be used to start application processing which will not admit to further port execution access. Interaction with a node actively running an application may also be arranged but at the cost of periodically polling, as described below.

BOOT (a n nn) loads code into the target node, starting at address **a** in both target and bin for **n** words, from the current binary output area identified by **nn**. This need not be the same node as the target.

FOCUS forces the target to call only the port on which the current IDE path is talking to it. This remains in effect until the target is directed to execute elsewhere.

VIRGIN forces the target to call its default multiport execution address to un-do the effects of **FOCUS**.

IO DATA UP DOWN LDATA LEFT RIGHT place chip port addresses on the host's stack.

RB! @B !B RA@ RA! @A !A @+ !+ R+ R+* R2* R2/ RINV RAND ROR RDROP RDUP ROVER Execute single F18 instructions in target, using target's stack, and update stack display in panel. Note: These functions may be omitted in favor of **A[opcodes ;]]** **RINS**

5.1.4 Advanced External IDE Uses

When ripping a path out, all wire and end nodes (if any) are left in their default multiport executes. This is necessary so that, for example, paths may be crossed so long as they are not used concurrently to cause conflicts. There are occasions, such as when starting nodes that send unsolicited code or data to other nodes, when this is not appropriate. In those cases the following two functions are useful:

UNFOC Conditions -hook to work as stated above..

FOC Conditions -hook to leave every wire and end node focused back toward the root. This leaves the path completely impenetrable to anyone other than the root node, and further hooking and starting of nodes will be necessary if that path is ever to be permeable again.

The IDE is factored such that it may become a tool of other utilities. Examples are the **SELFTEST** and **AUTOTEST** routines, the IDE boot and flash burning tools used for polyFORTH, and even the **HOST** and **TARGET** load blocks. Until these conventions are fully documented here, please read the code just identified to see how it is done. In the meanwhile here are general points:

1. **HOST** begins by loading the necessary code, and then sets up the communication environment.
2. If the IDE is to be used in an automated fashion, it's usually correct to set **?MUTE** nonzero. This will prevent chatty displays of the routing every time it's changed, and will allow use of remote instructions without updating the stack display gratuitously.

5.1.5 Configuring IDE for Target Environments

As shipped, the IDE is configured for the environment of the EVB002 Evaluation Board. This will do for many other circuits using our chips, however the target environment may require attention to several things. Most of these have to do with devices connected to GA144 I/O pins or to code being loaded from flash.

- **Flash Code:** In some cases it is necessary to enable flash booting even when planning to work with the chip using the IDE. For example, some I/O initialization may be always necessary, such as switching the 18-bit parallel Address and Data buses from output to input mode, or placing other pins in states other than their default weak pull-down. The existence of such code may require adjustment of high timing in the IDE's definition of **RESET** (10 ms as delivered, to accommodate one such board we use in development.) Additionally, if a watchdog circuit is active in the target environment, this flash boot may need to load code into at least one node to keep that circuit from resetting the chip; this will require changing IDE paths from the defaults. Finally, in some cases a more extensive amount of the chip or chips is booted but we still want to keep the external IDE usable on nodes not (yet) programmed. In this case more extensive IDE path changes will be required. In either case, it may prove difficult to devise single IDE paths that can reach all nodes one might wish to program.
- **Connected Devices:** If our output GPIO pins are high at the time of reset, the amount of time it takes for these to be restored to low state after the GA144 is reset can be rather long if the capacitance on the pin is high. Our inherent pin capacitance is on the order of 2.8 pF but that is without connection to a PCB or to driven device(s). The output circuit is typically an RC first order lag, with the weak pull-down transistor looking like a 50k resistor. The time constant of this circuit is on the order of 500 ns per 10 pF. This is more important for on-board reset circuits than it is for the external IDE, but it should still be taken into consideration if extremely high capacitances are involved. Another property with similar results might be the time required by a device to restore its outputs to resting conditions upon removal of an enable or other signal due to resetting the GA144. In all such cases the **RESET** low time must be long enough to ensure that code in the chip, particularly boot code, does not awaken at end reset and see a pin high when it should be low. The default low time in the IDE's definition of **RESET** is 5 ms as delivered.

5.1.6 Working with Two Chips using External IDE

The IDE may be used to interactively debug software on both chips of the Evaluation Board as though they were a single chip with 288 nodes (two of which are invisible), using just the serial interface for the host chip.

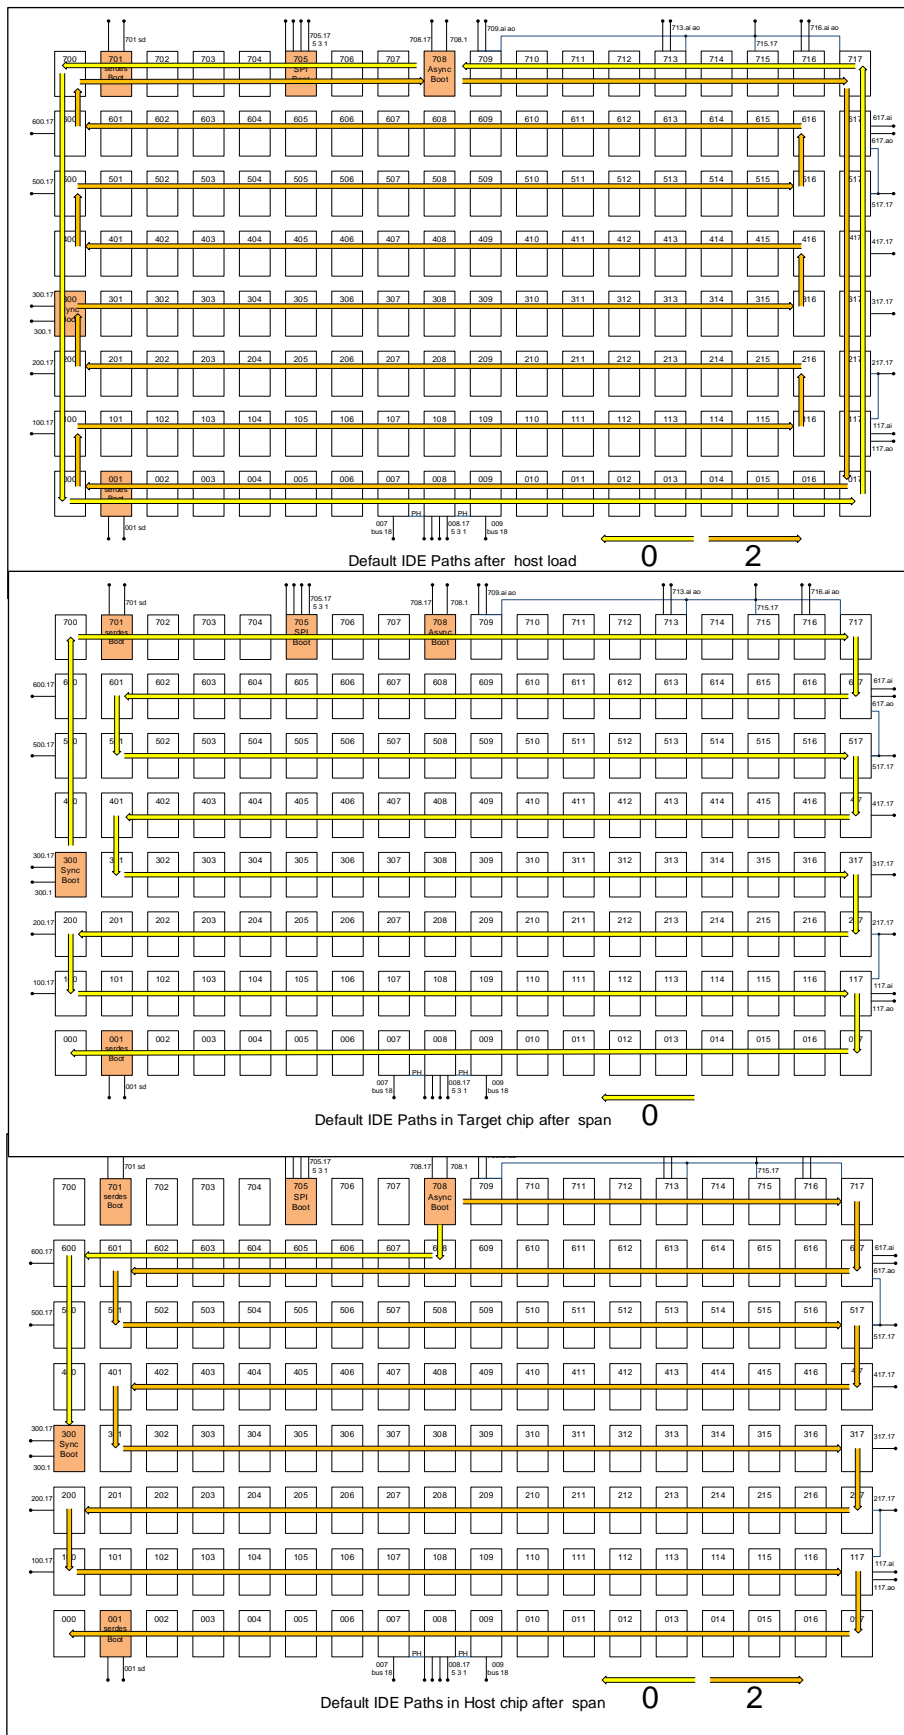
BRIDGE LOAD compiles a host IDE version capable of operating on both chips. Initially its environment and function is identical with **HOST** and you should establish connection with the host chip using **TALK** if it isn't already established.

SPAN extends the bridge IDE to encompass both chips. In order to successfully invoke **SPAN** the edge nodes of the host chip starting with 707, proceeding to the left to 700, and downward to and including 300 must be accessible for programming; it is typically used immediately after loading **BRIDGE** and establishing host connection unless something such as polyFORTH needs to be set up on the host chip first. **SPAN** resets the Target chip and programs node 300 on each chip as a transparent bridge for carrying port communications between them, using 2-wire synchronous communications. New default paths 0 and 2 are set up to cover both chips. Thereafter, nodes 300 are dedicated to this purpose until the chips are reset.

With the port bridge built, the **up** ports of node 400 on each chip are logically connected as though they were a simple COM port. Port read/write communications, such as IDE, are basically transparent across this connection except that data transfers take 100 or more times longer. The current version of the bridge supports flow control, so unlike the old version node 400 may determine whether the bridge has data for our chip or is ready to send a word to the other chip by examining **io**: If the bridge is sending us a word, node 400 can see **up** writing, and after we write to up from node 400 we will not see **up** reading again until the word was accepted by 400 in the other chip. *However, you may not infer by seeing **up** reading that node 400 on the other chip is at this time reading its **up** port. Any programming method that depends upon such awareness will have to be implemented in some other way.*

The default paths after **SPAN**, shown below, facilitate full access to the Target chip with or without the polyFORTH virtual machine present on the Host chip. Extended node numbering is supported in the form cyxx where c is zero-relative chip number; thus nodes 000 through 717 are on the Host chip, while nodes 10000 through 10717 are on the Target chip. The IDE is aware that a connection from 400 to 10400 exists through their up ports. ***You may only have one path hooked through the bridge at any time! Attempting to violate this rule will simply hang.***

5.1.7 Default External IDE Paths



After loading HOST, BRIDGE or TARGET these paths exist in the chip to which the serial connection has been made. Path 1 is available for general use at that time.

After SPAN path 0 is generally used for target chip, starting on host and passing thru the bridge to cover the entire target as shown here

... and path 2 is generally used for the host chip. Path 1 is available for general use. Path 0 by default avoids passing through the top row of nodes on the host because in polyFORTH environment node 705 is programmed for SPI flash operations.

Because paths 0 and 2 share nodes one should be unhooked before the other is hooked through those nodes. The port bridge only requires special code in nodes 300 and 10300.

5.2 *Internal IDE*

5.2.1 Terminology

5.2.2 Using the Internal IDE

Internal IDE Vocabulary

6. Preparing Boot Streams

While the external IDE may be used to load and interact with F18 code in any chip for hardware or software testing without requiring any other hardware or initialization on the part of the chip, higher level modules such as the polyFORTH Virtual Machine may require additional hardware and the loading and initialization of many nodes as well as perhaps external SRAM in a practically simultaneous manner. This is best done by generating a **boot stream** to do all of that.

Unless equipped with special ROM or booted using SERDES, our chips must be booted after reset by having a suitable **boot stream** made available to one of its **boot nodes** (such as async serial, 2-wire synchronous serial, or SPI flash nodes that are enabled for boot after reset.) A boot stream consists of one or more **boot frames**, which are data structures defined and processed by the boot nodes of our chips. Each boot frame contains zero or more words of data, a starting memory or port address at which the data are consecutively written, and a jump address to which control is transferred by the boot node after the frame has been processed. Every boot node defines a **concatenation address** to which a boot frame may jump to process another frame. After reset, a chip may boot itself from a slave device such as SPI flash memory, or it may wait to receive boot stream(s) on one or more of its enabled interfaces. For example, a daisy chain of subsidiary chips may be booted by a master chip using frames stored in the master's mass storage or read from its boot device (such as an SPI flash). As another example, the external IDE works by transmitting boot frames into the chip under test. arrayForth hosts can build and transmit streams of boot frames into a target chip for efficient booting (the old IDE based boot mechanism has been deprecated in arrayForth-3.)

6.1 The Streamer Utility

This utility is a set of tools for generating and packaging **boot streams** that consist of one or more **frames**. In a boot stream, each frame has a **header** to be interpreted by a **boot node**, and will end either by returning to a simple function in the boot node such as its **warm cold** or its appropriate **concatenation address**; in some circumstances the boot node will be given code to execute, which may read additional data from the boot stream and dispose of it in some way (for example, copying the polyFORTH nucleus into external SRAM) before completing its job and returning to such a place in the boot node, perhaps with a new **starting address for the next frame** in the case of SPI flash. Streams may also be generated for insertion into a node attached to the Snorkel, in which case no boot headers are necessary and special methods must be used if that first node is itself to be programmed.

The Streamer operates in two phases. In the first, it composes the entire stream as instructed, one 18-bit word per 32-bit cell in memory (in pF/144 we do this using unallocated space in the nearly 1 Megaword of extended memory.) In the second phase, the boot stream is converted into a form suitable for whatever medium is being used (for SPI it is a stream of 18-bit values expressed as a succession of bytes; for async serial each word is inverted and shifted into a 3-byte form equipped for auto-baud of each word. For delivery through the Snorkel, it is left unchanged.)

A minimal use of the Streamer consists of the following elements:

STREAMER LOAD	Compile the utility
nn STREAM[Specify root node and clear
{a COURSE}	Override default path for that node
FRAME[Begins a frame suitably for root node
boot description ---	see 6.1.1, Boot Descriptor Language (BDL)
]FRAME	Ends a frame suitably for root node and,
	if a forward reference was made, resolves
	it at the current aligned flash location.
]STREAM	Completes a stream and packs it as needed

Here is the full STREAMER vocabulary, followed by the Boot Descriptor Language:

STREAM[(nn) If usage from root node *nn* has been defined, selects the header formulation and final packaging suitable for that node as well as a default path whose address is returned by **ENTIRE** after **STREAM[** has been used. Clears the entire stream buffer.

The node numbers defined for this use are at the time of this writing **207**, **708**, **300** and **705**. For node 708 we use asynch serial header and 3-byte asynch word packing. For node 300 we use synchronous serial header and don't pack the stream. For node 207 we generate no headers and do no packing for Snorkel injection into the chip. And for node 705 we use SPI flash headers, compress each 8 18-bit words into 18 bytes (9 16-bit words), and set special end-frame behavior for forward references because at the end of loading each stream node 705 must be instructed to read the next stream starting at a byte boundary.

COURSE (a) Overrides the default path selected by root to be used as **ENTIRE** at any time. In a multi-frame stream it will often be necessary to change the path, for example because an earlier frame some of the chip's nodes have been programmed. It is the nature of boot streams that the first node loaded will be the last one in the path. The first node in a path *must* be the number of the node from which it originates. The address given must be the tick of an **ARRAY** holding a sequence of *nn* numbers ending with -1. Default paths for the standard root nodes are named **df708** thru **df207**.

FRAME[begins generation of a frame. All descriptor tables are initialized with their default settings. The root node is either not programmable at all (node 207) or partly programmable (memory settings and /P are the only valid operations on a boot node itself)

]FRAME Ends a frame. Calculates the frame length and generates the appropriate frame header if any, followed by all port and memory pumps to load the nodes in the current **COURSE**.

/ROOT (jmp a n bin) is used outside of **FRAME[]FRAME** encapsulation to generate a boot frame for the root node, loading that node with *n* cells from the given *bin* and jumping to the address *jmp*.

FORWARD is used immediately after **/ROOT** when generating a stream for SPI flash to indicate that the program just loaded will access flash before resuming stream processing. When this is done the address register in the flash loses synchronization with bit stream reading. In order to start up correctly again, the stream is padded to the next zero modulo 8 18-bit word index in the overall flash image, and the code loaded in the root node is patched to provide the absolute byte address for continuation in a standard manner. See definition of **FORWARD** and the source code it is used with.

]STREAM ends a stream. Converts the stream in place to the appropriate form for the intended node of origin.

.STREAM dumps the current stream, which must be in unpacked numeric form.

?STREAM audits an unpacked stream against a reference stream, see below.

?FLASH audits a packed flash stream against a reference stream, see below.

STREAM (- a n) sf ONLY returns origin and length in octets of the converted stream. Normally used within other tools.

STREAM (- da n) pf/144 ONLY returns origin and length in octets of the converted stream. Normally used within other tools.

6.1.1 Boot Descriptor Language (BDL)

Automated loading *touches* all nodes in a defined path starting at the relevant root node. Automated **SOFTSIM** loading touches all nodes in both chips. By default, a stream does as little as possible to each node touched. One item is by necessity pushed onto each stack, and register **A** is altered. Register **P** is set to the appropriate IDLE multiport address for that node, and **B** is set to the address of **IO** for all nodes not otherwise loaded. BDL is used to specify non-default treatments using statements that start with the word **+NODE** and continue with phrases describing the needed initialization of memory and/or registers in the node. Nodes may appear in any order since these statements are simply filling tables for later use. For the same reason, **+NODE** statements are cumulative in that one **100 +NODE**

phrase may set the memory loading and starting address for node 100 while another **100 +NODE** phrase interpreted later may override the starting address specified earlier. The following words constitute the BDL. Their use in context is described here and in the chapter on **SOFTSIM**; examples may be found in the distributed source code.

+NODE (nn) Selects table entries for node number nn, in cyxx notation. The values in the table entries are not changed by **+NODE** and so their values will be default unless a previous **+NODE** phrase has been interpreted for the same node.

/RAM (bin) Loads all of RAM from the given bin in cyxx notation. *By default nothing is loaded.* Partial loads may be described using the following words. Up to three RAM load descriptors may be specified for each node. Such loads are applied to the node in the order encountered in the boot description, so that if there are overlaps the last descriptor will overwrite earlier ones. When a RAM load of all 64 words is specified, any previously declared RAM loads are deleted and will not appear in the stream. Thus it is possible to preset a default background, such as of Ganglia, and override that for nodes that won't be able to serve as Ganglia, without extending stream length.

/SOME (s d n bin) Loads part of RAM, **n** words from the given **bin**, starting at address **s** in the bin, are loaded into the node's RAM at address **d**.

/PART (a n bin) Loads part of RAM, equivalent to **a a n /SOME**.

/B (d) Specifies an initial value for register **B**. *By default B is set to IO.*

/A (d) Specifies an initial value for register **A**. *By default a port call instruction is left in A.*

/IO (d) Specifies a value to be loaded into the **IO** register. *By default IO is not altered.*

/STACK (<n double values> n) Specifies up to ten values to be pushed onto the data stack, with the rightmost value on top. For example **30 20 10 3 /stack** produces the same effect as though a program had executed code **30 20 10**.

/RSTACK (<n double values> n) Specifies up to nine values to be pushed onto the return stack, with the rightmost value on top. For example **30 20 10 3 /stack** produces the same effect as though a program had executed code **30 20 10**.

/P (a) Specifies an initial value for register **P**. Default value is **xA9** which is the routine **warm** in every node's ROM. When the node is started, **warm** jumps to the appropriate multiport execute for the node's position in the array. If you wish to leave the boot routine enabled in a boot node that is touched by an automated loading procedure, specify its **P** value as **xAA** as is done by reset for such nodes

!ND (nn) Macro to load node nn with code from same bin.

+ND (a. b. p nn) Macro to load node nn with code from same bin and initialize registers.

ITS (nn _ - a) Returns the value of the label whose name follows in bin **nn**.

RIGHT LEFT UP DOWN IO LDATA RDATA Register names returning double addresses in I/O space.

EAST WEST NORTH SOUTH Register names like the above, but leading in cardinal directions for the node selected by **+NODE**.

6.1.1.1 BDL Macros

Several higher level BDL functions are also provided:

/RINF initializes return stack with all -1s. Used to support infinite <op> <op> unnext unnext loops.

+wire (s d nn) places a perfect 1-word wire in node *nn* with source and destination ports *s* and *d* respectively.

+dly (s d nn) Sets node *nn* up as a delay line from ports *s* to *d* .

6.1.2 Stream Structure

The components used for building frame content differ from those used in the External IDE. This section includes information for calculating worst case frame length. For calculations based on node quantity, we assume a bridged pair of chips containing a total of 288 nodes, two of which (300 and 10300) are inaccessible because they are the bridge, and the boot node itself which is loaded in a separate frame; so the frame will visit a new of 285 nodes in the full, 2-chip path.

A frame begins with a 3-word header unless it isn't needed. The first node receiving data via the frame body is given a focusing call that must be present in the stream; subsequent nodes get their focusing calls from the pump in the preceding node. So we begin with 1 or 4 words and add to it modules of the following kinds.

6.1.2.1 Port Pumps

A frame begins with one port pump for each node except the last in the path. The port pump is still five words but uses **A** rather than **B** to point to the port at which the pump is directed; thus the port pump clobbers A, and it also has to push one word on each stack. For *n* nodes in path after the root, *n*-1 pumps are required; for two chips *n*-1 is 284 so 1420 words of pumps add to the 4 above for 1424. Structure within the stream:

04DAF	@p dup a! @p (preceded by focusing call)
12xxx	Literal, the focusing call given to next node
lng-1	Literal, number of words following pump -1
2FAB2	>r !
05A72	begin @p ! unnext

Following the sequence of port pumps, there is a Memory Load module and a Post-Load Initialization module for each node beginning with the last node in the path and proceeding backward along the path to and including the first node after the root. When a node given a port pump is finished with pumping, the next thing in the stream will be its memory load(s) (if any) and initialization (at least of P and IO.)

6.1.2.2 Memory Loads

Each memory pump is five words long followed by the text to be loaded. In the typical case there is one memory pump and 64 words of text for a total of 69 words. Complicating this, we permit combining up to three memory loads for a node, in which case one might expect still a max of 64 words of text but with two additional pumps, bringing the total to 79 words. 285 of these would be 22515 words, for a total thus far of 22939 although a smaller number is more common because normally there is only one memory load in effect for each node. Structure:

04A12	@p a! @p
<adr>	Literal, start address of this RAM load
lng-1	Literal, number of words following -1
2E9B2	>r
05872	begin @p !+ unnext

It is possible to describe an even worse case in which there would be three fully overlapping 64-word memory loads for each node, but this is not worth considering. We do make provision for all nodes to be initially loaded as a background default by such things as the Ganglia, but to avoid letting this run up the size of the frame we check in the RAM load specifiers for full 64-words descriptors and if one is encountered all previous RAM load descriptors for that node are deleted so that they will not be present in the frame.

6.1.2.3 Post-Load Initializations

Initializing IO costs 4 words; A and B cost 2 words each and P costs 1. Initialization of each stack takes 2 words per value. We could fight for a reduction on the stack initializations but the benefit would appear to be marginal. The initialization components are structured as follow:

Set IO	04BB2	@p b!
	0015D	Literal, address of IO
	05BB2	@p !b
	value	Literal
Set A	04AB2	@p a!
	value	Literal
Set B	04BB2	@p b!
	value	Literal
Push R	048B2	@p >r
	value	Literal
Push S	049B2	@p
	value	Literal
Set P	10xxx	Jump to start address

P is (and must be) always set, so the minimum module is 1 word; for consistency with colorForth, we also always initialize B, to IO if nothing else is specified in the BDL. To initialize everything possible we have 9 words for the discrete registers and 38 words for all of the stacks, for a total of 47 words worst case and 13395 for two chips. This brings us to 36334 words so far for a practical worst case 2-chip frame.

That number pushed us across a 64k byte boundary and so we have chosen to reserve the first 128 kB of flash media for boot stream text.

6.1.2.4 Root Node Programming

The root node may simply be programmed as the last step in a complete boot stream. Earlier in the stream it may be necessary to program the root node to do other things such as for example speeding up the flash clock timing, or copying data from flash to an external SRAM as is done in booting eForth or polyFORTH.

Such intermediate programs for the root may be generated using /ROOT which does not employ BDL, or included in the BDL for a stream in which case the streamer generates two frames for the boot path (one for all nodes but the root, and one for the root itself.) When such a program will be accessing other flash during its operation and then resuming stream processing, the word FORWARD should be used after]FRAME in a flash boot stream to pad the stream to the next even byte boundary in flash, and to patch the root node program for a forward reference to the padded restart address.

Only one memory load may be specified for a root node, and it must have the same source and destination offsets if the load is specified by /SOME. Although it is possible to initialize more than RAM and P in a root node, this process would be complex and we have not built any particular such mechanism into the BDL. None of this is relevant for node 207, but it is true for any real boot node.

6.1.2.5 Special Considerations for SPI Flash

The SPI node processes standard boot frame headers which may be concatenated so long as the flash is not being accessed in any other way during booting. This stipulation is necessary because the end of a boot frame may occur on any even bit boundary within the flash, hence 0, 2, 4 or 6 bits into an addressable byte. A straight concatenation may begin on the next bit boundary to be read. A standard BDL frame will actually generate two concatenated boot frames: One to load all nodes but the root, and the next to load the root node and begin execution. The second frame is actually generated using /ROOT but is derived from the limited BDL allowed for the root node. Initialize P to the concatenation address [705 ITS spi-exec] if the root is not going to be otherwise accessing flash and there is to be another concatenated frame.

In complex flash boots there may be several logical steps, and some of them may involve programming the root node to operate on the flash. Examples include reading the polyFORTH nucleus from flash and writing it into external SRAM, and accessing a table of parameter values to be written at specified places in various nodes before booting the rest of their memory. In these cases the bit alignment of the end of the previous boot stream is lost and a new starting address must be written into the root node's memory so that after that step is complete the normal boot code may be started at a byte-aligned place following the end of the previous boot stream. This address, within the flash, is only known after the boot stream image has been padded, and `FORWARD` is used to patch the root node program with this address information.

To build an SPI boot stream that is unpacked for debugging purposes, use the phrase `0 S-END !` within `STREAM[...]STREAM`.

6.1.2.6 Special Considerations for Serial boot

When preparing a 2-chip boot stream for serial delivery through node 708, it's infeasible to visit node 000 on the first chip as part of the main stream. So, if you need to program this node in any way it must be done as part of the bridge construction. The 2-chip serial pF stream generator we provide loads a ganglion into node 000. Copy the BDL for your node 000 here if you need it.

6.2 Transmitting Boot Streams

You must generate a boot stream for the particular root node you plan to present it through. While the BDL may be the same for chips to be booted through various nodes, the text of the boot streams differs because the paths necessarily differ for each root node.

The vocabulary for injecting these streams depends on the system you're running on and on the method to be used. These methods are as follows:

6.2.1 From saneFORTH

6.2.1.1 Asynchronous Serial

sF uses its external IDE to inject asynchronous boot streams into node 708 of the HOST chip

S-ORG S-LNG HOST LOAD !STREAM resets the chip and pumps the stream just built into it through boot node 708. Returns when all words have been transmitted. There is no flow control so completion does not prove success.

!stream is used when the serial port is already set up and the chip has already been reset. Used for second and subsequent streams when multiple streams are being injected.

!NUCLEUS transmits the polyFORTH nucleus from block zero of the pF serial disk when initializing RAM over serial port.

The asynchronous boot stream can also span both host and target. Upon completion the external IDE remains loaded and, if the stream in question left IDE code in node 708 with proper P setting, the external IDE may still be used with care and due considerations for what paths are feasible after the stream has been loaded. Use `Talk` instead of `TALK` to use existing serial port and avoid resetting the chip.

6.2.2 From pF/144

6.2.2.1 Internal Streams rooted at node 207

Streams may be injected, at very high speed, directly into the host chip after being built:

!SNORK pumps the stream just built out through the Snorkel into node 208 or 307 as indicated by the path in effect at the start of the stream. Returns when all words have been absorbed by at least the first node in the path.

Here is an example that instructs node 715 to emit an approximately 75 MHz signal on pin 715.17, assuming that nothing but ganglia have been loaded into any of the nodes along its path:

```

1581
0 ( Descriptor test)
1 ASM[ # 715 NODE ERS # 0 org
2   begin begin !b unext unext  >BIN ]ASM
3
4 0 ARRAY MYP  207 ORGN 210 TO 710 TO 715 TO -1 ,
5
6 207 STREAM[   ' MYP COURSE
7   FRAME[ 715 +NODE 0 1 715 /PART
8   -1. -1. -1. -1. -1. -1. -1. -1. -1. 9 /RSTACK
9   x20000. x30000. 2OVER 2OVER 2OVER 2OVER 2OVER 2OVER
10          2OVER 2OVER 10 /STACK IO /B 0 /P
11   ]FRAME ]STREAM
12 !SNORK

```

This example instructs node 500 of the target chip to emit a similar signal on its pin 500.17, assuming that the bridge has already been installed and activated, and that there are no obstructions along the path. Note the use of ORGN to make the hop between nodes 400 and 10400:

```

1582
0 ( 2-chip test)
1 ASM[ # 10500 NODE ERS # 0 org
2   begin begin !b unext unext  >BIN ]ASM
3
4 0 ARRAY MYP  207 ORGN 407 TO 400 TO 10400 ORGN 10500 TO -1 ,
5
6 207 STREAM[   ' MYP COURSE
7   FRAME[ 10500 +NODE 0 1 10500 /PART
8   -1. -1. -1. -1. -1. -1. -1. -1. -1. 9 /RSTACK
9   x20000. x30000. 2OVER 2OVER 2OVER 2OVER 2OVER 2OVER
10          2OVER 2OVER 10 /STACK IO /B 0 /P
11   ]FRAME ]STREAM
12 !SNORK

```

For the next example, we begin with a flash boot that does NOT load anything but polyFORTH and, if desired, the bridge. This demonstrates that polyFORTH can load the ethernet cluster from within a live chip. It is a torturous procedure because we must initially do a 9 LOAD for serial clock (HOME 9 LIST and edit is the best way to go) :

```

1583
0 ( Add Ether to pF)
1 STREAMER LOAD 1 CONSTANT ?CLK 1 CONSTANT ?ETH
2
3 0 ARRAY MYP  207 ORGN 208 TO 108 TO 110 TO 10 TO 17 TO
4   617 TO 616 TO 116 TO 115 TO 415 TO 414 TO 114 TO 111 TO -1 ,
5
6 207 STREAM[   ' MYP COURSE
7   FRAME[ ( Clock/Ether) ENIC 1+ LOAD ]FRAME ]STREAM
8 !SNORK
9

3983
0 Assumes flash boot only has polyFORTH and, if desired, bridge.
1   This loads and activates the ethernet NIC Mk1 and the 10MHz

```

```

2    clock.  Because of the latter several steps are involved:
3
4    1.  Reset and hit space.
5    2.  HOME 9 LIST and edit to use serial clock.  HI
6    3.  AFORTH 1583 LOAD ... and link should become active.
7    4.  RELOAD hit space
8    5.  HOME 9 LIST and edit to use Ethernet clock.  HI
9    6.  ETHER LOAD
10
11 ETHER may in fact be loaded on top of AFORTH, or presumably
12 vice versa, but conflicts on double constants may exist.
```

6.2.2.2 Asynchronous Serial to Another Chip

6.2.2.3 Synchronous Serial to Another Chip

6.3 Burning Flash

By default, when generating a stream for node 705, both saneFORTH and pF/144 pack the stream into bytes that may be written to a flash memory starting at its absolute location zero. This section describes the procedures for writing that stream image into a local or remote flash.

6.3.1 Burning Flash from saneFORTH

6.3.2 Burning flash from pF/144

Code is provided for burning one's own flash (for example, the SPI flash that boots the Host chip on an Evaluation Board) or for burning the flash on another chip with which pF is communicating using external IDE.

6.3.2.1 Burning pF's Own Flash

Burning one's own flash is straightforward; all we need do is to copy the stream image to the front of the flash starting at absolute block zero (normally mapped at 40 DRIVE). After generating the stream, simply execute the following word to rewrite flash boot on your own system:

WRITE-FLASH copies the packed boot stream just generated to the boot area of the current system's flash, at absolute block 48000 and regardless of the current **OFFSET** .

Warning! This procedure assumes the flash boot area begins at absolute block 48000 in the **UNITS** mapping that is currently active on the pF/144 system. If you are using some other mapping you will need to attend to the definition named **)FLASH** .

6.3.2.2 Burning Flash on Another Chip

6.3.3 Erasing Flash

Flash needs to be erased before being written. Unless otherwise noted, all of our operations that write flash take care of the necessary erasure. If you need to erase flash without writing it, study the code in **BULK** which uses **FERS** with absolute starting block number and count on the flash (remember that 0 DRIVE starts at block 128 absolute.)

The smallest unit of flash that can be erased on the Evaluation Board is 4K bytes, so erasing always starts on a 4K byte boundary and the number of bytes erased will be rounded up to the nearest 4K as well.

6.4 Auditing Streams

When working with new stream structures it's useful to compare the stream just made with a reference. Mechanism for doing this is included in `STREAMER`. Standard areas involved in these operations are 4740, where a packed flash stream may be saved, and 2340, where an unpacked stream of up to 60 blocks (14760 words) may be saved.

UNPACK (n) unpacks the given number of blocks of flash stream from 4740 to 2340.

STASH saves the streamer's current stream buffer to 2340, limited to 60 blocks.

?STREAM dumps the current *unpacked* stream buffer in hex, highlighting each word that differs from the reference stream in the area at 2340. In saneFORTH the reference stream is at 2340 on the serial disk used by pF/144.

?FLASH dumps the current stream buffer in hex *packed for flash*, highlighting each word that differs from the current boot flash. In saneFORTH the reference stream is at 4740 on the serial disk used by pF/144.

.LOADING available after loading block 1619. Displays what the most recently generated unpacked stream does based on the content of the `STREAMER`'s tables.

6.4.1 Getting Streams from Flash **pF/144 ONLY**

The simplest way to obtain a stream made by colorForth is to burn the stream into flash from colorForth, boot to it, and then copy it to the comparison area. With standard mapping (SPI flash for 24000 blocks at absolute zero, followed by 24000 blocks of serial disk, followed by 128-block flash boot area), this can be done as follows:

1. 0 DRIVE DISKING LOAD
2. 48000 24000 4740 + nn BLOCKS (to place it on serial disk)
3. 48000 4740 nn BLOCKS (to place it on SPI flash disk).

6.4.2 Getting Streams from colorForth

When converting an existing body of colorForth based code to aF-3, it is prudent to use the colorForth stream for loading that code as a reference. The following steps accomplish this:

1. In colorForth, **compile** all of the code of interest.
2. Say **bnamed mystream** (or other name of your choice) to output file name.
3. Edit the stream generation block to abort with a bad word after the word `stream`
4. Generate stream, let it abort and note the address and count found there.
5. Say **32768 nnc + nnc + 4800 wback** to write a 4800-block file starting with the stream. Only the first few blocks are relevant, but this simplifies file mapping into the sF system.
6. Move this file into the `pf` directory
7. On the sF system, in block 149 temporarily replace **OBJ-AF3** with your file's name, **FLUSH** and **RELOAD**.
8. Use utility not written yet and conventions not made yet to capture and save the stream.
9. In block 150, restore the **OBJ-AF3** file name, **FLUSH** and **RELOAD**. If you are using the new project packaging mechanism, you will probably need to edit your `custom.txt` file instead.

10. Additional steps to be determined.

11. Remove the change you made to the stream generation block in colorForth.

7. Simulation Testing with SOFTSIM

SOFTSIM (the Software Simulator) is a program that simulates the actions of Green Arrays computers, *in two GA144 chips*, at a high level. The charter of this simulator is to achieve the same results in values but not in timing, either relative to wall time or between nodes, as would a real chip. There are coarse delays simulated for long operations such as memory read/write, but these are only provided to aid in perspective.. Tutorial follows.

7.1 The Display

After starting aF3, type **SOFTSIM LOAD** at the command line. The command prompt window will be enlarged (you may need to select a smaller font in order for this to succeed; we are not permitted to exceed the screen dimensions with that enlargement.) Here is an example of the display you should see. In the default layout shown below, there are three sections which help you "drill down" into the chip. In the upper right corner is a **full chip view** showing two GA144s as 8x18 arrays of computers, with chip 1 on top. On the left is an **overview** showing the major status and registers of a rectangular section of the chip's nodes (4x8). Between these sections is a detailed view of a **focus node**. The lower right-hand part of the display is a normal 25x80 FORTH terminal for control and programming.

The characters in the full chip view are # for nodes in the overview, 0 for the *focus node*, and = for all others. Background of nodes in this view are blue if suspended, yellow if active, and red if the node is presumed "dead" after executing an undefined operation. These colors also apply to node numbers in the overview and focus node sections.

The screenshot displays the SOFTSIM interface. The top section shows a 3D-like representation of two GA144 chips, each an 8x18 array of nodes. The left section provides an overview of a 4x8 node section, showing node numbers and their status (suspended, active, or dead). The middle section shows a detailed view of a focus node, displaying its registers and internal state. The bottom right section is a FORTH terminal for control and programming. The terminal shows the command 'SOFTSIM LOAD' and the resulting display. The terminal also shows the command 'SOFTSIM Help for sF/Win32 System' and the resulting help text.

```

V      V      V      V      V      V      V      V      00000  p= 25 io=04291 b=r-1 a=0015D  io  = = = = =
302a11 303a11 304a11 305a11 306a11 307a11 308a11 309a11  pins= 0010 data=  = = = = =
4fetch 4fetch 4fetch 4fetch 4fetch 4fetch 4fetch 4fetch  >.<  i=01A8A @b ! a .  2AAAA  = = = = =
1115A5 1115A5 1115A5 1115A5 1115A5 1115A5 1115A5 1115A5  loc= 24 slot=0  @b  2AAAA  = = = = =
>0a11< >0a11< >0a11< >0a11< >0a11< >0a11< >0a11< >0a11< 2AAAA  = = = = =
a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA  SUSPENDED 2AAAA  = = = = =
b io b io b io b io b io b io b io b io b io b io  1D 134A9 call A9 2AAAA  = = = = =
o200AA o200AA o200AA o200AA o200AA o200AA o200AA o200AA  1E 134A9 call A9 2AAAA  = = = = =
r2AAAA r2AAAA r2AAAA r2AAAA r2AAAA r2AAAA r2AAAA r2AAAA  1F 134A9 call A9 2AAAA  = = = = =
t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA  start 20 04B12 @p b! @p . 2AAAA  = = = = =
s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA  21 001F5 @b + + ; 2AAAA  = = = = =
^      ^      ^      ^      ^      ^      ^      ^      22 0015D @b + ex ; 2AAAA  = = = = =
V      V      V      V      V      V      V      V      23 2A9B2 a! . . . 2AAAA  = = = = =
202-1- 203-r- 204-l- 205a11 206-r- 207-u- 208a11 209a11  cmd 24 01A8A @b ! a . 2AAAA  = = = = =
0 @ 0 @ 0 @ 4fetch 1 lb 2 @b 2 @b 0 @b 0 @b 0  25 2FDB2 >r @p . . 2AAAA  = = = = =
i03B40 i03B40 i03B40 i115A5 i00000 i01D93 i115A5 i115A5  26 00007 @b and @b @p 2AAAA  = = = = =
>0 21 >0 21 >0 21 >0a11< 0-r-< 0 38 >0a11< >0a11< 27 01B03 @b lb @b dup 2AAAA  = = = = =
a00175 a001D5 a00175 a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA  28 007C2 lb 2/ 2/ . 2AAAA  = = = = =
b -r- b -l- b -u- b -u- b -u- b -u- b -u- b io b io  29 3EAB2 and a! . . 2AAAA  = = = = =
o204AA o204AA o204AA o200AA o204AA o200AA o200AA o200AA  2A 02CAA @r > a! . 00007 0.5796 ~us 414 cyc
r2AAAA r2AAAA r2AAAA r00032 r00175 r0001C r001D5 r00175  2B 00724 ! jump 24 2AAAA  = = = = =
t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA  last 2C 134A9 call A9 2AAAA  = = = = =
s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA  ^      ^      ^      ^      ^      ^      ^      ^
102-r- 103-l- 104 2E 105-l- 106-r- 107-d- 108-l- 109-r-
0 @ 0 @ 4fetch 1 lb 2 @b 2 @b 0 @b 0 @b 0 @b
i03B40 i03B40 i1B424 i05B05 i09B05 i09B02 i01A55 i01A55
0 21< 0 21< 1 2F 0 04< 0 3E< 0 1A >0 07 >0 11
a001D5 a00175 a001D5 a001D5 a2AAAA a001D5 a001D5 a00175
b -l- b -r- b io b -l- b -r- b -d- b -l- b -r-
o204AA o204AA o304AA o378AA o27CAA o200AA o240AA o250AA
r2AAAA r2AAAA r2AAAA r00032 r00175 r0001C r001D5 r00175
t2AAAA t2AAAA t20954 t011DC t00000 t20000 t2AAAA t2AAAA
s2AAAA s2AAAA s012A8 s00002 s00000 s15555 s2AAAA s2AAAA
^      ^      ^      ^      ^      ^      ^      ^
V      V      V      ^      ^      ^      ^      ^      +ph 0010 +ph
002rd1 003rd1 004rd1 005-d- 006-d- 007 40 008-r-1 009-r-
4fetch 4fetch 4fetch 1 lb 1 lb 0unext 0 @b 0 @b
1115B5 1115B5 1115B5 i05B40 i05B40 i1DB0A i01A8A i0040A
>0rd1< >0rd1< >0rd1< 0 03 0 03 0 41 >0 25< >0 27
a2AAAA a2AAAA a2AAAA a2AAAA a2AAAA a00141 a0015D a00141
b io b io b io b -d- b -d- b -l- b -r-1 b -r-
o202AA o202AA o202AA o272AA o372AA o102AA o04291 o052AA
r2AAAA r2AAAA r2AAAA r2AAAA r2AAAA r0001D r2AAAA r2AAAA
t2AAAA t2AAAA t2AAAA t2AAAA t2AAAA t3557F t2AAAA t00003
s2AAAA s2AAAA s2AAAA s2AAAA s2AAAA s00141 s2AAAA s2AAAA
^      ^      ^      ^      ^      ^      ^      ^
00E1F 04770

```

7.1.1 Meaning of the Data Displayed

For simplicity, the simulation assumes a fixed cycle time. All instructions begin at a cycle boundary, including resuming after suspension or completing the delay for a long operation. All parts of the display show conditions as they are before any node has begun executing its next cycle. This means the opcode shown is about to be executed if the node isn't suspended, and registers/stacks are as that next opcode will find them. While an instruction fetch is occurring, the opcode reads **fetch**. Various instructions involve time delays, in which case the opcode shown will be the one that is still executing during the delay (the delay indication differs between overview and focus node views.)

Please remember that this is not actually how the chip works! The purpose of SOFTSIM is to assist in visualizing and debugging code at a very high level, not to accurately simulate the timing of a chip made of asynchronous computers, each of which will due to process variation be running at a slightly different speed, and responding to stimuli that are synchronized only with the sources of those stimuli, not with any fixed clock. If your code respects the rules pertaining to multiport read and write operations, it should give the same results in SOFTSIM as it does in the real chip... but not in the same amount of time, and certainly not with synchronization of events between nodes that are not talking to each other or to the same internal or external central points. If on the other hand your code is pushing those rules, it could appear to work in SOFTSIM but fail in the real chip, or vice-versa. It is important to always keep this in mind when making use of SOFTSIM.

Likewise, I/O testbeds, discussed later, must be written and used with the understanding that the simulation time scale is intentionally and practically not reliable. Attempts to stress the chip with external inputs whose speeds approach the chip's limits will easily produce misleading results. Using the simulator to test logic is fine; explore timing limits with real chips and apply margins to account for variation in chip performance due to process, voltage, temperature and aging.

Where addresses are shown in these displays, they are formatted either as two digit hex numbers for RAM or ROM; or a three character string for I/O addresses. In the latter case **all** means **rdlu**. Other forms are self-explanatory.

The display shown below represents "cycle" 505 of chip 0 loaded with polyFORTH and the Ethernet NIC, set to use a 10 MHz crystal, but with no I/O testbeds (so external SRAM is not connected and what's being read from it is garbage). The details of the sections, and the vocabulary for arranging the display as you like, are discussed in the following sections.

7.1.2 "Time" and Cycle Count

Under the full chip views are two numbers. The first, labeled **~us**, is an extremely rough estimate of the elapsed time since simulated RESET. The second, labeled **cyc**, is the number of simulation cycles that have run since RESET. Because, as noted above, "cycles" do not exist in the real chip either within a node or between nodes, the number of cycles is an abstraction. The approximate elapsed time is accumulated at 1.4 ns per "cycle", so it too is an abstraction. That said, the time is likely within an order of magnitude of reality so it can be useful.

```
0.5796 ~us
414 cyc
```

7.1.3 The Overview Section

This section shows a high level summary of what's going on in each of 32 nodes in an 8x4 array. The array may be moved around on the chips; the nodes currently displayed are shown on the full chip view as the character **#**. As an example let's look at node 008 in the example above.

The lines, starting from 1 on top, indicate: 1) that we are driving pins 17,5,3,1 as shown; 3 high means SRAM read. 2) This is node 008, not running, and the "address bus" is **r-l**. 3) Next op is **@b** in slot 0. 4,6,8,9,10,11) Registers **I** (instruction), **a**, **io**, **r**, **t**, **s**. 5) no delay, register **P** is **x25** (op came from **x24**). 7) **b** register, interpreted symbolically as an address, is **r-l**. The reason the node isn't running is that it is suspended, waiting on a fetch from **RIGHT** or **LEFT**.

Refer to the table in the next section for the complete legend for nodes in this view.

```
0010
008r-1
0 @b
i01A8A
>0 25<
a0015D
b r-1
o04291
r2AAAA
t2AAAA
s2AAAA
```

I/O pins or ports are shown on the top and/or bottom of the overview. GPIO and other normal pins are depicted as 0 or 1 for low or high with color attribute indicating pin mode (white on blue high impedance input; white on grey input with weak pull-down; white on magenta output.)

Each node's depiction shows a list of the values of registers and opcode names representing the current state of that node. The state depicted is *before* an instruction is executed; the opcode shown will be executed next, and all the registers are shown as they will be when the opcode begins execution. When the F18 is suspended waiting for port or pin wake-up, the operation shown is the one that's suspended. For certain fetch or store operations there is a multi cycle delay, in which case the operation shown is the one whose execution is not yet complete. If an instruction word is being fetched, the current opcode will read **fetch**. Each node shows twelve lines. From top to bottom, using node 007 from the above display as an example, they are:

Row	Example	Name	Description
1	<i>C pppp</i>	Port r/w and pins.	<i>C</i> indicates reading (V) or writing (^) the port leading to the "North". <i>pppp</i> is blank unless the node has pins. +ph and -ph indicate phantom pin high or low. Pin states are shown with pin 17 on left and pin 1 on right. Each pin is shown with its voltage as high (1) or low (0). The mode of the pin is shown by color and background mapping. Output is white on magenta. High impedance input is white on blue. Input with weak pull-down is white on grey.
2	<i>yxxbus</i>	NODE/Abus	<i>yxx</i> is node number within either chip. As noted above, it will be black on yellow if active, white on blue if suspended, or white on red if locked up. The "address <i>bus</i> " value shown (symbolically if in the I/O range) is that used for the most recent (or ongoing) memory operation.
3	<i>sOpcod</i>	Slot/Opcode	The slot selector and next/current opcode to be executed in the next cycle. The opcode is fetch if we are doing an instruction fetch.
4	<i>i03B0A</i>	I	The instruction register
5	<i>WdpppE></i>	M/P	<i>W</i> and <i>E</i> indicate writing or reading the "West" and "East" ports using <> to show read/write direction similarly to use of ^ and V for North/South. <i>d</i> is the delay counter. If it's >1, the operation shown will not complete this cycle. <i>ppp</i> is the P register (program counter) value, defining the next address to be fetched/stored by @p !p or instruction fetch.
6	<i>a00115</i>	A	18 bit pointer register A.
7	<i>b -1-</i>	B	10 bit pointer register B, symbolic if known.
8	<i>o102AA</i>	IO	The write-only part of the IO register.
9	<i>r2AAAA</i>	R	Top of return stack.
10	<i>t00141</i>	T	Top of data stack. If carry latch is set, "*" replaces "t".
11	<i>s0015D</i>	S	Second on data stack.
12	<i>Cddddd</i>	Pins and/or ports	<i>C</i> indicates reading (^) or writing (V) the port leading to the "South". <i>ddddd</i> if nonblank is the value of the data register/bus/pins. When white on blue the port or device is in input mode. When white on magenta, the port or device is in output mode.

7.1.4 The Focus Node Section

Between the overview and the full chip view lies a detailed depiction, expanding on the overview, of the current *focus node*. Field depiction and color coding are the same except as indicated. Here is node 8 at the moment of the above example.

The full **cyxx** node number is shown, highlighted for activity or lock-up. Registers p, io, b and a are shown on the top line; p and b are shown with symbolic addresses when relevant. If a is in the I/O range, its full hex value is also shown symbolically.

Pins and data in the next line are the same as those in the top and bottom rows of the overview.

The next line shows the I register with its disassembly.

The address, slot and opcode shown reflect the instruction to be executed in the next cycle, or the previous instruction being continued due to suspension (for I/O) or delay (for operations like memory). Note

that in normal operation P advances after each fetch or store in memory using P, so unless the node is executing from a port, "p" and "loc" will normally differ. Beneath that line may appear **SUSPENDED**, **DELAYING** or **LOCKUP!**.

To the left of these lines is a dot. If there are any port communications going on, the symbols used for the same purpose in the overview appear on the cardinal sides of the dot. "Arrows" pointing to the dot are reads, away writes.

Beneath these lines is a dump of 16 contiguous cells of memory showing labels if any, address, value of the data and a disassembly of that value interpreted as an instruction word. If the address in P is visible, it is highlighted black on yellow.

The right side of this view shows the return stack on top, growing downward, and data stack below, growing upward. Carry is shown in the same way as in overview: Label for register t changes to * when carry set.

Operator functions are provided to select a focus node, and to toggle between the current and most recent focus nodes. For any given node, by default the memory dump tracks P to ensure the next cell of the instruction stream is visible. Operator functions exist to override this by specifying the starting address which is then frozen, and to release this freezing of the dump origin.

7.2 Loading Code to Simulate

When SOFTSIM is loaded, node RAMs are only initialized as they typically appear on power-up (x15555) and the nodes are all set to their P values as of RESET (the simulator presently uses xA9, warm, for non-boot nodes, for simplicity). If you allow this to run, all nodes will, in not many ns, be **SUSPENDED** (asleep) as is the case with the real chip.

At present, there is only one fully supported method for initializing the chip with code to run. SOFTSIM directly understands and processes Boot Descriptor Language (see 6.1.1). So, after loading SOFTSIM, load some BDL. For example, see block 1674, which loads polyFORTH virtual machine, SRAM cluster, and the Ethernet NIC. This body of code will not do much without I/O testbeds but it is a place to start. See below for another example.

```

00008  p= 25 io=04291 b=r-l a=0015D io
        pins= 0010 data=
>.<    i=01A8A @b ! a .           2AAAA
        loc= 24 slot=0 @b         2AAAA
        SUSPENDED                 2AAAA
                                   2AAAA
                                   1D 134A9 call A9      2AAAA
                                   1E 134A9 call A9      2AAAA
                                   1F 134A9 call A9      2AAAA
start   20 04B12 @p b! @p .       2AAAA
        21 001F5 @b + + ;         r2AAAA
        22 0015D @b + ex ;
        23 2A9B2 a! . . .         t2AAAA
cmd      24 01A8A @b ! a .         s2AAAA
        25 2FDB2 >r @p . .       2AAAA
        26 00007 @b and @b @p     2AAAA
        27 01B03 @b !b @b dup     2AAAA
        28 087C2 !b 2/ 2/ .       2AAAA
        29 3EAB2 and a! . .       2AAAA
        2A 02CAA @ r> a! .       00007
        2B 0B724 ! jump 24        2AAAA
last    2C 134A9 call A9         2AAAA

```

7.3 Operating the Simulator

At any time, you may type **HELP** to display a help screen with reminders of these operating procedures. The vocabulary for operating is as follows:

OVVIEW enables the IJKL keys to move the 8x4-node overview rectangle around. You will see feedback in both the full chip view and the overview area. Any key other than IJKL exits this dialog.

SEE (nn) selects a new Focus Node (cyyxx) and remembers the previous one.

OTHER toggles between the current focus node and the one most recently focused.

MEM (a) sets new starting address for memory display of current focus node and freezes it.

-MEM un-freezes memory display of current focus node so it will again track P.

Z takes a single step of all nodes and updates entire display.

FAST (n) runs *n* steps updating only the display of "time" and cycles. Full display update when done.

SUPER (n) like **FAST** but does not update the display at all until done.

SS Sets for operator-initiated single steps by striking any key except ENTER, which punches out.

BREAK (a s nn) sets breakpoint for node *nn* when next op to execute is address *a* slot *s*.

-BREAK (nn) clears any breakpoint for node *nn*.

7.3.1 About Breakpoints

Each node may have one and only one breakpoint defined. When any node is about to execute the instruction from the address and slot given, SOFTSIM will stop before starting to execute the entire cycle in which that instruction would begin (or continue) execution and display that node as the focus node.

If you specify an instruction that may suspend, it will continue to stop simulation before its execution has started and will not re-breakpoint during suspension. The same is true of operations such as memory fetch/store that can be delayed. Some thought is justified in choosing breakpoints.

7.3.2 Initialization

Block 1698 is an example of loading SOFTSIM, loading chip 0 with polyFORTH and the Ethernet NIC, using the TB-SRAM testbed that simulates a 2Mb external SRAM that's initialized with the polyFORTH nucleus. It will run up to the point of waiting for an autobaud space key in node 200.

Another example is block 1737 which loads SOFTSIM, assembles a simple program, loads it into node 000, and sets up the view for simulation.

7.4 Testbeds

Testbeds in SOFTSIM are pieces of code that simulate external circuitry well enough to exercise the F18 code. Several are provided and by default compiled with SOFTSIM. Each testbed provides a vocabulary to instantiate it for one or more interfaces. For example, **TB-SRAM** instantiates a 2MB external SRAM interfaced with the 40 pins of nodes 7, 8 and 9.

You may write your own testbeds for use with SOFTSIM. The linkage is via **nTB** for each node. This vector is @EXECUTEd after each instruction cycle has been executed on a node but before the end-cycle updating, suspension or delay processing have been done.

7.5 Interactive Testing with Softsim

You are invited to study the code of SOFTSIM and to make use of that information if you wish to probe details of a node's state that are not displayed in our views. The vocabulary for doing this is not documented presently because SOFTSIM in arrayForth 3 is still under development and thus is subject to change.

8. Practical Example

We have chosen a simple application to act as a practical example of how to develop and test an arrayForth program. The only parts we will need are those included in your EVB001 evaluation kit so you should be able to reproduce our results exactly. Our application is a simple PWM algorithm generating output to an LED.

Although simple, the PWM we will demonstrate uses a nontrivial approach. Many PWMs divide a fixed interval into a low period and a high period such that their sum is a constant period. The output value is the ratio of high to low time. The fixed maximum update rate derives from the period chosen. The resolution derives from the number time units the period is divided up into, usually a power of two.

PWMs have the benefit of generating an analog value from a digital output which is linear, assuming you can feed a perfect integrator. All PWMs force a trade-off between resolution and update rate. Usually this trade-off is fixed by the designer for any given application by choosing an inner loop timing interval and a number of intervals in the major period.

The algorithm we will demonstrate has several benefits over the classical design. The value presented for output is represented as a binary fraction between 0 and 1. The precision of the output is not affected by the inner loop update frequency which should always run as rapidly as attainable by the selected hardware. The higher the inner loop frequency the faster any given output will reach its desired average value. The maximum output rate is determined by period of the inner loop times the power of two represented by the least significant bit you have decided is important.

8.1 Selecting resources

For our example we will be using node 600 from the host chip because its output is easily accessible. We must move the jumper on J39 from 1-2 to 2-3 to expose host 600.17 output. This would affect automatic MMC access but will not interfere with access to the boot flash.


For brightness we will use 3.3v provided by one of the FTDI chips. Pin J7 is from the USB chip we will be using so it makes a good choice. We have chosen to solder one of our LEDs between J7 and J8-3 as a com input activity light. By connecting one of our clip leads to the anode of this LED we can pick up the 3.3v safely.

To minimize measurement interference we placed our scope between ground and J39-1. We have soldered stake pins to one of the ground areas near the prototyping region on our EVB001 and connected scope ground there.

We will use the default IDE hook path 0 which runs the perimeter counter clockwise. To make sure that node 705 is available right after reset be sure to keep J26 1-2 jumpered (for SPI no-boot select) whenever testing.

8.2 Wiring

See the adjacent image for a wiring example. We twisted the clip leads together loosely and connected one lead from the J7 USB LED anode to the anode of our free standing test LED. The other color clip lead runs from J39-1 to the test LED cathode. You may choose to place one of the resistors supplied between 3.3v and the LED to limit current when the LED is powered. We have chosen to leave it out in this demo as the measured voltage is just below spec and it simplifies the setup.



We can verify our wiring and check our assumptions by simple interactive use of the IDE.

1. Make sure the No-Boot jumper J26 is installed to avoid conflicts which might cause hangs.
2. Simply type **host load panel** to load the IDE.
3. Then type **talk 0 600 hook upd** to gain access to node 600 and display its stack.
4. Switch to hex input (see **Error! Reference source not found.**) and type **20000 io r!** to place the pin in strong pull down and the LED should illuminate. A glance at the scope should show the pin to be near 0.5v due to maximum current draw through the LED.

Try typing the following in succession and observe the voltage on your scope:

- **10000 io r!** (weak pull-down)
- **30000 io r!** (drive pin high)
- **0 io r!** (high impedance)

8.3 Writing the code

The sample code we will show you has been placed in block 842. This is part of an open range of blocks that you may use freely. To make sure that our work is automatically run through the F18 compiler we have also placed a load for it into block 200. (You should follow the same pattern when you begin your own projects. If your project spans many blocks it's a good practice to use the first block to load all the others and load only this first block from 200.) Blocks 200 and 842 in the arrayForth distribution contain these things to facilitate your walking through this exercise.

When picking a sequence for loading your blocks the safest one is to load each server node sometime before their respective clients. As a matter of definition a pair of nodes have a server/client relationship if the server trusts jumping or calling the shared port and the client feeds instructions to the same port. In practice this relationship does not change dynamically. By loading the servers first then names defined in them are available for making instruction words that the client can feed back to them.

Our sample is only one block and one node long. It begins with an identifying comment and then shows that it is F18 code for node 600 beginning from location zero. If this code were node-independent, you might want to leave out the **600 node** phrase and specify that from the block loading this one. The code is divided into three sections, each one in turn more aware of the others. We will describe the code in the order of its writing, as if wrapping the onion.

<pre>pwm demo for host node 600 pol checks for ide inputs and calls down when noticed. rtn is the return point from a down call and is used by upd as an re-entry point. cyc begins the actual pwm code. upd is the ide entry point for initial start or output update.</pre>	<pre>842 list pwm demo 600 node 0 org pol 000 @b 2000 dw and if ... 003 ... down b! @b push ex rtn 006 ... io b! then 008 drop cyc ie- 1FFFF and over . + -if ... 00C ... 20000 !b pol ; ... 00F then 10000 0 !b pol ; upd 012 xex- drop push drop 100 ... 014 pop pop iex- rtn ; 016</pre>
--	--

The core function is the three lines beginning with the comment "cyc". This code expects two stack items: an increment value and an error accumulator. It implements the inner loop of the PWM algorithm by calculating and sending a new value to the output pin. The hex number 20000 sets the output to strong pull down which will turn on the LED. All other output values will not cause significant current flow. A hex value of 30000 would select strong pull up and is not useful for this application. The hex value 10000 sets weak pull down and turns the LED off. The commented 0 would select tristate output which also turns off the LED but permits the pad to float higher. By toggling which of these values is commented one can rapidly compare the consequences. We will discuss how this is done in the next section.

The algorithm used is essentially an adaptation of the classical Bresenham line interpolation algorithm (or at least one understanding of it). PWMs are good at adjusting average power by controlling the duty cycle of a current source or sink. Power sinks such as LEDs or motors are examples of good candidates for PWM control. Because the switch to the energy source is either full on or off they are more efficient than typical analog control methods. The following analogy should help us to visualize this use of the algorithm.

Think of a pixel as being the smallest unit of energy that we can control; full power times the shortest time period we can cycle the algorithm. We can map the set of all positive slopes onto the set of binary fractions between 0 and 1 by thinking of the slope as the ratio of power on to power off time. A vertical slope is full power and a flat slope is zero power. An infinitely thin line leaving the origin with rational ratio will pass between many pixels without striking them dead center. There is an error term that maintains the amount by which each ideal point is missed as the line goes by. The slope, as a binary fraction, is added to this error term. Each time there is an overflow a one is output. Each time there is no overflow a zero is output. The remainder less than one left in the error term is always carried forward to the next interval.

At 0.5 duty cycle there is a perfect square wave at maximum frequency. Above 0.5 the high pulses begin to concatenate, separated by low pulses of the minimum width. Below 0.5 the low pulses concatenate between lone high pulses. At 0.25 there is a pulse train at half the maximum frequency. At 0.125 the frequency is a quarter of maximum. Above 0.5 for complementary slopes (where complementary is defined as $1-x$ and x is 0.5 to a positive power), the signal frequency changes just as it does for those x below 0.5, except that the output signal is inverted. For slopes below 0.5 represented by more than one 1 bit in their binary fractional forms, the output is made up from all contributing frequencies interspersed. Because only a single overflow is possible in each cycle, each frequency is magically merged at its own unique phase. At slopes of either 1.0 or 0.0 the error term never changes and the output either saturates or stops.

In our implementation the hex number 20000 represents 1. The number 10000 is a half and so on. At the beginning of "cyc" the **1ffff and** removes any present overflow bits from the error term so that the next new one can be detected. The phrase **over . + -if** adds the slope in S to the error term in T and tests the overflow bit. In the case of overflow **20000 !b** maximizes the output current. For the non-overflow case **10000 !b** minimizes the current flow. If our algorithm never had to represent but a single slope then at this point each of the two output phrases would simply jump back to the **cyc** point. We want our code to entertain new inputs as well as to accept debugging illumination requests so instead we jump to the command monitoring function called **pol** above.

When the code is running in diagnostic mode it will be loaded from an IDE "wire" coming from node 700. Examination of the G144 quick reference poster shows that 600 and 700 are connected by their **down** ports. When an IDE command is issued across the wire the first instruction is a focusing call that limits the target's program counter to only a single port decode (in case the target had been executing a multiport fetch before). At the completion of an IDE command, other than the call command, a return instruction completes the command and returns the target to its original task. For an idle target node this will be a multiport execute. In our case we will be executing the PWM rather than a port fetch. The three lines of code starting with **pol** serve to poll for IDE commands and if one is detected then we turn control over to the port completely.

The first line of code fetches the IO port value, masks out the down write bit and if the result is zero (no write pending) it jumps to the **drop** which restores the stack to the values expected by "cyc". If a write is pending we need to give up control to the port but need to leave a return path to the PWM in case the intentions of the IDE are temporary. We assume there is a focusing call to **down** sitting in the port and this instruction needs to be removed by reading. We also must emulate executing this call but with a return address back to ourselves. The phrase **@b push ex** pulls the call from the port, pushes it to the return stack where **ex** (execute) performs a co-routine jump to the address pushed as part of the call. The new return address replaces the call address on the return stack. In this way when the IDE completes it will return us to the **io b!** phrase on the last of these three lines, restoring B.

The final requirement that must be met for this demonstration code is to help the IDE in setting up and changing the operating conditions. The IDE provides for pushing a literal item onto the target data stack but it does not support insertion of an item to replace S in a single operation. The two lines of code called **upd** provide this function. The hex literal 100 at the end of the first line is known to occupy location 13 in ram. We will code an IDE script word in the next

block called **seed** which will store a value from the arrayForth stack at location 13 and then call **upd** to inject that value into the PWM as a new slope.

As soon as the code is written we add the phrase **842 load** into block 200 so that the next time we type **compile** this block will be included in the compilation. Once we do so we will observe that all our grey numbers which were initially 001 have all been modified by the F18 compile to reflect the program counter at that point in the program. If you mistyped any names or used a name before defining it the compiler will abort with question mark appended to the word you typed (such as **compile**) to perform the compilation. If you respond with **e** the editor will take you to the point in the block where the error was first detected. If all compiles well the next step is to generate some IDE script to help us install and test the code.

8.4 The IDE script

We believe that interactive development is not merely the responsibility of some esoteric, third-party, software development platform. We believe it is primarily a mindset. The choices and tools presented by the development platform must simply not conflict with the proper mindset. The mindset cannot be enforced by the tools. The following precepts can be considered to be part of the mindset.

- Make small changes.
- Save often.
- Test every change.
- Make sure all steps are repeatable (such as rebuilding everything before each test).
- Choose development paths that do not preclude testing for significant intervals.
- Shun tools which delay your feedback because they serve to distract you.

The reason colorForth keeps all source code memory-resident and compiles all code from pre-parsed tokens is to encourage recompiling often. The use of scripts encourages repeatability and in colorForth you cannot even make a definition without first committing it into a block. The arrayForth word **compile** supports these precepts by quickly rebuilding all F18 object code and also by reloading whatever test environment you have configured for the current stage in your development. To support the latter function, whenever you load testing tools they should begin with the phrase **0 fh orgn !** which directs **compile** to reload that test environment as part of its job.

In block 844 you will find the test script we have built for exercising the PWM demo. The template for this script was copied from the **host** block that you loaded earlier when we typed **host load** as part of assumption checking for node 600 and our wiring. We encourage you to read the shadow for the **host** block. In our present case we have deleted the sample definition and the **canon load** and added two script definitions on top the standard serial IDE functions. We could also, if we had wanted, defined here a custom wiring path to replace the standard one but the defaults will be adequate for such a simple program.

configure ide for demo testing.	844 list
no canonical opcodes	demo ide boot empty compile serial load
use the 'remote' ones	customize -canon 0 fh orgn !
seed loads pwm 'rate' and re-/runs cycle.	a-com sport ! !nam
run selects node 600 target, loads the pwm	seed n 13 r! 12 call upd ;
into it and starts it with a default value.	run talk 0 600 hook 0 64 600 boot
	upd ?ram panel 0 lit 18000 seed ;

The word **seed** uses IDE remote commands to place its argument on top of the template literal in node 600 **upd** and to restart the PWM at the **upd** entry point. The panel stack display is also updated so that, if you are viewing the panel, you will see your new argument there. It is helpful when outputting a sequence to see what you put out last. On the other hand outputting a new seed does not force panel display as this would be presumptuous and could easily be considered a distraction.

The word **run** forces a reset to the host chip and reloads node 600 via path 0. It displays the panel and starts the PWM with an initial seed value of 0.75 duty cycle. Note that you would not need to use this word after a **compile** if you had not made changes to the F18 code. If you had only added or changed some script function you would be good to go.



To compile this script code for the first time and get it hooked in you type **844 load** this time only. In the future simply use **compile** and you will also load these definitions.

So now that your PWM and test code is loaded, it is time to power up the scope, type **run** and use **seed** to observe the effects of different values upon the waveform and the LED. Note that although the waveform energy is a linear function of duty cycle, your eyes do not perceive the LED intensity as linear. In a dark room you can observe the smallest value of 1 as well as increments of 1 but as soon as it becomes significantly bright you can no longer resolve such small differences. Your eye has some kind of a logarithmic response and responds better to powers of two or less. Perhaps a Fibonacci ramp would be more pleasing.

8.5 Output Observations



The output waveforms shown in this section demonstrate the behavior of the PWM. We will use the term *bit cell* when referring to the smallest unit of output and the word *frame* to refer to the period determined by the frequency contribution of the least significant one bit in the slope value. For consistency all the waveforms shown have a frame size of 64 bit cells. The size of our bit cell is determined by the both the algorithm and the particular speed of the node and chip under test. For our case the bit cell time is approximately 59.5ns and the frame time comes out to approximately 3.81us. The output values of 1/64 and 63/64 have been chosen because they are easy to sync a scope to, because they are complementary, and because they highlight the effect and shape of a single bit cell. The last value of 33/64 demonstrates how two frequency components merge.

The first image shows a frame of mostly zero (higher voltage). You can see how the output rings and then floats up gradually.



The second image zooms in on the one bit at the perimeter of the frame. We see a strong negative going pulse that arcs back up to a stable voltage just before releasing.

The third image is of a frame of 63 one bit cells and a single zero bit cell. You can see how the frame begins with a strong low going pulse that quickly stabilizes to a firm value and ends with a short ring as it is released for one bit cell.




The fourth image is a zoom into the frame transition of the previous capture and shows the single zero bit cell being replaced by the long string of ones.

The last image is of the value 33/64. It contains two complete frames. In each frame you may count 33 one bit cells and 31 zero bit cells. Observe that this packing is accomplished by joining two one bit cells in each of two locations. There is also a change of phase in the second half of each frame.

8.6 Further Study

Please see the hybrid DAC function **-dac** described in the *G144A12 Chip Reference*. That algorithm, present in ROM of the Analog nodes, uses duty cycle variation to enhance the resolution of the 9-bit current sourcing DACs.



9. Commissioning a New G144A12 Board

This section describes the simplest procedures for "bringing up" a new PCB with one or more G144A12 chips. If your hardware differs, use these as examples to develop your own procedures.

9.1 G144A12 Standalone, No Flash or SRAM

9.1.1 sF as Host

Arrange and configure an FTDI serial interface. Use EIDE interactively to explore and debug. Use serial boot streams to program.

9.1.2 pF/144 as Host

Not supported yet.

9.2 G144A12 Embedded, Flash Only

9.2.1 sF as Host

Take the steps as for Standalone above. **External flash burning is not supported yet.** Ensure board set for no-boot.

9.2.2 pF/144 as Host

Not supported yet.

9.3 G144A12 polyFORTH Capable

9.3.1 sF as Host

One method is presently supported:

1. Take the steps as for Standalone above and verify the chip is alive with EIDE. Ensure board set for no-boot.
2. Bring up polyFORTH (no Ether or clock) using serial boot stream, 1- or 2-chip.
3. **SERIAL LOAD PLUG <space>** to autobaud
4. **20 DRIVE HI** (from serial disk; default **UNITS** table maps SPI flash at zero). If you have recompiled nucleus on that serial disk or have changed clock selection in block 9 you may need to think about what you are doing.
5. **If you are commissioning an EVB001 board with old flash**, take these additional steps:
 - a. Edit block 61 There's a line selecting new and old flash that starts with 1 for new. Change this to 0 and **FLUSH**.
 - b. **COMPILER LOAD CHIP LOAD**
 - c. **EMPTY TESTING LOAD INSTALL** (now zero has an old-flash nucleus.)
 - d. Repeat steps 2 thru 4 above so you are now running on a system that can write old flash, and proceed with step 6.
6. **AFORTH** and make flash boot stream using 1668 or 1671.
7. **WRITE-FLASH** if not present in the load block you just used.
8. **?FLASH** to verify that it was written.

9. **BULK LOAD 0 DRIVE 24000 0 9 BLOCKS** to copy at least the first 9 blocks of serial disk to the flash disk area so there is a nucleus to boot. It's better to copy all 4800 blocks if you have the large flash. **Use DISKING if you are on an EVB001 board or otherwise running old style flash.**
10. **If you are running on an EVB001 or otherwise using old style flash**, edit block 129 on both flash and serial disk. Insert **EXIT** after **DISKING LOAD**. The **BULK** utility only supports new style flash; by making this change you may use procedures that involve BULK without concern. Alternatively, edit block 10 on both images to remove the definition of **BULK** entirely.
11. Set board to flash boot and Reset board.
12. Hit space & say **HI** (repeat 4 above if you only copied 9 blocks), and you are running pF/144 booted from flash.
13. Your board is now self-maintaining using either serial or flash disk as you wish.

9.3.2 pF/144 as Host

Not supported yet.

10. Reference Material

10.1 F18A Code Library

GreenArrays has written, tested and used a considerable Library of generally useful F18A code; this varies from single utility nodes, such as the Perfect Wire, to multi-node clusters such as the Ethernet NIC. Some of the code is useful anywhere on the chip, while other code such as, for example, the SRAM cluster, depends on the I/O and geometric properties of the nodes for which it is designed. Every bit of this code is distributed as source. Where practical, such code is pre-assembled and is distributed with the system as object code in virtual bins (1600 to 3200) so that it may be used readily by tools such as the IDE or included in applications via boot descriptors. In other cases, such as the Ethernet Cluster or the polyFORTH Virtual Machine, enough nodes are involved that it makes more sense to assemble it for the target nodes when needed by an application being built.

This section documents the code, and its usage, for each such element in the Library.

10.1.1 Perfect Wire (any node)

It's often necessary to use a node as a simple wire to connect ports that are not physically adjacent. In the simplest case, when the use is permanent (until chip reset) and the data move in only one direction, there is a demonstrably perfect solution: A single instruction word program that, once fetched from memory, runs entirely in the Instruction Register **I** and is an infinite loop with minimal latency and minimal jitter (one word of every 262144 transferred takes an extra unext time.)

The instruction word may be written as `A[begin @ !b unext unext]`. This single word is provided as location zero of Virtual Bin 1802 and may be placed in an application by the following descriptor macro:

+wire (*s d nn*) Places perfect wire in RAM location zero of the given node. Initializes **a** and **b** to the source and destination addresses **s** and **d**. Initializes the return stack with nine -1s.

10.1.2 64-word Delay Line (any node)

A node may be programmed to serve as a 64-word delay line with throughput on the order of 8 memory cycles (~40 ns in theory, 52 ns in practice) per data word, and this performance remains the same no matter how many such nodes are placed in series (although latency from input to output when a word is pushed into a quiescent line does depend on the number of nodes).

Each data word is pushed into the first delay line node surrounded by two instruction words, and each data word between nodes and emitted at the far end of a delay line is also surrounded by the same two instruction words. No Virtual Bins are used because no code in RAM or ROM is involved at all. Delay line nodes are only committed to this use in the sense that they require special initialization of the data stack (with instructions) and registers **a** (as any valid RAM address), **P** (as the source port) and **B** (as the destination port). So, at any time, a delay line node may be recommissioned via port execution (by default, we set it up with the source port address in **P**, but any multiport address may be used if desired.)

To place a delay line node in an application, use one of the following descriptor macros:

+dly (*s d nn*) Initializes **P** and **B** of the given node to the source and destination addresses **s** and **d**. Sets **A** to zero. Initializes the data stack with 5 copies of this instruction pair:

`A[!b @p @]` in T and `A[!b !+ !b]` in S.

+dx (*nn n*) Places **n** delay line nodes starting with node **nn** and moving horizontally (to the East if **n** positive, to the West if **n** negative). Source and destination ports of all nodes including the first and last proceed in the direction indicated.

+dy (nn n) Places **n** delay line nodes starting with node **nn** and moving vertically (to the North if **n** positive, to the South if **n** negative). Source and destination ports of all nodes including the first and last proceed in the direction indicated.

The following code shows one way to insert a datum into the delay line:

```
: shove ( n)    <port> b!  @p !b A[ !b @p @ ]] ,
    !b @p !b ; A[ !b !+ !b ]] ,
```

At the output end, here's one way to receive a datum:

```
: suck ( - n)    @ drop @ @ drop ;
```

As you can understand by studying the above, this is an excellent example of the remarkable capabilities of the F18A computer. The "program" for a delay line node is the pair of instructions pushed into its source port and is executed in that port; these two instructions surround the incoming datum. What this program does is to send, from its own data stack, exactly the same two instructions to the next node; however the datum between those two instructions is the oldest word in this node's RAM, which is replaced by the incoming data word.

10.1.3 Synchronous Port Bridge (300 and others)

An almost-transparent port bridge in bin 1904 uses nodes 300 on two chips to simulate a COM port between the UP ports of nodes 400 on each chip. The exception about transparency is that a node 400 cannot tell by inspecting its IO register whether the other chip's node 400 is actually reading its UP port. This bridge code is used by External and Internal IDE, by boot streams, and afterward by applications in 2-chip environments. For discussion, see 5.1.6 above.

10.1.4 Ganglia Mark 1 (any node)

10.1.5 Ganglia Mark 2 (any node)

By default we like to fill all nodes with this versatile messaging fabric, which is capable of exchanging up to 262k word payloads and replies between any two nodes on one or more chips that can be reached via any path through contiguous nodes running the ganglia and connected by COM ports or equivalent, such as the synchronous port bridge or any other bridging mechanism that supports flow control. These Ganglia are documented fully in App Note AN017, *Ganglia Mark 2*.

Object code is stored in Virtual Bins 1714, 1715, 1716 and 1717, one for each of the four physical orientations of F18A nodes. To fill both chips with ganglia, use the descriptor macro **GANGLIA** as an *early step* (before any memory load descriptors are processed) in a new frame definition.

10.1.6 SRAM Mark 1 Cluster (7, 8, 9, 107, more)

Code for providing services to, and coordination among, up to three Masters for a 1 MWord external SRAM. See AN003, *SRAM Control Cluster Mark 1*, for detailed documentation of usage and of internals.

10.1.7 Snorkel Mark 1 (any SRAM client node)

Code for a programmable DMA channel that may act as one of the Masters for the SRAM cluster. See AN010, *The Snorkel Mark 1*, for complete details.

10.1.8 IDE Components (708, 300 + any nodes)

The External IDE loads these components into the nodes it uses to reach into a chip and touch any nodes to which a path of available nodes exists.

10.1.8.1 IDE Asynchronous Interface (708)

Code for node 708 to provide services for external IDE. See AN019, *Interactive Development Environment*, for discussion of internals.

10.1.8.2 IDE Synch Interface (300)

Code for node 300 to provide services for external IDE. See AN019, *Interactive Development Environment*, for discussion of internals.

10.1.8.3 IDE Special Wire (any node)

Code for the nodes, if any, between the root node and the end node of an external IDE path. See AN019, *Interactive Development Environment*, for discussion of internals.

10.1.8.4 IDE End Node (any node)

Code for the final node, if any, between the root and target nodes of an external IDE path. See AN019, *Interactive Development Environment*, for discussion of internals.

10.1.9 Boot Stream Components

Boot stream generation is done entirely by distributed utilities. The components should be changed only if absolutely necessary and then only with great care.

10.1.9.1 SPI Flash Speed-up (705)

The ROM code for SPI flash booting uses, initially, a very slow timing parameter so that it can interoperate with slow devices. The first boot frame included in a boot stream for use from flash should begin by dropping a short program into node 705 and executing it to adjust this parameter for a higher speed. This program is provided in Virtual Bin 1608 for use by GreenArrays utilities.

10.1.9.2 Synch Boot Master

This code for node 300, in bin 1901, can transmit boot frames to a GreenArrays synchronous boot node. It's used as a step in the set-up of a Port Bridge and for other purposes.

10.1.9.3 Port Bridge Set-up

This code for node 400, in bin 1907, is used ephemerally along with the Synch Boot Master in node 300 and a copy of the code for the Synchronous Port Bridge in node 500. Once all three of these nodes have been loaded it initializes node 300's I/O, resets the other chip via node 500 on pin 500.17, loads the port bridge code into node 300 of the other chip and starts it up, then loads the same code into our chip's node 300 and starts it as well. After this is completed nodes 400 and 500 may be re-loaded with application code, filled with Ganglia, or whatever else you desire.

10.1.9.4 SRAM Initialization (specific nodes)**10.1.10 polyFORTH Cluster (specific nodes)**

This body of F18 code implements a virtual machine that runs a pseudo-instruction set from external SRAM. It includes a 2-node basic VM with four nodes of extension coprocessors, serial I/O through nodes 100 and 200 with associated wiring, and SPI flash mass storage. Dependencies include the SRAM Control Cluster Mk1, Snorkel Mk1 and Ganglia Mk2. Support is provided for the Ethernet NIC and for the various 10MHz frequency references it supports. For details about this body of code, please see DB006, *G144A12 polyFORTH Supplement*.

10.1.11 Ethernet NIC Cluster (specific nodes)

This body of F18 code, documented in AN007, implements a 10baseT Full Duplex autonomous NIC that operates on data structures in external SRAM as a virtual DMA device. The object code is temporarily stored in bins corresponding with the nodes in which the code resides when loaded. In due time it will be moved to library bins.

10.1.12 ATS Mark 1 (Serial IDE based)

The Automated Testing System is documented in DB007. The bulk of its F18 test code, in the form of "creeper" packages, is stored in bins 2000..2208, although a couple of components are in bin "row" 1900.

10.1.13 ATS Mark 2 (On-chip based)

10.1.14 eForth (specific nodes)

10.1.15 Documented Examples

10.1.15.1 Practical Example used above

10.1.16 Additional Test Code

10.2 Bin Assignments for Rev 03b4

"Bins" are receptacles for object code and are numbered like nodes in the chips. Bins 000 through 717 contain code that is by default destined for the nodes on the host chip; future releases will vacate this range completely. Bins 800 through 1517 map onto the target chip. Bins 1600 through 3117 are reserved for utility and tool code, and are managed by GreenArrays. Their purpose is to allow this code to be compiled and available for use at any time without interfering with application code. This table documents the assignments in effect as of the release identified above; yellow highlight indicates conversion to and validation of af-3 source.

1600 IDE async root node	1900 ats/ide analog	2200 ATS + (Mark)
1601 IDE sync root node	1901 ats/ide sync boot master	2201 ATS Ret&Data Stacks v2 (Mark)
1602 IDE wire node	1902 ats/ide test frame	2202 ATS Data stack v2 (Mark)
1603 IDE end node	1903 ats/ide test frame	2203 ATS + v2 (Mark)
1604 ide all nodes template	1904 Sync Bridge (with flow)	2204 ATS and (Mark)
1605 Snorkel Mk1	1905 ats/ide uut bridge debug	2205 ATS and v2 (Mark)
1606 SST25WFxxx Flash node 705 pF	1906 ats/ide tester bridge debug	2206 ATS xor (Mark)
1607 New flash helper 706	1907 framer bridge builder	2207 ATS xor v2 (Mark)
1608 spi speedup function	1908 Sync Bridge (no flow)	2208 ATS inv (Mark)
1609 18-bit flash r/w for 705	1909	2209
1610 18-bit flash helper for 706	1910	2210
1611 8-bit flash r/w	1911	2211
1612 Flash ops for new devices	1912	2212
1613 flash erase	1913	2213
1614 SRAM Node 107 interface (Mk1)	1914	2214
1615 - Node 007 SRAM Data Mk1	1915	2215
1616 - Node 008 SRAM Ctl's Mk1	1916 ATS empty creeper	2216
1617 - Node 009 SRAM Adr Mk1	1917	2217
1700 polyFORTH Stack (106)	2000 ATS Port test	2300 <reserved creeper tests 4>
1701 - Stack down (006)	2001 ATS Port test with 2*	2301
1702 - Stack up (206)	2002 ATS RAM extensive test (Jeff)	2302
1703 Bitsy (105)	2003 ATS T&S register test	2303
1704 - Bitsy down (005)	2004 ATS T&R register test	2304
1705 - Bitsy up (205)	2005 ATS Data stack registers	2305
1706 Serial tx (100)	2006 ATS Return stack registers	2306
1707 - Rx (200)	2007 ATS LoVD RAM test (Mark)	2307
1708 - Interface (104)	2008 ATS Stacks 1 (Mark)	2308
1709 - Wire (102 and others)	2009 ATS Multiply Step & shift	2309
1710 Flash to sram for 705	2010 ATS T,S,A&R Data Path Tests	2310
1711 - Wire for 605 etc	2011 ATS GPIO pin test (no EVB)	2311
1712 - <unused>	2012 ATS Data stack (Mark)	2312
1713 - Temp SRAM code for 108	2013 ATS I/O reg latches	2313
1714 Ganglion nodes eee (Mk2)	2014 ATS Carry set/clear	2314
1715 - nodes eoo	2015 ATS @p test (Mark)	2315
1716 - nodes oee	2016 ATS R,I to Address data paths	2316
1717 - nodes ooo	2017 ATS ROM Checksums	2317
1800 <reserved for eForth/pF>	2100 ATS Parallel Port Pins	
1801 Ethernet DMA	2101 ATS @b !b (Mark)	
1802 Small tools (wire, etc.)	2102 ATS T-B (Mark)	
1803	2103 ATS Port-I (Mark)	
1804	2104 ATS Port-I Shorts 1 (Mark)	
1805	2105 ATS Port-I Shorts 2 (Mark)	
1806	2106 ATS Port-I Shorts 3 (Mark)	
1807	2107 ATS Port-I Shorts 4 (Mark)	
1808	2108 ATS Port-I Shorts 5 (Mark)	
1809	2109 ATS Port-I Shorts 6 (Mark)	
1810	2110 ATS P-R-P Call (Mark)	
1811	2111 ATS P-R-P Execute (Mark)	
1812 Flash to SRAM 8-bit	2112 ATS SERDES 001 Slave	
1813 Node 600 code (temporary)	2113 ATS SERDES 701 Master	
1814 Flash to sram for 705	2114 ATS Analog Basic checks	
1815 - Wire for 605 etc	2115 ATS 2* (Mark)	
1816 - SRAM interface for 208	2116 ATS 2/ (Mark)	
1817 - Temp SRAM code for 108	2117 ATS 2*.b (Mark)	

2400	2700	3000	GPS Correlator Mk1 Carr 1st
2401	2701	3001	" " Carrier 2nd
2402	2702	3002	" " Carrier Turnaround-1
2403	2703	3003	" " Carrier Turnaround 2
2404	2704	3004	" " Code
2405	2705	3005	" " Carrier special first
2406	2706	3006	" " Carrier special last
2407	2707	3007	GPS Strm delay fill
2408	2708	3008	GPS Strm delay empty
2409	2709	3009	GPS Osc Monitor (600)
2410	2710	3010	GPS Osc Agent (700)
2411	2711	3011	
2412	2712	3012	GPS TEMP input mux strm/data
2413	2713	3013	GPS TEMP Smart answer buffer
2414	2714	3014	GPS TEMP snork/gang sig feed
2415	2715	3015	GPS TEMP Sig gen for meas
2416	2716	3016	GPS TEMP Haleakala sF inject
2417	2717	3017	GPS TEMP absorber
2500	2800	3100	SRAM test helper node 108
2501	2801	3101	
2502	2802	3102	
2503	2803	3103	
2504	2804	3104	
2505	2805	3105	
2506	2806	3106	
2507	2807	3107	
2508	2808	3108	
2509	2809	3109	
2510	2810	3110	
2511	2811	3111	
2512	2812	3112	
2513	2813	3113	
2514	2814	3114	
2515	2815	3115	
2516	2816	3116	
2517	2817	3117	
2600	2900		
2601	2901		
2602	2902		
2603	2903		
2604	2904		
2605	2905		
2606	2906		
2607	2907		
2608	2908		
2609	2909		
2610	2910		
2611	2911		
2612	2912		
2613	2913		
2614	2914		
2615	2915		
2616	2916		
2617	2917		

000	TEMP	300	TEMP	600	TEMP
001	TEMP	301	TEMP	601	TEMP
002	TEMP	302		602	TEMP
003		303		603	TEMP
004		304	TEMP	604	TEMP
005	TEMP	305		605	TEMP
006	TEMP	306		606	
007	TEMP	307		607	
008	TEMP	308		608	
009	TEMP	309		609	
010	ETH Rx Control	310		610	
011	ETH Rx Byte swap	311		611	
012	ETH Rx Pack	312		612	
013	ETH Rx CRC	313		613	
014	ETH Rx Frame	314	ETH Tx Framing	614	
015	ETH Rx Parse	315	ETH Tx Mux	615	
016	ETH Rx Timing	316	ETH Wire **	616	ETH/CLK Osc Agent
017	ETH Rx Jitter buffer & Clean	317	ETH Tx Pin	617	ETH/CLK Osc Monitor
100	TEMP	400	TEMP	700	TEMP
101	TEMP	401	TEMP	701	TEMP
102	TEMP	402		702	
103	TEMP	403		703	
104	TEMP	404	TEMP	704	TEMP
105	TEMP	405		705	TEMP
106	TEMP	406		706	TEMP
107	TEMP	407		707	
108	TEMP	408		708	TEMP
109	ETH Slave of DMA Nexus	409		709	
110	ETH Master of DMA nexus	410		710	
111	ETH Tx Control	411		711	
112	ETH Wire **	412		712	
113	ETH Tx Unpack	413		713	
114	ETH Tx CRC	414		714	
115	ETH Link States	415	TEMP XTAL Start control	715	
116	ETH Wire **	416	TEMP XTAL Sequencer	716	
117	ETH Rx Pin	417	ETH Passive TX Oscillator	717	TEMP
200	TEMP	500	TEMP		
201	TEMP	501	TEMP		
202		502			
203		503			
204	TEMP	504	TEMP		
205	TEMP	505			
206	TEMP	506			
207	TEMP	507			
208	TEMP	508			
209		509			
210		510			
211		511			
212		512			
213		513			
214	ETH Tx Delay	514			
215	ETH Negotiation	515			
216		516			
217	ETH Rx Pull-down	517	TEMP Xtal Osc		

10.3 Examples to hang onto

A simple stream:

```

1581
0 ( Descriptor test)
1 ASM[ # 715 NODE ERS # 0 org
2   begin begin !b unext unext >BIN ]ASM
3 0 ARRAY MYP 207 ORGN 210 TO 710 TO 715 TO -1 ,
4 207 STREAM[ ' MYP COURSE
5   FRAME[ 715 +NODE 0 1 715 /PART
6   -1. -1. -1. -1. -1. -1. -1. -1. -1. 9 /RSTACK
7   x20001. x30002. x20003. x30004. x20005. x30006.
8   x20007. x30008. x20009. x30010. 10 /STACK IO /B 0 /P
9   714 +NODE 111. 222. 333. 444. 4 /STACK
10  ]FRAME ]STREAM
11 !SNORK

```

```

.STREAM 512
0 10175 4DAF 121D5 78 2FAB2 5A72 4DAF 12175 208,9
8 72 2FAB2 5A72 4DAF 12115 6C 2FAB2 5A72 210
10 4DAF 12145 66 2FAB2 5A72 4DAF 12115 60 310,410
18 2FAB2 5A72 4DAF 12145 5A 2FAB2 5A72 4DAF 510,610
20 12115 54 2FAB2 5A72 4DAF 121D5 4E 2FAB2 710
28 5A72 4DAF 12175 48 2FAB2 5A72 4DAF 121D5 711,12
30 42 2FAB2 5A72 4DAF 12175 3C 2FAB2 5A72 713
38 4DAF 121D5 2E 2FAB2 5A72 4A12 0 0 714,15
40 2E9B2 5872 9175 4BB2 15D 48B2 3FFFF 48B2 Regs
48 3FFFF 48B2 3FFFF 48B2 3FFFF 48B2 3FFFF 48B2
50 3FFFF 48B2 3FFFF 48B2 3FFFF 48B2 3FFFF 49B2
58 20001 49B2 30002 49B2 20003 49B2 30004 49B2
60 20005 49B2 30006 49B2 20007 49B2 30008 49B2
68 20009 49B2 30010 101B5 49B2 6F 49B2 DE 714
70 49B2 14D 49B2 1BC 101B5 101B5 101B5 101B5 713-11
78 101B5 101A5 101A5 101A5 101A5 101A5 101A5 101A5 710-208

```

11. Maintenance

This section covers procedures for maintaining the sF and pF/144 systems.

11.1 *saneFORTH* Maintenance

11.1.1 Bootstrapping

sF is normally started on a Win32 platform via one of three Windows executables, named in the following forms:

- **sF<vers>-aF3.exe** the normal production nucleus, with 16 MB of system memory. This executable must lie in a path whose name contains no embedded spaces because it parses its command line with simple code. Upon boot, this autoloading nucleus loads block 534 which loads block 9, defines code for accessing the command line, parses it, and INTERPRETs any FORTH code following the first space (which should be the one following the name of the executable). This program is normally started using a supplied shortcut by the name arrayForth 3 G144A12 whose command line has AFORTH appended.
- **sF<vers>-glow.exe** the normal production nucleus, but with 1.6 GB of system memory to support the GLOW design tools (which are redacted from publicly released systems.) Normally started using a supplied shortcut by the name arrayForth 3 with GLOW. As distributed the shortcut has no interpretive code.
- **sF<vers>-panic.exe** is a standard saneFORTH nucleus which does not autoloading; you will need to type HI to load block 9. It has its own shortcut as well, named arrayFORTH 3 PANIC. It is provided to give a clean saneFORTH environment that may be used, for example, to fix mistakes that have been made in either the 9 LOAD or in the AFORTH code.

11.1.2 Nucleus Maintenance

Because this is a hosted system, the nucleus that is run for saneFORTH is necessarily packaged into the Windows executable. The sF nucleus source, as well as the Target Compiler and other utilities employed to test nuclei and to generate Windows executables, are proprietary to FORTH, Inc. and/or ATHENA Programming, Inc. and GreenArrays has not been licensed to distribute them. Therefore these sections are redacted and users of arrayForth 3 are unable to maintain the underlying x86 nucleus. Special versions, if necessary, must be requested from GreenArrays.

11.2 *polyFORTH/144* Maintenance

pF/144 is fully equipped to maintain and alter both its Virtual Machine and its polyFORTH nucleus. These procedures are fully documented in *DB006, G144A12 polyFORTH Supplement*. The shadows for pF/144 nucleus source, in the index page at 60, constitute its primary documentation.

12. Appendix: Microsoft Windows® Platform

12.1 Windows arrayForth Requirements

The host system requirements are as follow:

- Microsoft Windows with Win32 environment. We have tested this system on Windows 2000 Professional, Windows XP Professional, Windows Vista and Windows 7, 8 and 10. See below for Windows 8/10 issues.
- Sufficient physical memory and swap file to run at least one instance of a program that is capable of requiring commitment of 1.6 GB of virtual memory (if you run the GLOW version). Actual requirements depend on how much memory you actually access, which should be considerably less than this.
- Display capable of at least 1024x768x24 bit graphics. More may be required by some components.
- Standard PC keyboard.
- At least one RS232 interface with COM port drivers if you wish to use the Interactive Development Environment to communicate with GreenArrays chips. For the EVB001 and 2 Evaluation Boards, this is done with direct USB cables to the FTDI chips on the board, allowing much faster communications than are reliable with RS232 electrical interfaces.

12.2 Installation

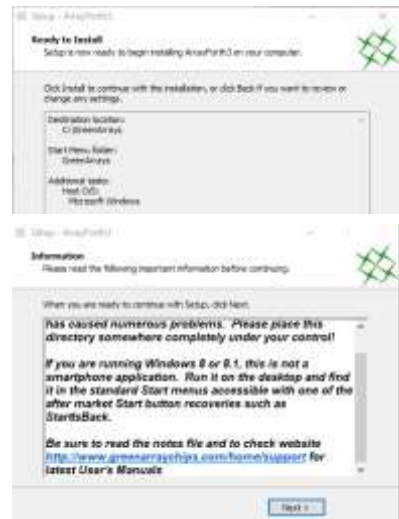
Starting with release 01b, arrayForth is distributed as a conventional, executable Windows installer. Use it as follows:

1. Download and run the installer, or run it directly from your browser if the browser supports that operation. You will see a greeting dialog that looks like this. Click the "next" button.
2. You will now be asked to read and accept our Standard Terms and Conditions for Delivery of Free Software. If you click the "I accept the agreement" item, the "next" button will be enabled and you may proceed.
3. The installer asks you to select a location at which arrayForth should be installed. By default this is a directory c:\GreenArrays. If you don't want it to be there, please change the destination in this dialog; the shortcuts made by the installer will be configured for this directory and it will be simpler for you to place it where you like initially rather than to move it later. You may install multiple instances as you wish. Hit next when done. Don't use Program Files because Windows handles updated files bizarrely.
4. Next, the installer asks where it should place the program's shortcuts. By default there will be a new Start menu group called



GreenArrays in which shortcuts will be generated for this instance of arrayForth 3 and for a file explorer view of the entire GreenArrays directory structure, facilitating your finding of your files. You will need to get at them later. Specify where you want these shortcuts and hit next.

5. The next dialog is important. It begins with a check box which, if checked, will instruct the installer to save backups of your existing source files (.cf for arrayForth and .blk for polyFORTH) in a \backup folder. This will be a subdirectory inside the arrayForth directory structure. It is advisable that you check this box if you are installing an update on top of an existing installation; you may then, later on, use a copy of your previous working .cf file as the backup source file in order to merge your work with the updated system. The second area asks you to select one of a set of buttons to indicate which supported platform you are using. This will affect mainly what shortcuts the Installer will generate. For Windows and Mac Parallels platforms, the shortcuts will be in the Start menu and will assume a Windows compatible environment. For Wine environments, the shortcuts will be on the desktop and will use different scripting files to start the arrayForth program. Select the appropriate items and hit next.
6. The next dialog is the familiar "Ready to Install" summary of what the Installer proposes to do based on your input in the preceding dialogs. If there are unpleasant surprises, use the back button and correct your input as needed. When you are satisfied with what you see here, hit Install. You will then see a progress meter dialog that should complete very quickly.
7. You will then see a dialog which purports to convey Important Information... and indeed it might. Should any given release require that you take any special actions or be aware of any changes in procedure or usage that might be at odds with the current editions of manuals like this one, this may be your only chance to learn of such before encountering it. Please take the time to read what is written here so you will not later regret having failed to do so!
8. Finally you will see a dialog about Completing the Setup Wizard, which will by default offer you the opportunity to read our standard text file that documents changes in this version. Again this is highly recommended reading; if you must forego it at this time, the file is present in the arrayForth directory when you have the time.



This concludes the initial phase of installing arrayForth. You should be able to run arrayForth using the installed shortcut and will only need to take further steps if you intend to communicate with real GA144 chips, such as those on the Evaluation Board.

12.2.1 Identifying and Configuring COM Ports

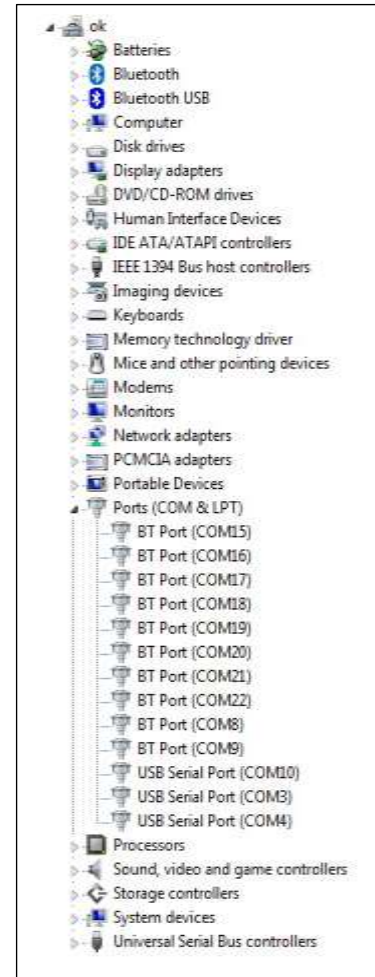
When arrayForth 3 runs, it sets basic COM port parameters like baud rates, framing and so on directly. These need not be set in script files (on Linux / wine platforms) nor in Windows COM port settings in device manager. If you find a situation in which arrayForth is unable to set the COM port up properly, please inform Customer Support right away.

If you plan to use COM port(s) to communicate with a GreenArrays chip or with the Eval Board, you will need to identify the COM port(s) in question. This procedure is designed for use with the Evaluation Board; to use arrayForth to interact with our chips on other boards, you will most likely be using some single USB to RS232 adaptor and so some of the steps may be omitted. *We recommend FTDI based communication devices for all communications with our chips, whether they be RS232 adaptors or embedded chips such as those used in our Eval Boards. This is because the*

FTDI drivers appear to be free of some common bugs in FIFO management that can create havoc with automated communications.

1. The EVB002 board has three USB connectors identified as A, B and C reading left to right along the top left corner of the board. Port A is by default used for IDE on the Host chip, port B is used for a serial terminal on polyFORTH or eForth, and port C is used for IDE on the Target chip. The following steps should be taken for each of these three independent USB to serial adaptors. If instead you are configuring your installation for use with one or more actual USB to RS232 adaptors, perhaps for use with other boards than the EVB002, you will need to take these steps for each of those adaptors and should do them completely with one adaptor at a time. If you intend to be using several adaptors it is a good idea to leave the ones you have already configured plugged in while configuring the next, so that the system does not try assigning all of them the same COM port number, as some do!

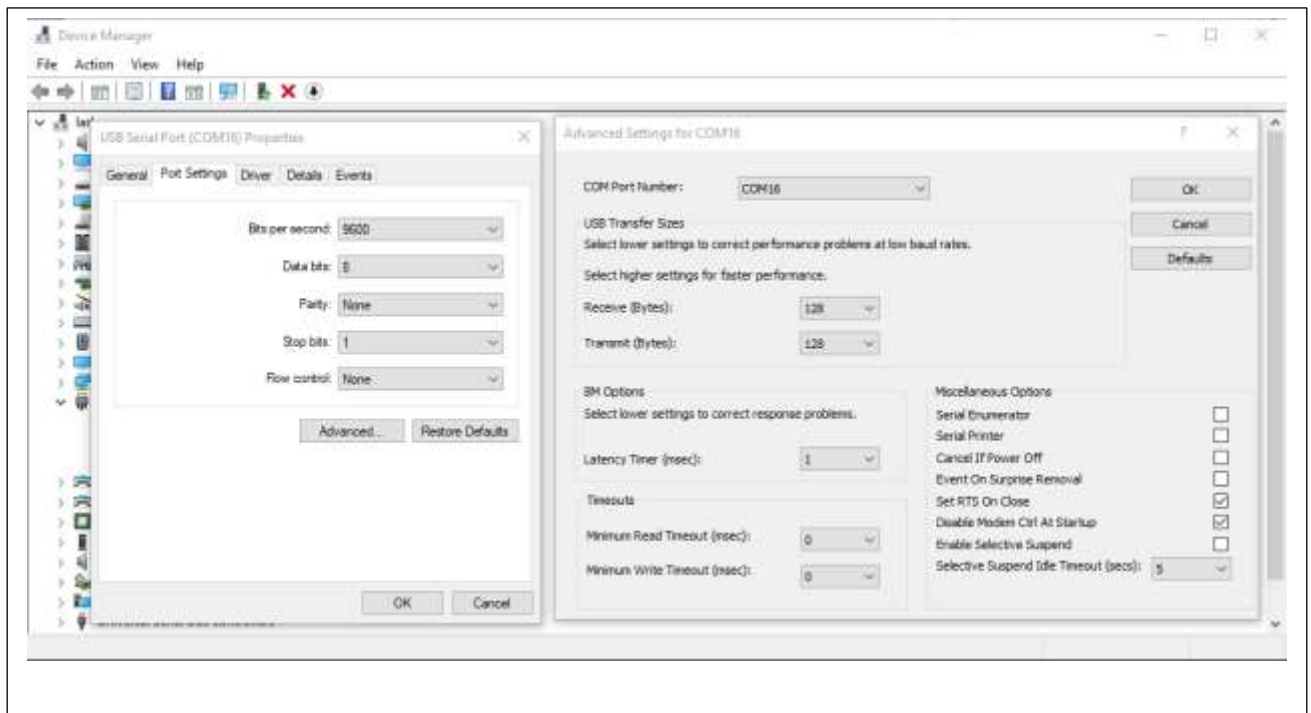
- a. Obtain Windows drivers for each USB serial adaptor you intend to use and appropriate to the version of Windows you are running. Although it is generally recommended to install drivers before first mating the device with the computer, Windows 7 seems to have relaxed this requirement. Nevertheless, even on Windows 7 you may have to connect to Windows Update or to go to the manufacturer of your adaptor for drivers.
- b. Plug in the USB to serial adaptor (or Eval board port) and go through the procedure for installing a new Windows USB device. This can vary from trivial to a grueling hassle. Eventually your device will be "ready to use."
- c. Find the Windows Device Manager. The exact procedure varies between versions of Windows but in general you can get there by right clicking "My Computer" desktop icon (or corresponding clickable area of Start menu) and selecting "Properties". This should take you to a display on which you may select "Device Manager", or perhaps "Hardware" and then "Device Manager". You may also find it in the "Control Panel." Persist until you are looking at a tree view, like this one, listing various classes of devices. Open the tree view for "Ports" which include COM and LPT port names. Your USB port should be listed here. If there are more than one and you are uncertain which is your new port, unplug the device while watching this display. It should update and the device in question should have disappeared. Now plug it back in and record its COM port number and which of your interfaces (such as Eval Board port A/B/C) it represents.



Regarding this relationship: Some USB devices have internal unique serial numbers and some do not. The default behavior of later Windows systems is to make a permanent association between a given vendor/device/serial number and a given COM port, regardless of which USB connection it is plugged into. This is a *good thing*. For devices lacking a serial number, Windows tends to assign the COM port to that vendor/device on a particular USB socket. Thus the COM port number can be expected to change if you move the same USB to RS232 adaptor from a direct USB socket on the computer to one on, say, a docking station. This is the stuff of madness but you will need to cope with it if you use such interfaces.

- d. Having identified the new port, double click it in the Device Manager and you should get a USB Serial Port Properties dialog box. On the front page / first tab you will see the manufacturer ID.

- e. If the manufacturer is FTDI, you have some things to do in order to optimize this port. On the Port Settings tab you should find an "Advanced" button and on pressing it should get a separate dialog box entitled something like "Advanced Settings for COMxx". If this is the case, please refer to the screen image below and do the following:
 - i. In the "USB Transfer Sizes" section, the default settings for Receive and Transmit are 4096 bytes. Reducing these to 128 bytes seems to improve performance in our uses.
 - ii. The next item called "Latency Timer (msec)" which defaults to 16. We get the best results in IDE and serial terminal performance with this value reduced to 1 ms.
 - iii. Please see "Miscellaneous Options." We get the best results from setting these as shown. Instructions in app notes relative to the EVB001 will assume you have made these same settings.



2. Decide what baud rate(s) you will employ (but do not enter it in the above property sheets). What is feasible depends on the electrical characteristics of the PC serial interface in use. Our default baud rate, 921600, is correct for the EVB002 USB ports (all factory testing of each Eval Board is done at this speed.) If you are using an actual RS232 interface then the maximum usable rate will depend on the quality of its chipset including its line transceivers. We have on rare occasion found RS232 interfaces that we could use at 460,800 baud, but more commonly the limit will be 115,200 or if you are lucky 230,400. Don't go any lower than 19,200 baud because our async nodes only have 18 bit time delay counters...
3. Start arrayForth 3 using the shortcut made by the installer. Say AFORTH CONFIG LIST and edit the port number for A-COM (IDE for host or only chip) or C-COM (target chip on Eval Board) as appropriate to reflect the ports(s) you have identified. Make sure the new port number is in yellow. Similarly, edit desired baud rates for these ports A-BPS and C-BPS if necessary; see above regarding appropriate baud rates to be used. Say FLUSH .
4. Shut down arrayForth 3 by saying FLUSH and RELOAD, then getting back to AFORTH .
5. You are now ready to begin IDE operations.

You may install multiple copies or versions of arrayForth 3 on the same machine; just make a separate directory for each, and name your shortcuts to avoid confusion. This may be useful if you are working with multiple chips, or even multiple projects.

12.2.2 Windows 8 and 10 Installations

Our software is designed to run on a desktop computer, not on a telephone. It is assumed that you will be using the desktop interface to Windows, and that you will have installed one of the after-market programs such as StartIsBack that restore easy access to the start menu directories which are maintained in Windows 8 and 10 but are hidden by the asinine configuration in which those systems are shipped. Contact our customer support hotline if you need advice on how to go about this.

12.3 Running arrayForth 3

Once you have taken care of the above chores, running arrayForth is simple:

1. Click the shortcut for the desired copy of the software.
2. When 25x80 blue screen screen appears, the system is ready to use. If it's the panic version, you will need to start with HI. For the G144A12 version, you will come up in AFORTH (HI AFORTH already executed) although if you find this consternating you may delete AFORTH from the shortcut. For the "with GLOW" version, HI has already been executed but AFORTH and GLOW must be loaded manually.
3. To exit arrayForth normally, simply say BYE to the interpreter.

You may run multiple copies of arrayForth so long as you don't exceed the virtual memory capacity of the operating system you are using. This can be convenient when talking to more than one chip. For example, in one step of our factory testing we run SELFTEST on both host and target chips simultaneously using two arrayForth sessions. Be careful to only edit your source base in one instance if they both use the same directory, otherwise FLUSH operations will overwrite one another.

12.4 Installing arrayForth 3 on an EVB002 Board

The above procedure has placed the x86 portion of arrayForth 3 on a PC or similar platform. To install the polyFORTH/144 portion on an EVB002 (or EVB001), you will need to duplicate steps we have already taken when commissioning your board at the factory. This factory procedure is documented in internal document PCBCOM, which is available on our website. For more information about that process, and for other options, please see section 9.3, G144A12 polyFORTH Capable, as well as section *Loading or Installing polyFORTH* in DB006, *GA144A12 Supplement to polyFORTH Reference*.

12.5 Installing arrayForth 3 on an EVB001 Board

Because the nucleus as distributed is configured for new-style flash, you will need to do some additional work, as indicated in the section of this manual 9.3 referenced above so it can write the old one on the EVB001. If you are planning to use AFORTH on the pF system you will need to accommodate the smaller flash; most likely you'll choose to use serial disk for development. If you wish to use the flash as your primary working image with arrayForth 3, you will need to reorganize the disk so that the aF3 tools fit into the small image. We don't really recommend doing this because it really does not fit, and the aF3 source code will surely grow in the future. Realizing the general futility of such a reorganization, we chose not to attempt it in the delivered system. In fact, it was the desire to enable the EVB to act as its own development system that motivated us to put the largest 24-bit addressable flash on the EVB002 that we could.

12.6 Customizing for a Project

We recommend that you keep your application code separate from the system code we distribute so that it may be backed up independently of the system, so that your application need not be merged into new system releases and need not move when we reorganize the system distribution. In the default distribution, we provide the files **Project** and **Projback** in the /sF directory, mapped into 4 DRIVE on both sF disk and pF serial. **PROJECT** is provided as a CONSTANT whose value is the first block (4800) of this area. These files are among those backed up into EVB002/backup during installation if you request it.

12.6.1 Managing Multiple Major Projects

As of rev 03c we've provided a more powerful method that we use ourselves. All the code, data, documentation, notes, etc. for each project are organized into a separate directory along with the shortcut to run aF-3 set up to work on that project. If your plans include this sort of activity, we recommend you consider adopting our method.

A new directory, **EVB002/Projects** organizes multiple projects, each of which has a subdirectory. This entire directory is backed up into EVB002/backup during installation if you request it. The default subdirectory (or project) is named **DEFAULT** and contains at least the following necessary files:

- **DEFAULT.src** 4800 blocks to be mapped into 4 DRIVE (4800) on both sF disk and pF serial.
- **DEFAULT-back.src** 4800 blocks to be mapped into 24 DRIVE on sF disk to be used as a backup (by sF) for DEFAULT whenever you wish to reconcile your changes into it.
- **OBJ-DEFAULT.src** Object code you compile for this project. Supplied with nothing but the **STOCK** code, distributed with aF-3, such as ROM and common tools. (**AFORTH STOCK LIST** shows the load block for assembling this object code.)
- **custom.txt** A text file of sF source that may be INCLUDED to configure sF for this project.
- **arrayForth 3 DEFAULT** A shortcut to run the appropriate sF executable, using EVB002/sF as working directory, and to interpret **custom.txt** from this directory via interpretive command line code at boot time.

To make a new project, we recommend you make a new directory EVB002/Projects/<newname> and copy the files from DEFAULT into it. Then rename four of the files, replacing DEFAULT with <newname>. Edit the properties of the shortcut; the "Target" (command line) initially reads

```
C:\Xga\arrayFORTH-3\sF\sF5b2-af3.exe INCLUDE ../Projects/DEFAULT/custom.txt
```

so you replace DEFAULT with your <newname> directory (project) name.

Finally, edit the text file custom.txt, interpreted immediately by the INCLUDE statement on the above command line that's evaluated immediately after block 9 is automatically loaded. The default file is as follows:

```
( Customize aF3 for Default project.)
( All files are in af3/Projects/DEFAULT
  or EVB002/Projects/DEFAULT.)

EMPTY ( 4800 CONSTANT PROJECT) FORTH GILD

4740 4740 u.O_RDWR u.O_BINARY + 1
      CALLED" ../projects/DEFAULT/DEFAULT.src"
4740 4740 u.O_RDWR u.O_BINARY + 5
      CALLED" ../projects/DEFAULT/DEFAULT-back.src"
4800 4800 u.O_RDWR u.O_BINARY + 9
      CALLED" ../projects/DEFAULT/DEFAULT.src"

4800 4800 u.O_RDWR u.O_BINARY + 11
      CALLED" ../projects/DEFAULT/OBJ-DEFAULT"
```

```
8 ( Enable) 4 8 +WRT  
AFORTH
```

That directory (Projects/whatever) can hold everything involved with the project including other backed up versions, test data and notes, etc. We have found it convenient to store the custom shortcut in that same directory.

13. Appendix: Apple Mac® Platform with Windows

13.1 Mac Requirements

arrayForth can be run on Apple computers that are built from Intel x86 processors, by using Parallels, a product of Elements5, hosting a Microsoft Windows® operating system (XP or later). The following procedures have been tested on a MacBook running OS X version 10.6, Parallels 5.09, and Windows XP Professional. We also believe they should work using Windows installed on a Mac system using VirtualBox; if you install our software on this configuration, please let us know how it went.

Note: If you are new to Windows, be aware that you will need to change some of the default settings in Windows before it is practical to do some of these things. For example, by default the Windows file explorer hides file extensions like .bat and .exe ... and even hides whole files. If you are in this situation and have no idea how to do those things, contact our support hotline and we may be able to supply a "script" describing how to do it on the version of Windows you have installed.

13.2 Installation

Because Parallels (and VirtualBox) are Virtual Machine environments running actual Windows software, the procedures are almost identical with those on a Windows system. The differences we know of are as follow:

1. You should download the installer executable into some path actually rooted in C: in the windows environment's file system. Working directories are not correct when Windows programs are executed from places that are not, such as [\\.\psf\Home\Documents](#) as one example.
2. When first plugging an adaptor or Eval Board port, go through the procedure for installing Windows drivers, associating the device with Virtual Machine rather than Mac OS when you are given that choice.

13.3 Running arrayForth

Again, running arrayForth is the same as on the Windows platform with these exceptions:

1. Be prepared for keyboard surprises and check the Parallels documentation to explain anomalies. For example, to use the F1 key you may have to hold down the Fn key as a shift.

14. Appendix: unix Platform including Mac OS X

If you have an Intel x86-based unix system, you may be able to run arrayForth using the Wine subsystem. Wine is an open source implementation of the Windows API that runs on BSD Unix, Linux, Mac OS X and Solaris systems. Its home page is at WineHQ.org. The following has been tested in a freshly installed Ubuntu Linux version 11.04 with a freshly installed Wine obtained from the "Ubuntu software center" as "Microsoft Windows Compatibility Layer (meta package)."

14.1 Installation

Although Wine is not a Virtual Machine environment, most steps are the same as with the Windows platform. The differences are as follow:

1. Download the Windows Installer from our website. Save it in your Downloads directory.
2. Start Nautilus, the GUI file manager program, and navigate to your Downloads directory.
3. Right click on the icon for the setup program, select "Properties", select the "Permissions" tab on the resulting dialog, and check the box to "Allow executing file as program." (equivalent to `chmod u+x`)
4. Right click again and choose the menu item "Open with Wine Windows Program Loader." The installer should run as described above.

5. In the "Select Additional Tasks" window, check the box "Linux with Wine." When you finish the process, there should be two new icons on your desktop. The first is the arrayForth icon named "arrayForth EVB001 <vers>" that starts arrayForth, and the other is a folder icon named "GreenArrays" which opens the install directory using a GUI program called "Wine File".
6. What Wine calls drive C: in this window can be found in the Ubuntu file system as `/home/<yourlogin>/.wine/drive_c`.

14.1.1 Identifying and Configuring COM Ports

1. The procedure for identifying COM ports in unix is quite different than in Windows; contact Customer Support for assistance or suggestions. If you plug USB serial devices in one at a time they will by default be assigned consecutive unix device names of the form `/dev/ttyUSBn` where `n` starts at zero. The capability of Windows to identify a particular USB device by serial number does not seem to exist. Not only does the order in which currently inserted USB devices matter, but also the amount of time a device was unplugged before it is plugged back in seems to matter as well. You will probably be best served by plugging cables into your machine in a fixed sequence and following that sequence, without skipping cables, every time. We recommend that EVB001 users make a habit of plugging the three cables for USB ports A, B, and C into the computer in that same sequence every time; if port C is to be plugged in, make sure A and B have been plugged in or C will be assigned to one of the ports normally used by A or B. Follow this procedure for each of the three USB ports A, B, and C in order:
 - a. Plug in the first USB to serial adaptor and run `dmesg` at the `bash` command line. You will see the name of the device listed near the end of the report, most likely `/dev/ttyUSB0`. If you do not see such a line, unplug the adaptor and plug it back in, waiting several seconds between steps; log writing seems to be out of phase some times.
 - b. When this has been done, do the same for adaptors or cables B and C in order, one at a time.
 - c. In the shell, navigate via `cd ~/.wine/dosdevices`. Typing `ls` there should show you `c:` and `z:`. If you have a real serial port its device name will be `/dev/ttyS0`. You will need to create a link for each of the USB ports you have identified; we recommend you use the following pattern for the commands that do this:

```
ln -s /dev/ttyUSB0 com1
ln -s /dev/ttyUSB1 com2
ln -s /dev/ttyUSB2 com3
```

You can use whatever COM port numbers you like, but 1, 2, 3 are sensible for A, B and C. The script files supplied assume you have done so. After you have done this, in future you may check which devices are currently plugged in using the command `ls -l ~/.wine/dosdevices` where color coding of device names will indicate their status.

2. Now, explore the install directory using the desktop icon. You will see a file called `linuxwine.bat` which is the script normally used by the installer in making the arrayForth icon. Open this file with a text editor. You will see the following:

```
rem used ONLY for Linux/wine desktop shortcut.
```

```
z:\\bin\\stty -F /dev/ttyUSB0 921600 -parenb cs8 -cstopb -crtscts raw -echo
z:\\bin\\stty -F /dev/ttyUSB2 921600 -parenb cs8 -cstopb -crtscts raw -echo
```

```
Okad2-41b-pd.exe
rem howdy
```

The purpose of these lines is to set the baud rate and other parameters for the IDE connection(s) to USB ports A and C on the eval board. If you have had to use some other ports than USB0 and USB2 for these purposes, edit this file accordingly. You should now be ready to restart arrayForth and communicate with your chips.

The `linuxwine.bat` file may need editing for other configurations. In addition you may need to employ other scripting options we have worked out; contact customer support for advice if necessary.

14.2 Running arrayForth

Preferably you will use the desktop shortcuts. You may choose to create scripts to run arrayForth from the BASH or other shell:

1. Navigate to the directory containing the `.exe` file, `OkadWork.cf`, and `Okad.sh`.
2. Type `./Okad.sh` after editing it to do the `stty` commands for your ports.

15. Data Book Revision History

REVISION	DESCRIPTION
161231	Initial revision from arrayForth 02c base.
170219	Version passed out with preliminary aF-03a5+
180331	Documents aF-03a6. STREAMER works in both systems and reproduces flash boot made by colorForth. Chip can burn own boot flash; sF can inject serial boot streams.
181226	Documents aF-03b1. Integrated with GLOW for a single base. Crystal oscillators in base. Various fixes.
190521	Documents aF-03b4. Prepared for EVB002 release but without released/supported SOFTSIM.
190707	Documents aF-03b5. Beta-releases for IDE test of external SRAM assembly, SOFTSIM.
190921	Improve SOFTSIM documentation. Improve disk organization with dummy index page at relative 4740 so the same file may be used as project for both sF and pF. Add cardinal direction support in Assenbler and BDL. Describe and give example of using INCLUDE in shortcut command line. Document gross disk structures including new use of Project and Projback. Extend SOFTSIM to two chips. Add procedures for commissioning and using EVB001 with aF3 level systems.
200126	Document additional details for release 03c.
200831	Add "Getting Started". Correct a literal block number.

IMPORTANT NOTICE

GreenArrays Incorporated (GAI) reserves the right to make corrections, modifications, enhancements, improvements, and other changes to its products and services at any time and to discontinue any product or service without notice. Customers should obtain the latest relevant information before placing orders and should verify that such information is current and complete. All products are sold subject to GAI's terms and conditions of sale supplied at the time of order acknowledgment.

GAI disclaims any express or implied warranty relating to the sale and/or use of GAI products, including liability or warranties relating to fitness for a particular purpose, merchantability, or infringement of any patent, copyright, or other intellectual property right.

GAI assumes no liability for applications assistance or customer product design. Customers are responsible for their products and applications using GAI components. To minimize the risks associated with customer products and applications, customers should provide adequate design and operating safeguards.

GAI does not warrant or represent that any license, either express or implied, is granted under any GAI patent right, copyright, mask work right, or other GAI intellectual property right relating to any combination, machine, or process in which GAI products or services are used. Information published by GAI regarding third-party products or services does not constitute a license from GAI to use such products or services or a warranty or endorsement thereof. Use of such information may require a license from a third party under the patents or other intellectual property of the third party, or a license from GAI under the patents or other intellectual property of GAI.

Reproduction of GAI information in GAI data books or data sheets is permissible only if reproduction is without alteration and is accompanied by all associated warranties, conditions, limitations, and notices. Reproduction of this information with alteration is an unfair and deceptive business practice. GAI is not responsible or liable for such altered documentation. Information of third parties may be subject to additional restrictions.

Resale of GAI products or services with statements different from or beyond the parameters stated by GAI for that product or service voids all express and any implied warranties for the associated GAI product or service and is an unfair and deceptive business practice. GAI is not responsible or liable for any such statements.

GAI products are not authorized for use in safety-critical applications (such as life support) where a failure of the GAI product would reasonably be expected to cause severe personal injury or death, unless officers of the parties have executed an agreement specifically governing such use. Buyers represent that they have all necessary expertise in the safety and regulatory ramifications of their applications, and acknowledge and agree that they are solely responsible for all legal, regulatory and safety-related requirements concerning their products and any use of GAI products in such safety-critical applications, notwithstanding any applications-related information or support that may be provided by GAI. Further, Buyers must fully indemnify GAI and its representatives against any damages arising out of the use of GAI products in such safety-critical applications.

GAI products are neither designed nor intended for use in military/aerospace applications or environments unless the GAI products are specifically designated by GAI as military-grade or "enhanced plastic." Only products designated by GAI as military-grade meet military specifications. Buyers acknowledge and agree that any such use of GAI products which GAI has not designated as military-grade is solely at the Buyer's risk, and that they are solely responsible for compliance with all legal and regulatory requirements in connection with such use.

GAI products are neither designed nor intended for use in automotive applications or environments unless the specific GAI products are designated by GAI as compliant with ISO/TS 16949 requirements. Buyers acknowledge and agree that, if they use any non-designated products in automotive applications, GAI will not be responsible for any failure to meet such requirements.

The following are trademarks or registered trademarks of GreenArrays, Inc., a Wyoming Corporation: GreenArrays, GreenArray Chips, arrayForth, and the GreenArrays logo. polyFORTH is a registered trademark of FORTH, Inc. (www.forth.com) and is used by permission. All other trademarks or registered trademarks are the property of their respective owners.

For current information on GreenArrays products and application solutions, see www.GreenArrayChips.com

Mailing Address: GreenArrays, Inc., 821 East 17th Street, Cheyenne, Wyoming 82001

Printed in the United States of America

Phone (775) 298-4748 fax (775) 548-8547 email Sales@GreenArrayChips.com

Copyright © 2010-2019, GreenArrays, Incorporated

